
Temporal Analysis

Introduction to Temporal Analysis and Visualization



Objective



Objective

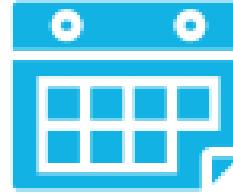
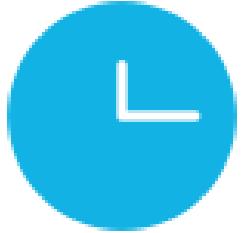
Describe temporal analysis

Temporal Analysis



“**Time** is an outstanding **dimension** reflected by Shneiderman’s Task by Data Type Taxonomy.”

Time-Oriented Data



- | Time oriented data is ubiquitous
 - Stock markets
 - Movie trends
 - Business
 - Medicine

- | Each data case is likely an event of some kind, with one variable being the date and time

Time Series



“A random selection of **4000 graphs** from 15 newspapers and magazines worldwide showed that between 1974 and 1980, **75% of these graphs were time series.**”

Time Series



| What questions can we ask of these visuals?

- Does a data object exist at a certain time?
- When does a certain data object exist?
- How long does a data object exist?
- How fast and how much does the data object change?
- What order do objects appear/disappear?
- Is there a cyclical pattern to appearances?
- Which objects exist simultaneously?

Time is....



Ordered

Continuous

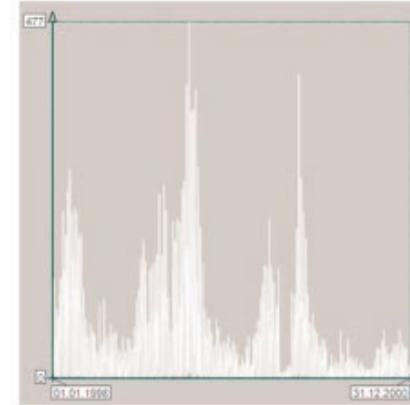
Cyclical

Independent
of location

Linear vs. Cyclical Time

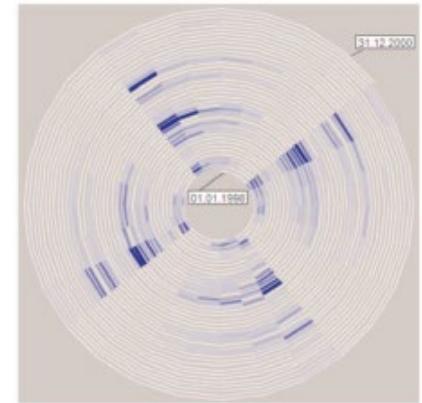
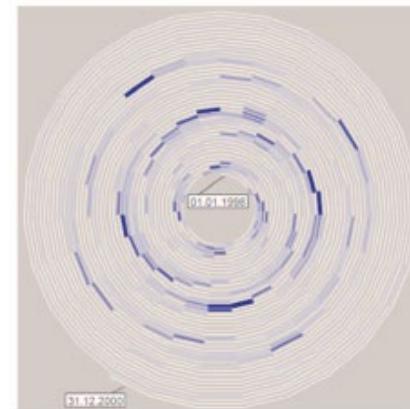
| Linear time

- One time point precedes another
- Time being ordered is closely bound to notion of causality



| Cyclical Time

- The ordering of points in a cyclic time domain would be meaningless
- Winter comes before summer, but also after summer





Temporal Analysis

Temporal Visualization

Objective

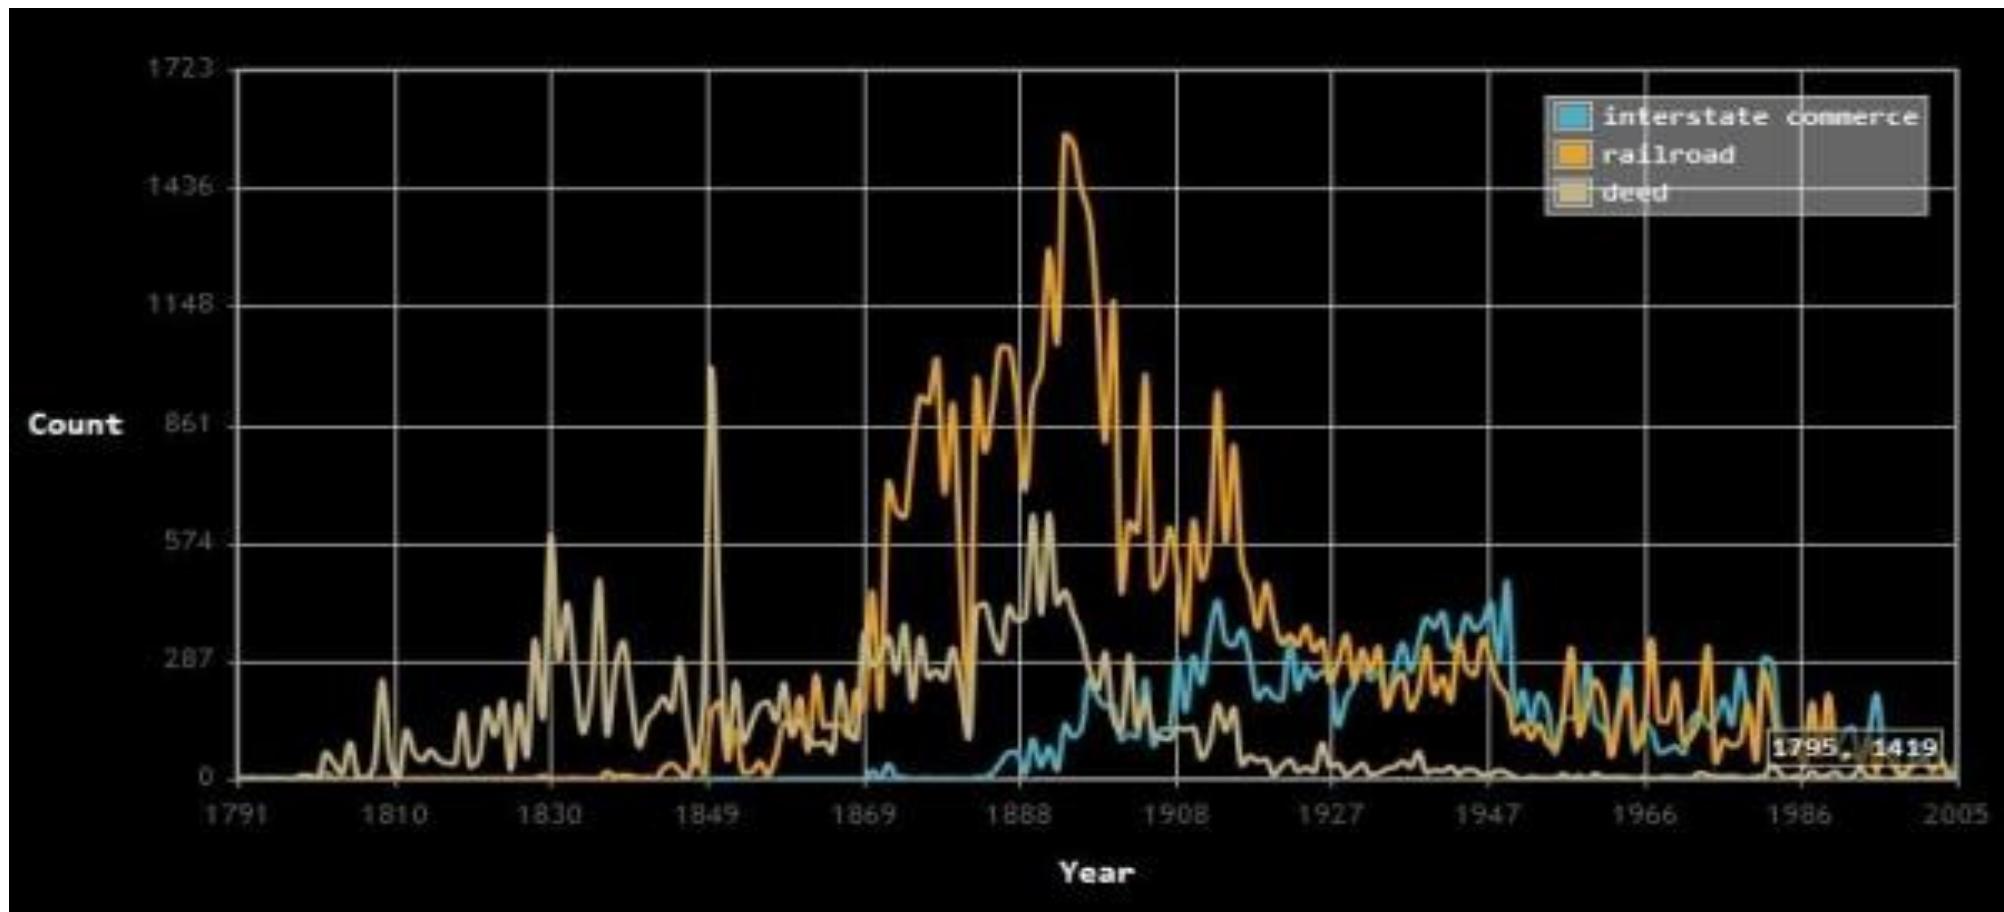


Objective

Apply methods of
temporal analysis

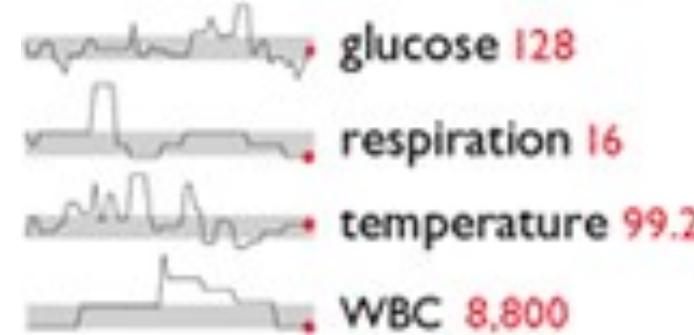
Linear Time

2D line graph with time on x-axis and value on y-axis

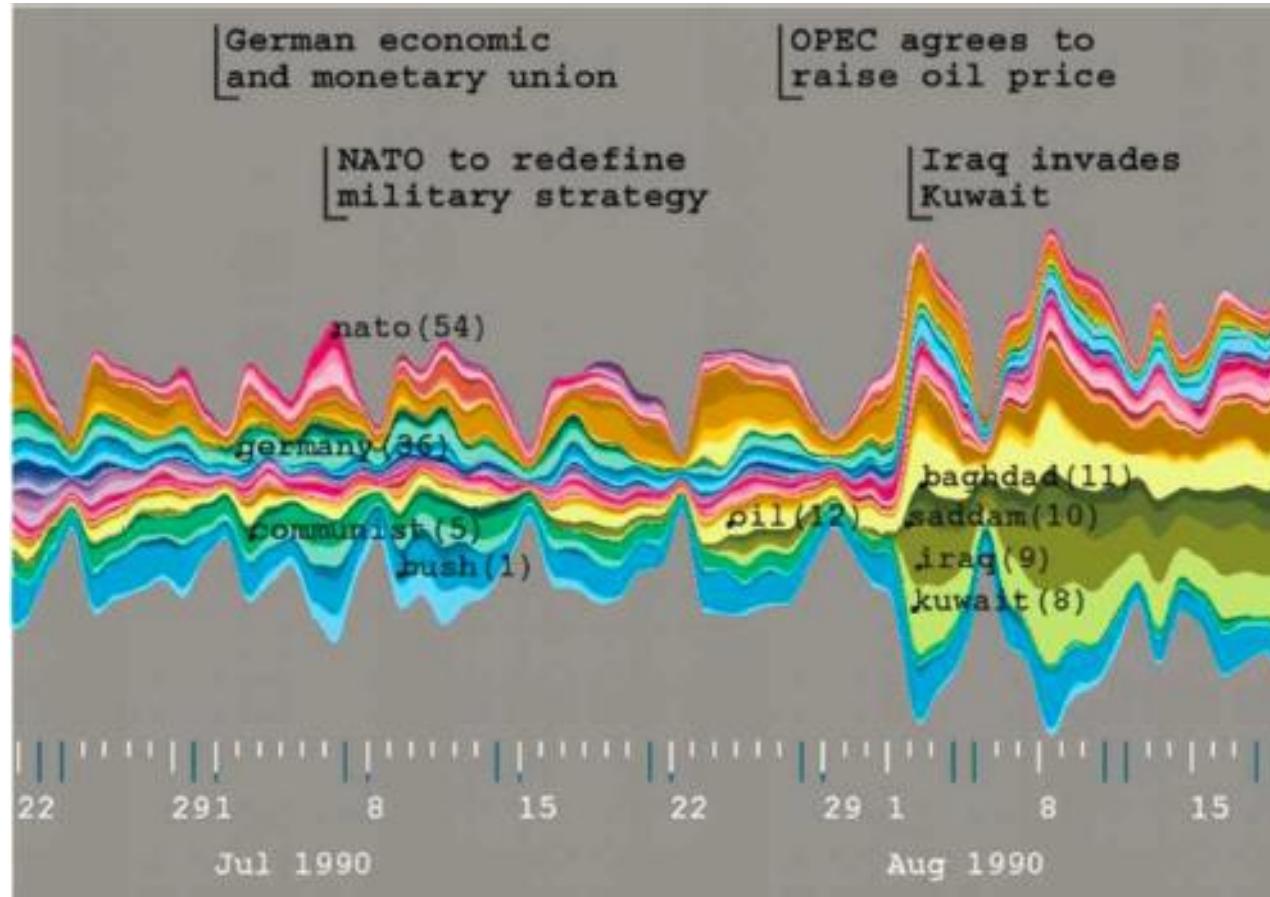


New Database Creates Time-Series Plots of Phrases in U.S. Supreme Court Opinions December 16, 2011 | Posted by: Scott C. Idleman:
<http://law.marquette.edu/facultyblog/2011/12/16/new-database-creates-time-series-plots-of-phrases-in-u-s-supreme-court-opinions/>

“Intense, simple, wordlike graphics”

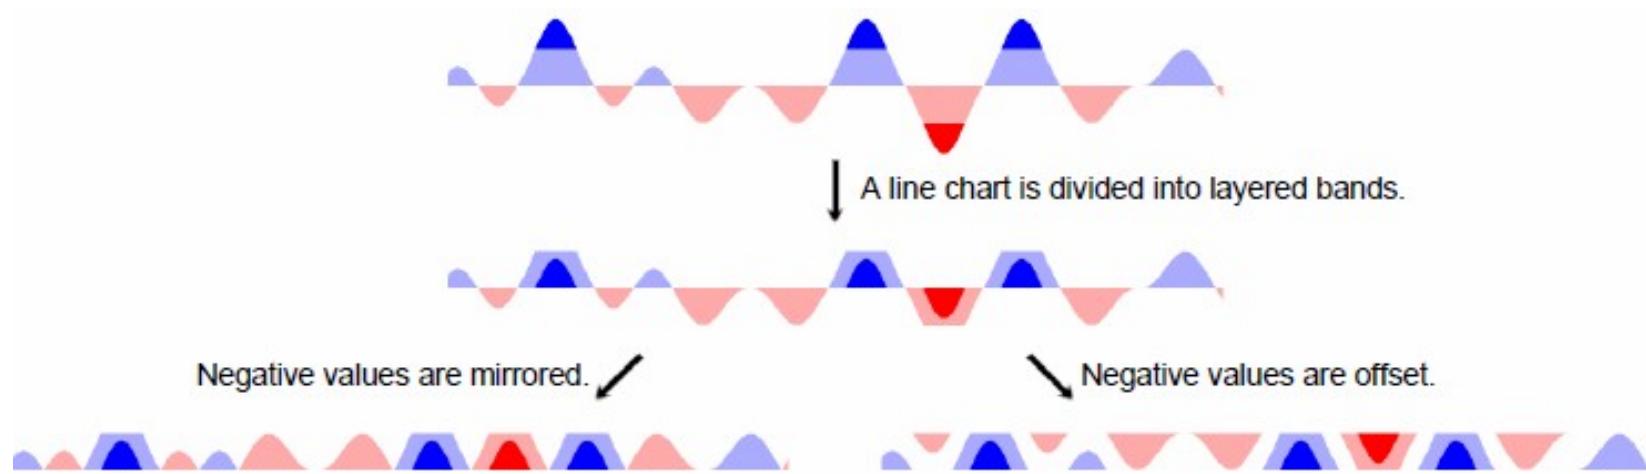


Theme River

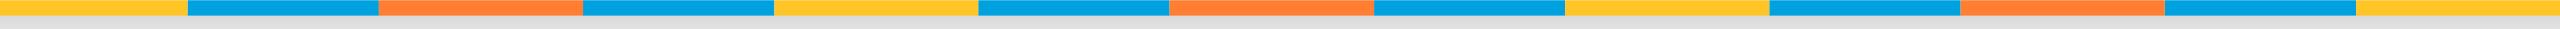


1– S. Havre, E. Hetzler, P. Whitney, and L. Nowell. Themeriver: Visualizing thematic changes in large document collections. IEEE Transactions on Visualization and Computer Graphics 8(1):9-20, 2002.* - Theme River Image from:
http://www.nytimes.com/interactive/2008/02/23/movies/20080223_REVENUE_GRAPHIC.html#

Sizing the Horizon



Linear Vs. Cyclical

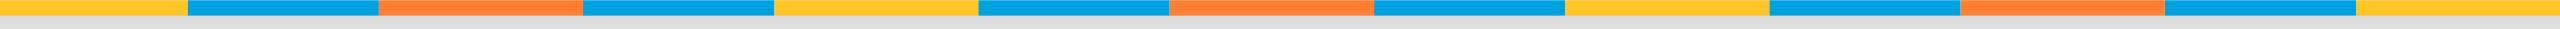


Trends are easily seen in a linear plot, but what about repeating patterns?

| Use spirals as they easily represent idea of repetition

| Finding periodicity may be easier than a bar chart

Spiral Graph



| Archimedean Spiral

$$r = a + b\theta$$

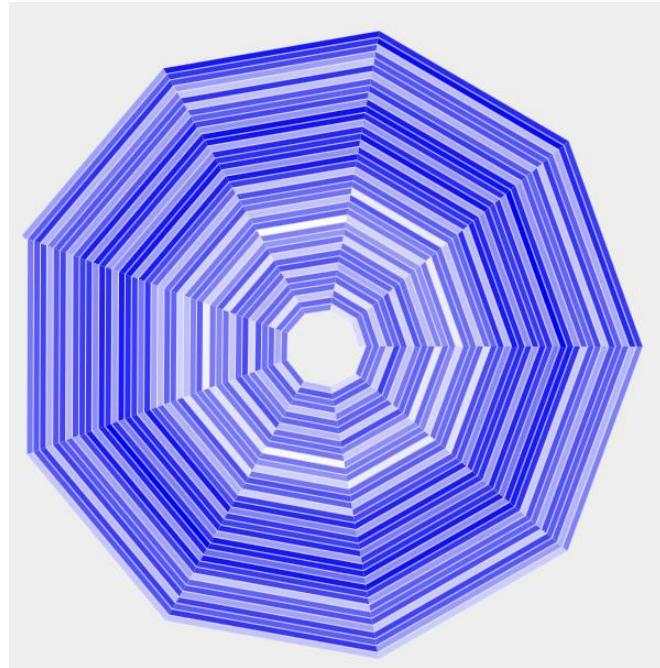
| Logarithmic Spiral

$$r = ae^{b\theta}$$

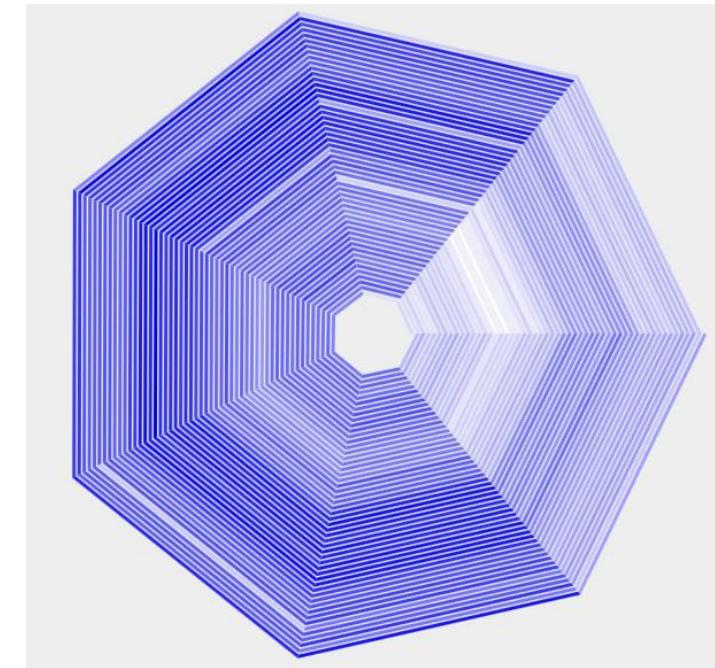
$$\theta = \frac{1}{b} \ln\left(\frac{r}{a}\right)$$

Spiral Graph

9 Day Spiral



7 Day Spiral

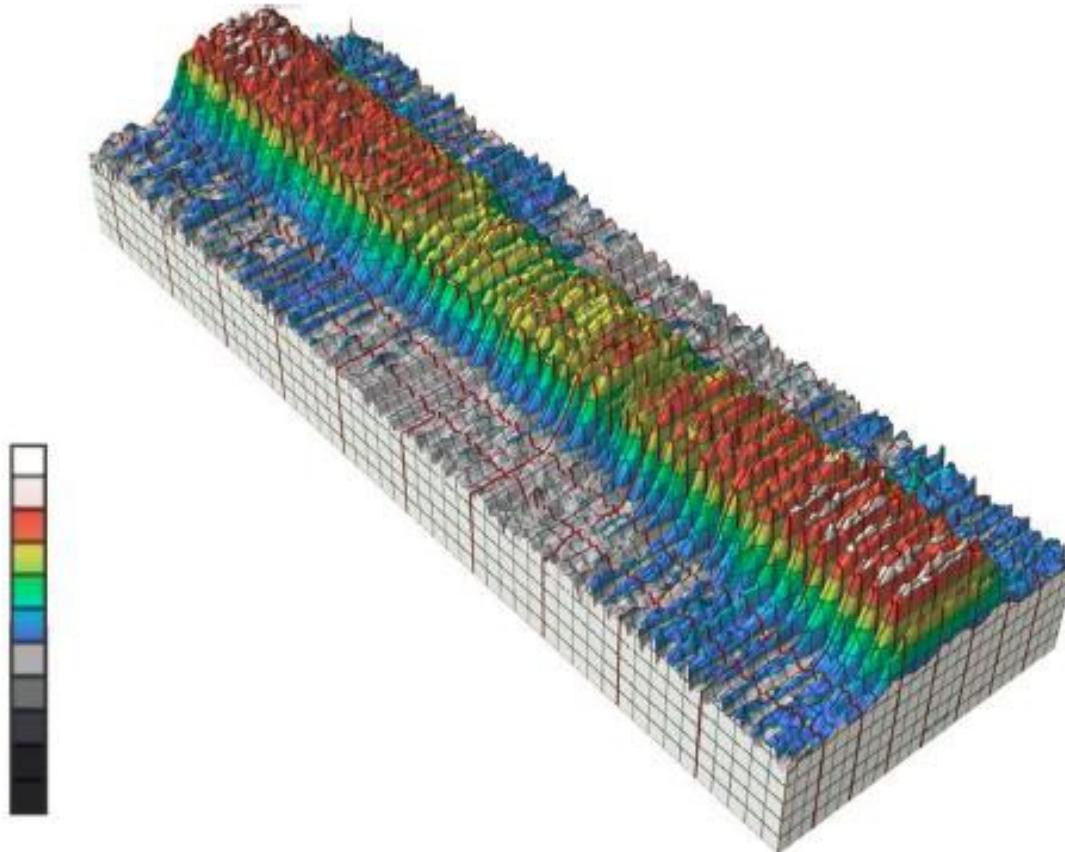


1 - John V. Carlis and Joseph A. Konstan, Interactive Visualization of Serial Periodic Data, Proceedings User Interface Software and Technology (UIST), pp. 29- 1998

2 - M. Weber, M. Alexa, and W. Müller, "Visualizing Time-Series on Spirals," Proceedings of the IEEE Symposium on Information Visualization (InfoVis '01), 7-14, Oct. 2001.

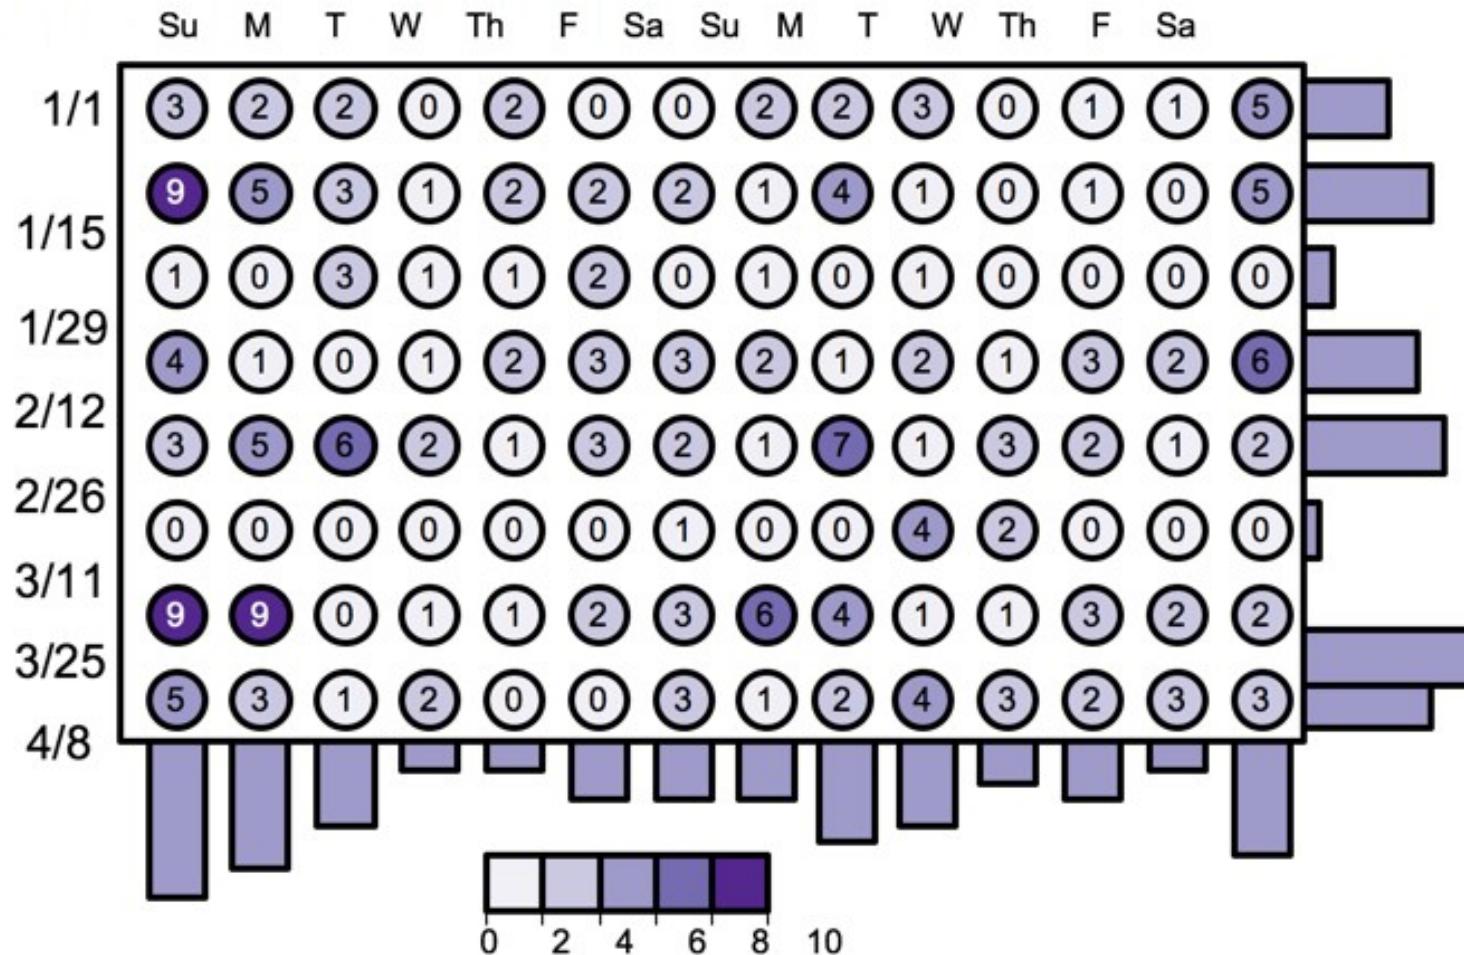
3 - Spirals for Periodic Data By Robert Kosara On August 7, 2011: <http://eagereyes.org/techniques/spirals>

Calendar Based Visualization



1 – J. J. Van Wijk and E. R. Van Selow. *Cluster and calendar based visualization of time series data*. Proceedings of the 1999 IEEE Symposium on Information Visualization, 1999.

Calendar Based Visualization



Calendar Visualization



Search Calendar

Today Feb 5 – 11, 2012 Day Week Month 4 Days Agenda

Calendar

CREATE ▾

February 2012 < >

S	M	T	W	T	F	S
29	30	31	1	2	3	4
5	6	7	8	9	10	11
12	13	14	15	16	17	18
19	20	21	22	23	24	25
26	27	28	29	1	2	3
4	5	6	7	8	9	10

My calendars ▾

- Ross Maciejewski
- Tasks

Other calendars ▾

-
- Contacts' birthdays a...
- Ross Maciejewski
- US Holidays

Sun 2/5 Mon 2/6 Tue 2/7 Wed 2/8 Thu 2/9 Fri 2/10

GMT-07 Yun Visit

9am

10am

11am

12pm

1pm

2pm

3pm

4pm

5pm

9:30 – 10:30 Mack meeting

10 – 11 Macie skills

10 – 11 VisLoh

11 – 12p Yifan

11:30 – 12 Tao

10 – 11 Jingxian

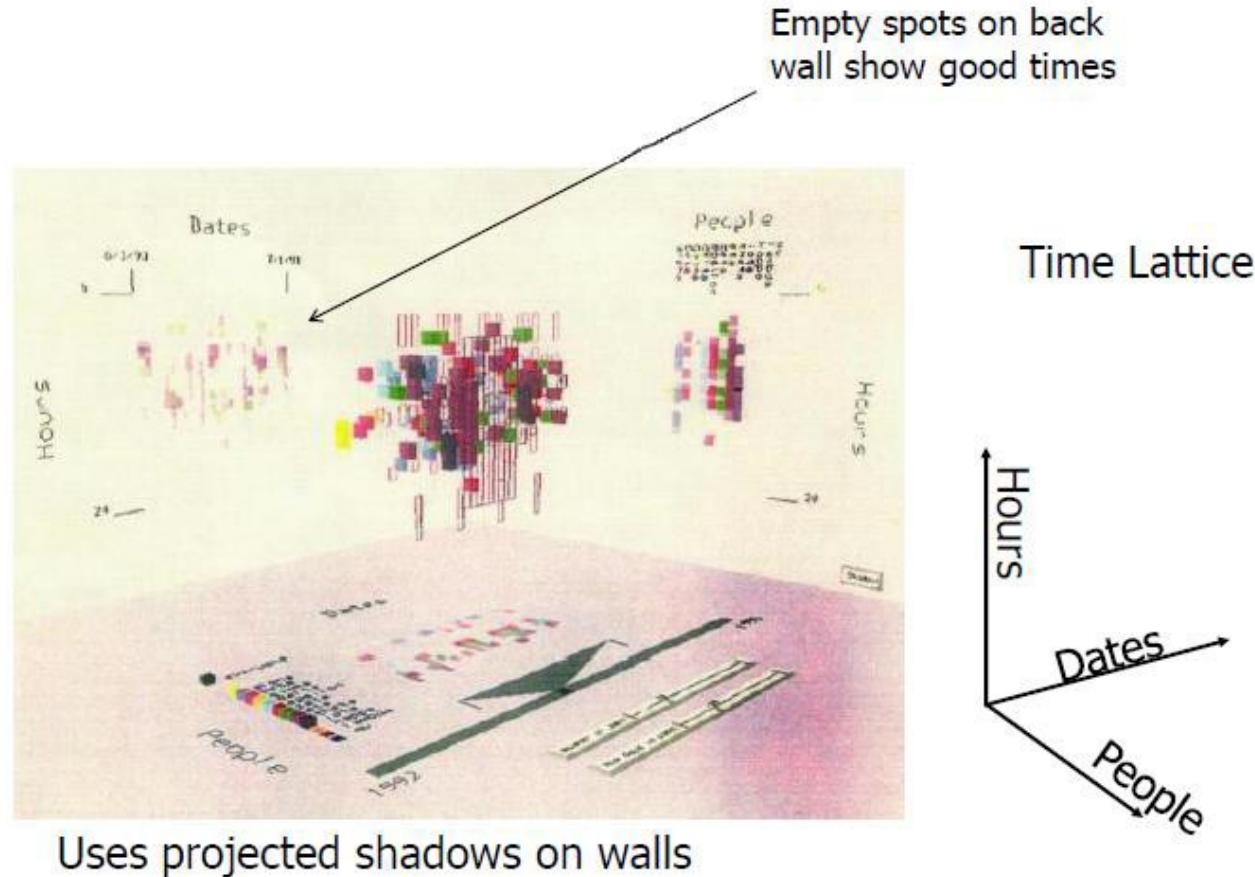
10:30 – 11 Nikunj

11 – 12p AJ

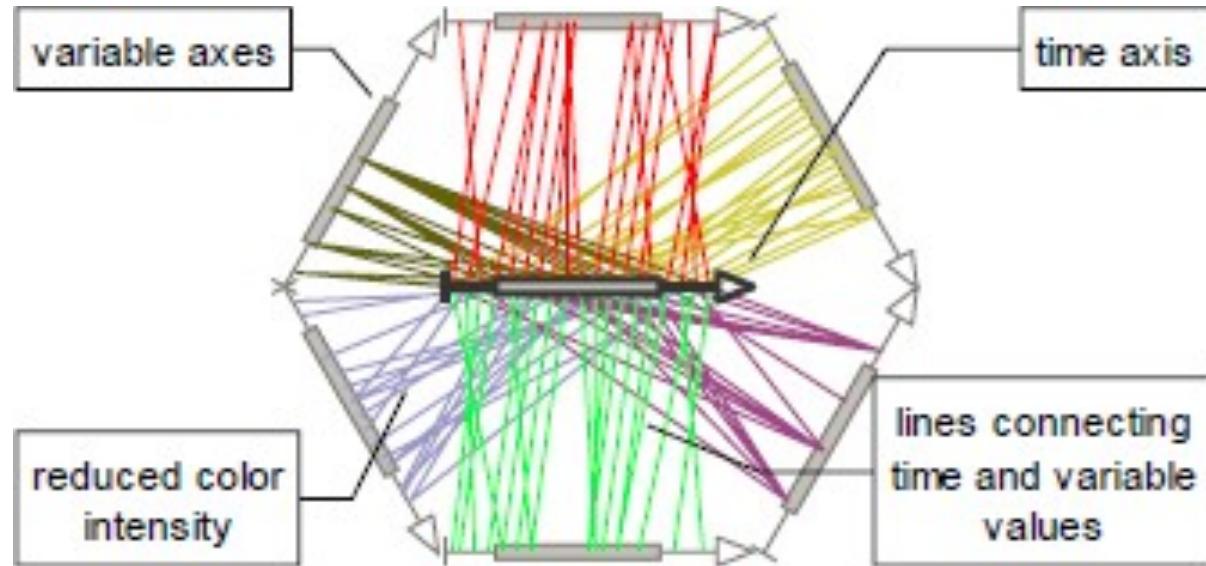
12p – 1p CIDSE Faculty Meeting (RYFNC)

1:30p – 2:30p STB1 281 F

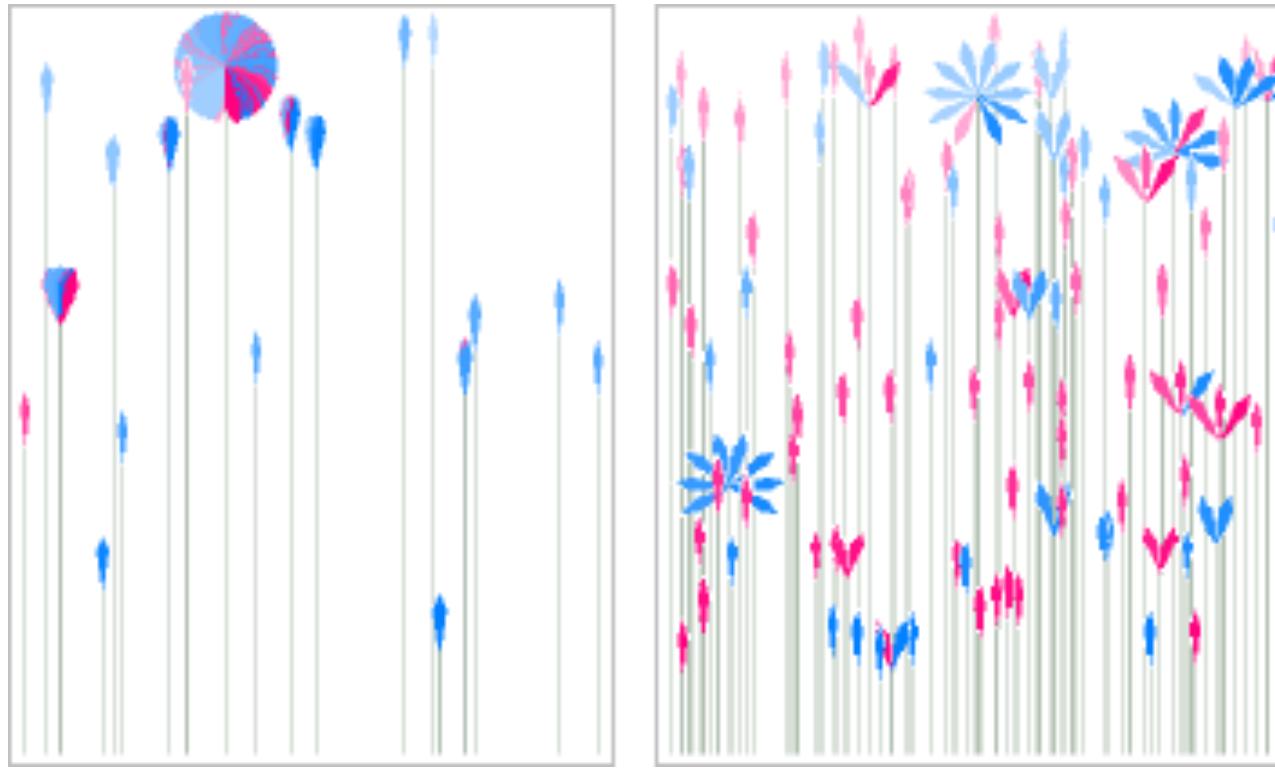
Spiral Calendar



Time Wheel

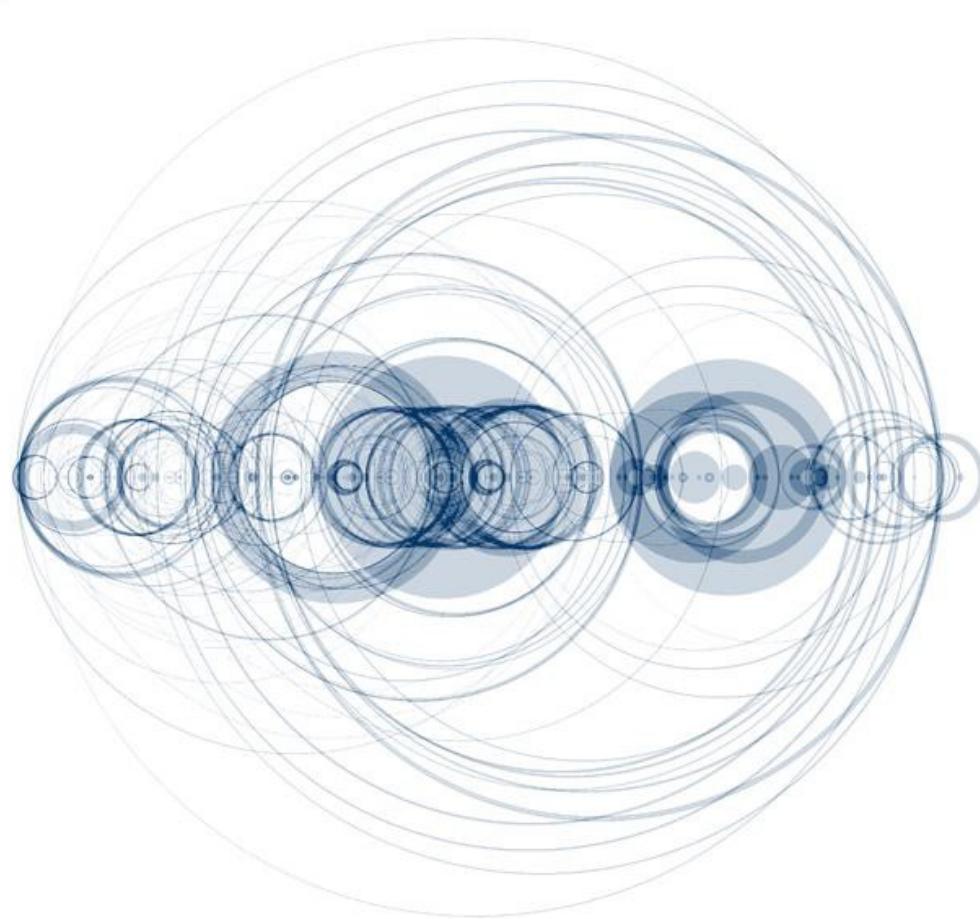


People Gardens



Rebecca Xiong and Judith Donath. 1999. PeopleGarden: creating data portraits for users. In *Proceedings of the 12th annual ACM symposium on User interface software and technology* (UIST '99). ACM, New York, NY, USA, 37-44.

Arc Diagrams



1– M. Wattenberg, “Arc Diagrams: Visualizing Structure in Strings,” InfoVis 2002. <http://www.bewitched.com/song.html>

Ordered Time vs. Branching Time



| Ordered time

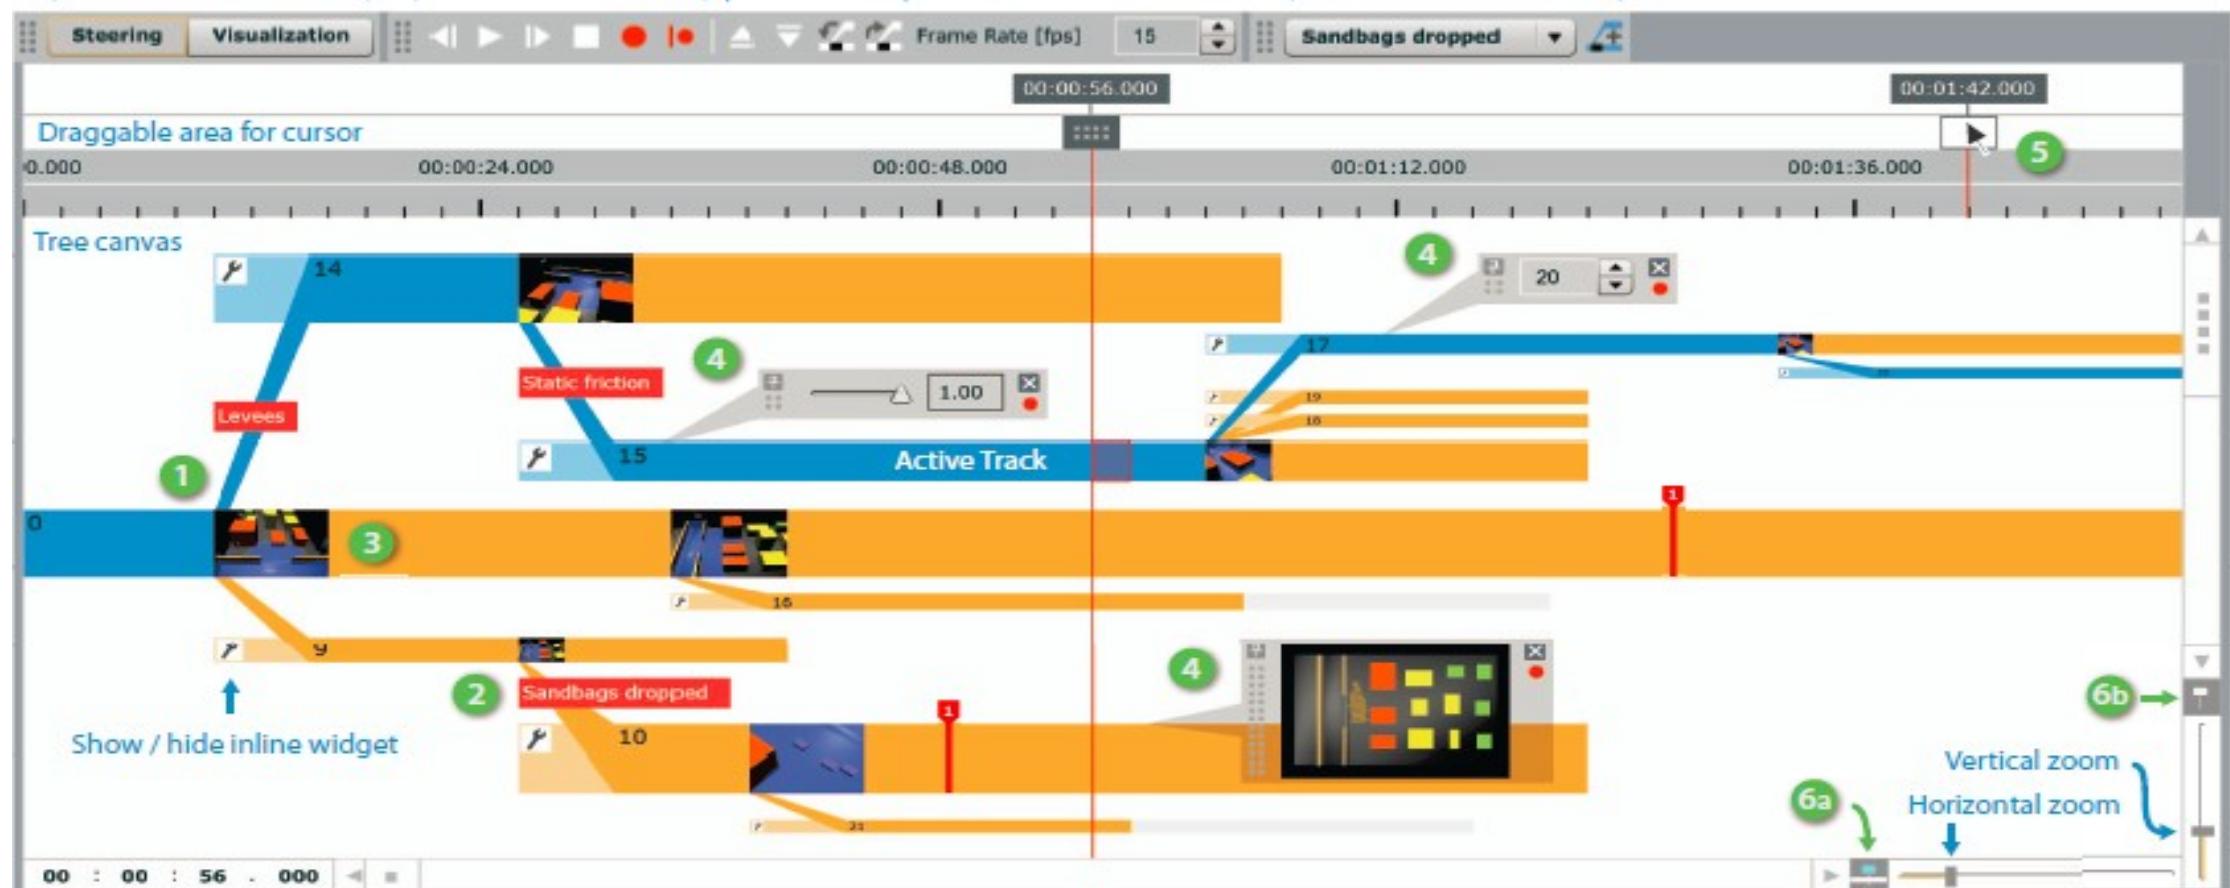
domains consider
things that happen
one after another

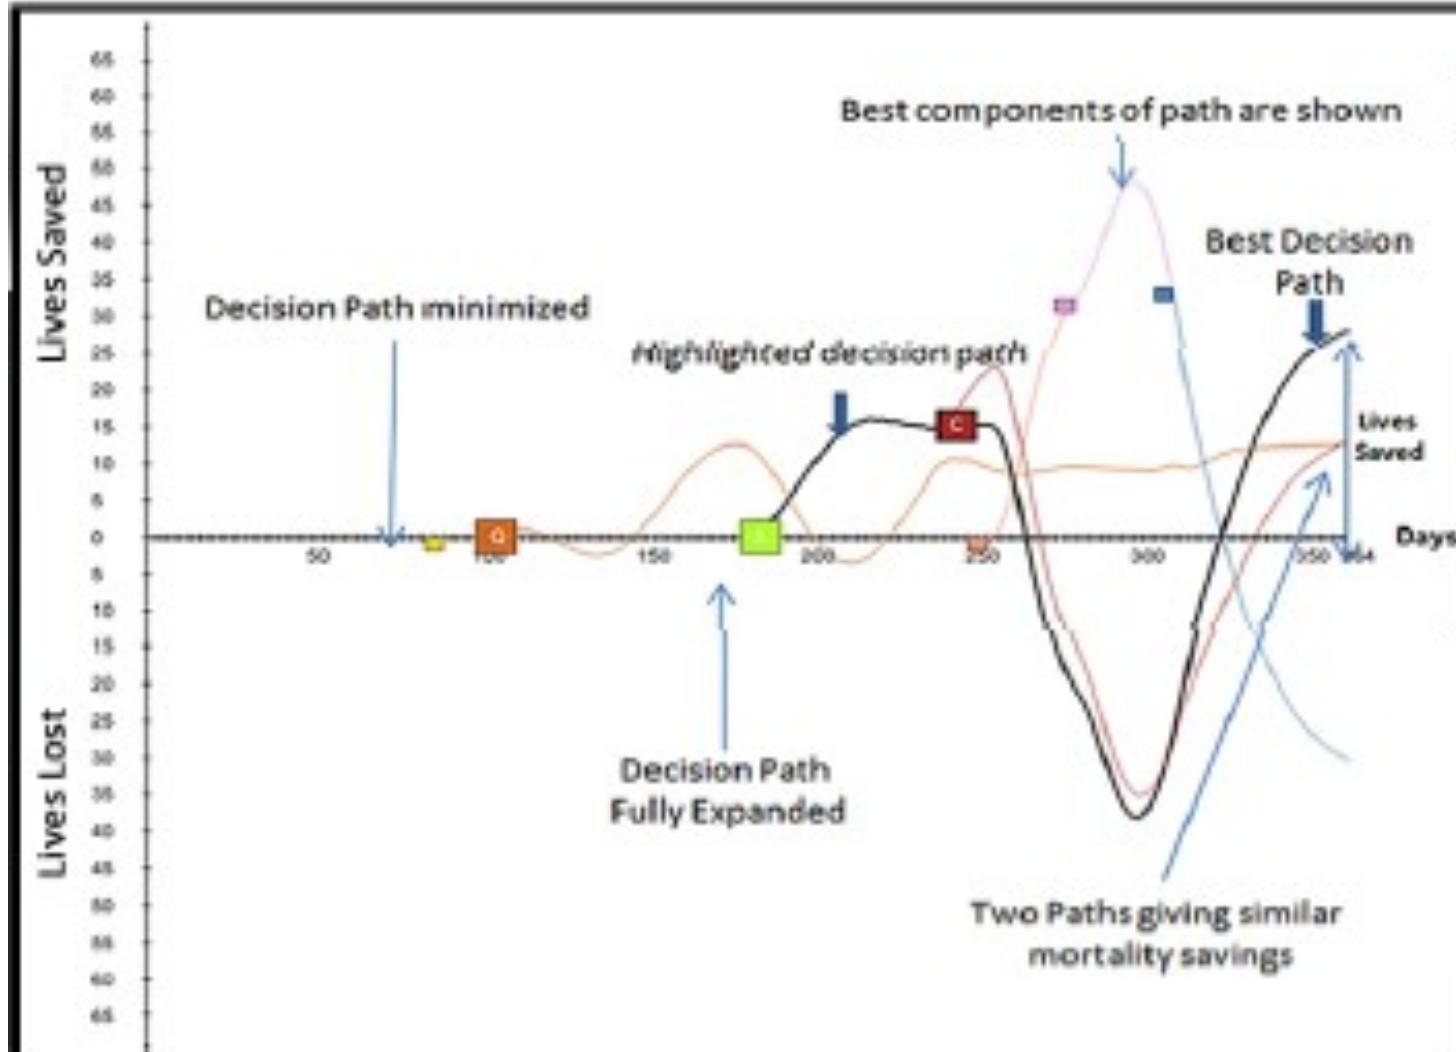
| Branching time

considers **multiple what-if scenarios**,
allowing comparison of
alternate scenarios

World Lines

Creating new simulation tracks through temporal branching





Design Principles



- | Show familiar visual representations whenever possible
- | Multiple views are more effective when coordinated through explicit linking
- | Provide side-by-side comparisons of small multiple views
- | Avoid abrupt visual change
- | Spatial position is strongest visual cue



Control Chart Analysis

Temporal Visualization

Objective



Objective

Apply methods of temporal
analysis

Data Mining



| Data mining domain has techniques for examining time series. Looking for

- patterns
- anomalies

| Enhance the visualizations

- show what is important

| Used in exploratory analysis

- “I think this looks interesting, show me similar trends.”

Typical Time Series Analysis



| Trend analysis

| A company's linear growth in sales over the years

| Seasonality

| Sales are higher in summer than winter

| Forecasting

| What is expected sales next quarter?

Control Chart Overview



Control Chart Components

- For temporal data, we can find anomalies using control chart methods
- Control charts consist of a statistic representing some measurement in time

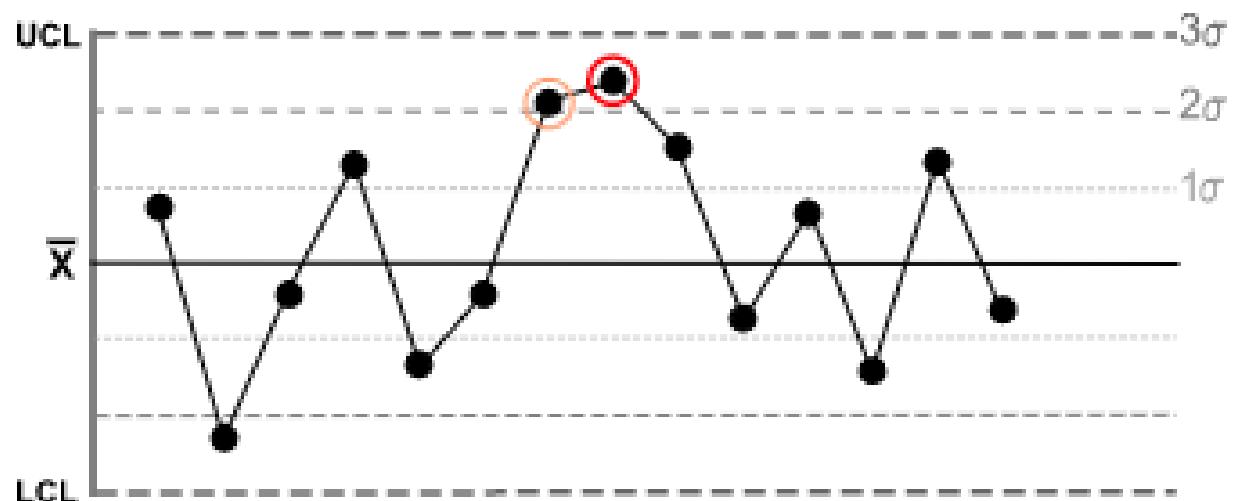
Calculations

- The mean and standard deviation of the statistic is calculated given all the available samples
- If the current value is greater than some pre-set number of standard deviations from the mean, then an alert is generated

What is a control chart?

- A graph used to study how a process changes over time
- Data are plotted in time order
- Always has a central line for average, an upper line for upper control limit and a lower line for lower control limit
- Lines are determined from historical data

Control Chart



When to use a control chart?



Controlling ongoing processes by finding and correcting problems as they occur

Predicting the expected range of outcomes from a process.

Determining whether a process is stable (in statistical control).

Analyzing patterns of process variation from special causes or common causes.

Determining whether the quality improvement project should aim to prevent specific problems or to make fundamental changes to the process.

Control Chart Model



| Upper Control Limit

$$\mu + k\sigma$$

| Center Line

$$\mu$$

| Lower Control Limit

$$\mu - k\sigma$$

Moving Average/Range Charts



Moving Average Chart

- monitors the process location over time
- generally used for detecting small shifts in the process mean
- control limits are derived from average range on Range Chart

Range Chart

- monitors the process variation over time
- should be reviewed before Moving Average Chart

Moving Average: Stock Market Closing Example



| Daily Closing Prices:

11,12,13,14,15,16,17

| First day of 5-day SMA:

$$(11 + 12 + 13 + 14 + 15) / 5 = 13$$

| Second day of 5-day SMA:

$$(12 + 13 + 14 + 15 + 16) / 5 = 14$$

| Third day of 5-day SMA:

$$(13 + 14 + 15 + 16 + 17) / 5 = 15$$

Exponentially Weighted Moving Average



| **SMA:**

10 period sum / 10

| **Multiplier:**

$(2 / (\text{Time periods} + 1)) = (2 / (10 + 1)) = 0.1818 \text{ (18.18\%)}$

| **EMA:**

{Close - EMA(previous day)} x multiplier + EMA(previous day)

The Lag Factor



| Shorter Moving

- nimble and quick to change

| Longer Lag

- Longer the moving average, more the lag

| Longer Moving

- Longer moving - slow to change

Differences between **simple moving averages** and **exponential moving averages**, one is not necessarily better than the other

Length of your **moving average** depends on your **analytical goal**

Sequences and Time Series

Introduction to Sequences

K. Selçuk Candan, Professor of Computer Science and Engineering



ASU Center for Assured and SCALable Data Engineering

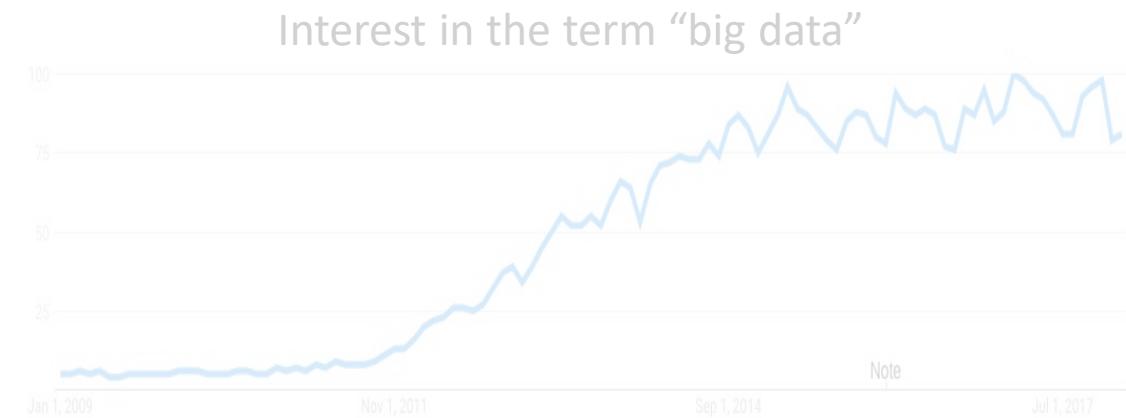


Strings, sequences, time series

| A *string* or *sequence*, $S = (c_1, c_2, \dots, c_N)$, is a finite sequence of symbols.

abcbbaabbaabcbbaaaabbc

| A *time series*, $T = (d_1, d_2, \dots, d_N)$, is a finite sequence of data values.



<https://trends.google.com/trends/explore?date=2008-12-19%202018-01-19&q=big%20data>

String/sequence matching and search

- Prefix search:
 - Find all strings that start with “tab”:
 - “table”; “tabular”; “tablet”;
- Subsequence search:
 - Find all strings that contain the subsequence “ark”:
 - “marketing”; “spark”; “quark”
 - Find all occurrences of “acd”:
 - “aabacdcdabdcababdacddcab.”
- Sequence similarity:
 - “table” vs. “cable”?
 - “table” vs. “tale”?
 - “table” vs. “tackle”?

String/sequence matching and search

- Prefix search:

- Find all strings that start with “tab”:
 - “table”; “tabular”; “tablet”;

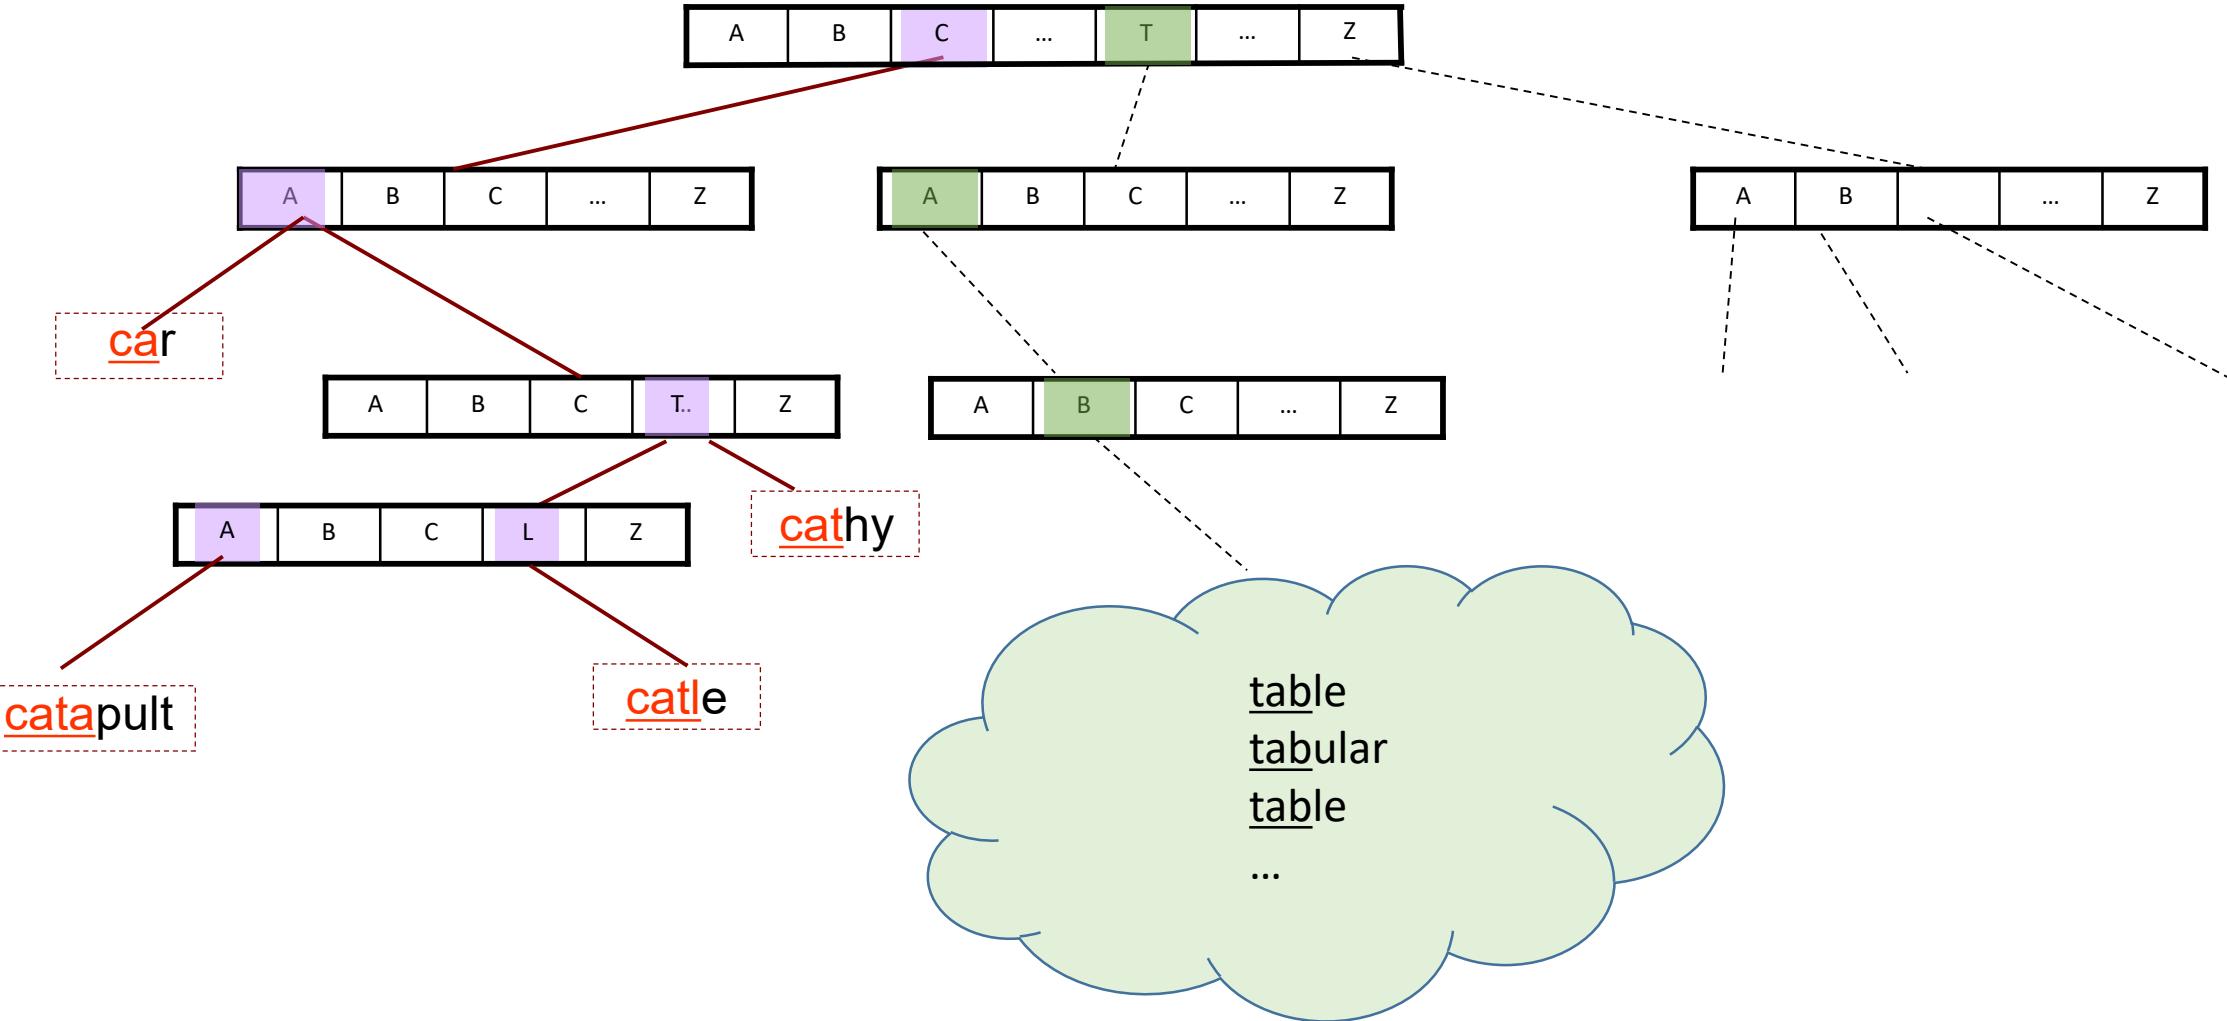
- Subsequence search:

- Find all strings that contain the subsequence “ark”:
 - “marketing”; “spark”; “quark”
 - Find all occurrences of “acd”:
 - “aabacdcdabcdabcdababdacdddcab.”

- Sequence similarity:

- “table” vs. “cable”?
 - “table” vs. “tale”?
 - “table” vs. “tackle”?

Trie data structure



String/sequence matching and search



- Prefix search:

- Find all strings that start with “tab”:
 - “table”; “tabular”; “tablet”;

- Subsequence search:

- Find all strings that contain the subsequence “ark”:
 - “marketing”; “spark”; “quark”
 - Find all occurrences of “acd”:
 - “aabacdcdabdcababdacddcab.”

- Sequence similarity:

- “table” vs. “cable”?
 - “table” vs. “tale”?
 - “table” vs. “tackle”?

Subsequence/Pattern Search

data: abcbbbaabbaabcbbbaaabbccbbbaabbaacbbbaabbccbcbbbaabcbbbaabab

pattern: abbcc

- Brute force approach:
 - scan the sequence, while aligning the pattern for each position in the sequence
- Given a sequence of length N, and pattern of length M
 - Cost: $O(N \times M)$
 - For the above example, cost: 60×5

Suffix Trees and Arrays

- Tries work well if we search for a prefix

- Suffix trees and suffix arrays

- Input text: a single long string
 - each position in the text gives a suffix

we are teaching suffix trees and arrays in the course



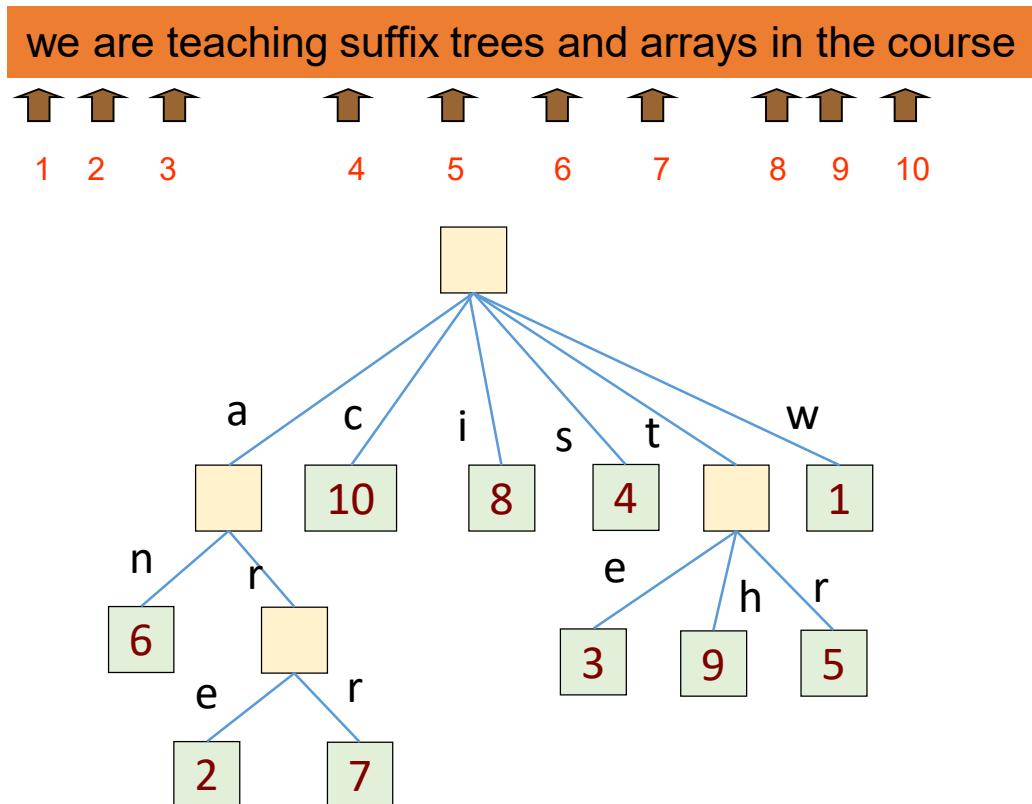
- alternatively, start of each word in the text gives a suffix

we are teaching suffix trees and arrays in the course



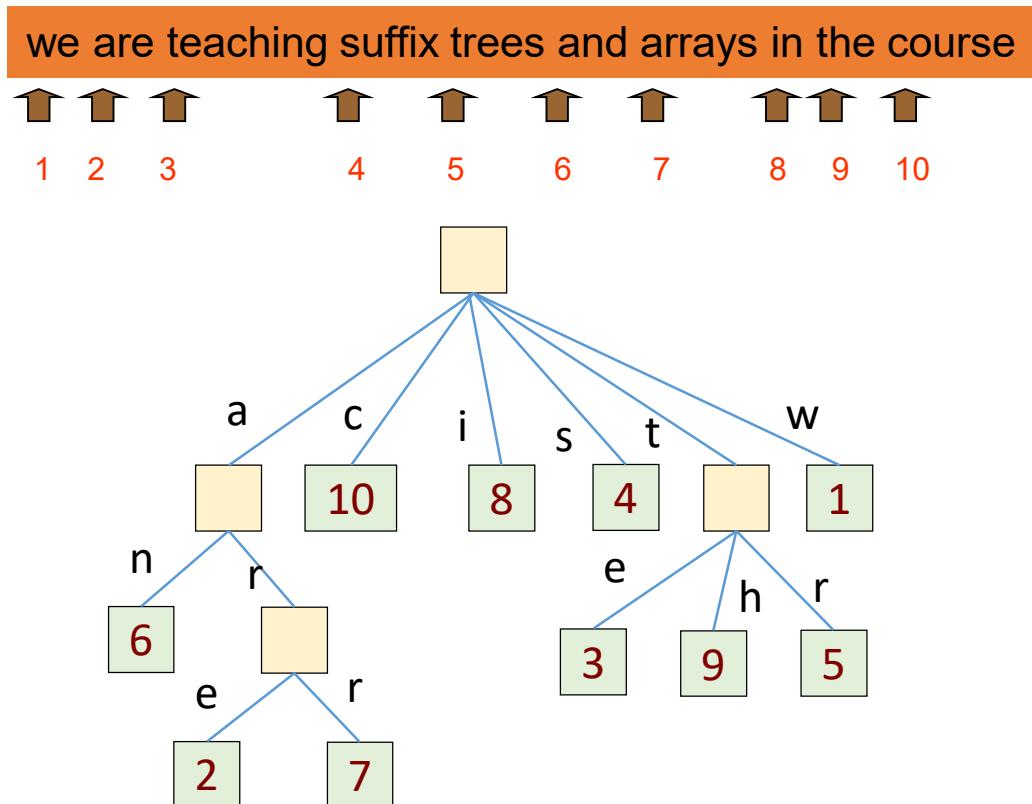
Suffix Trees

- Suffix trees
 - Input text: a single long string
 - each word position in the text gives a suffix



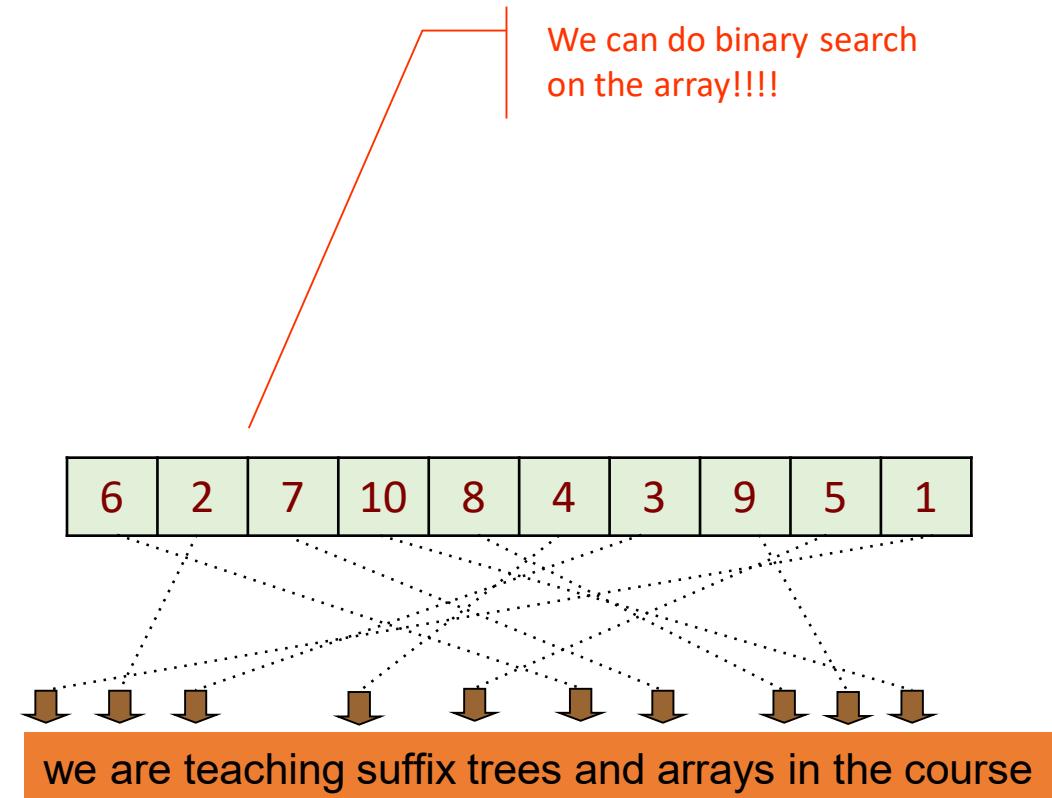
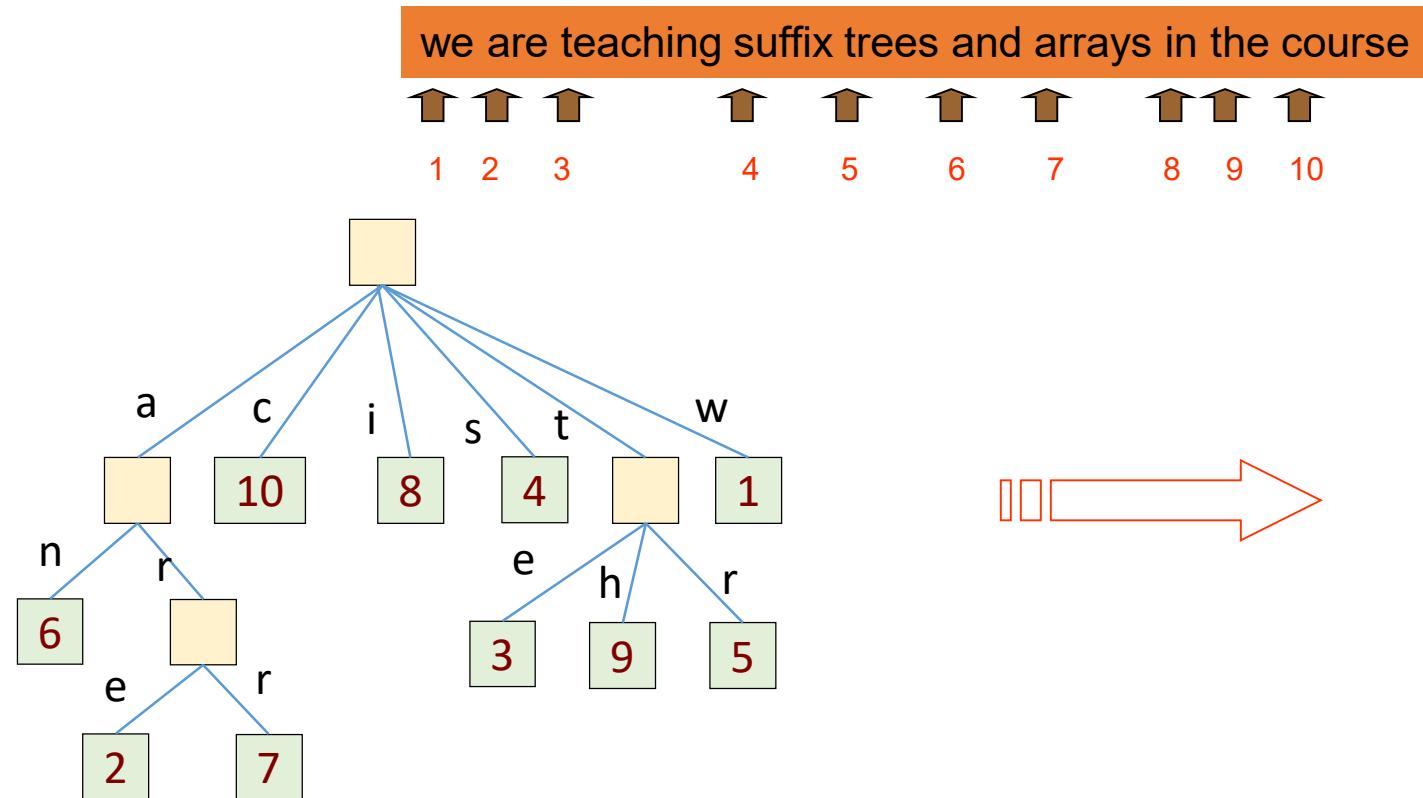
Suffix Trees

- Suffix trees
 - Input text: a single long string
 - each word position in the text gives a suffix



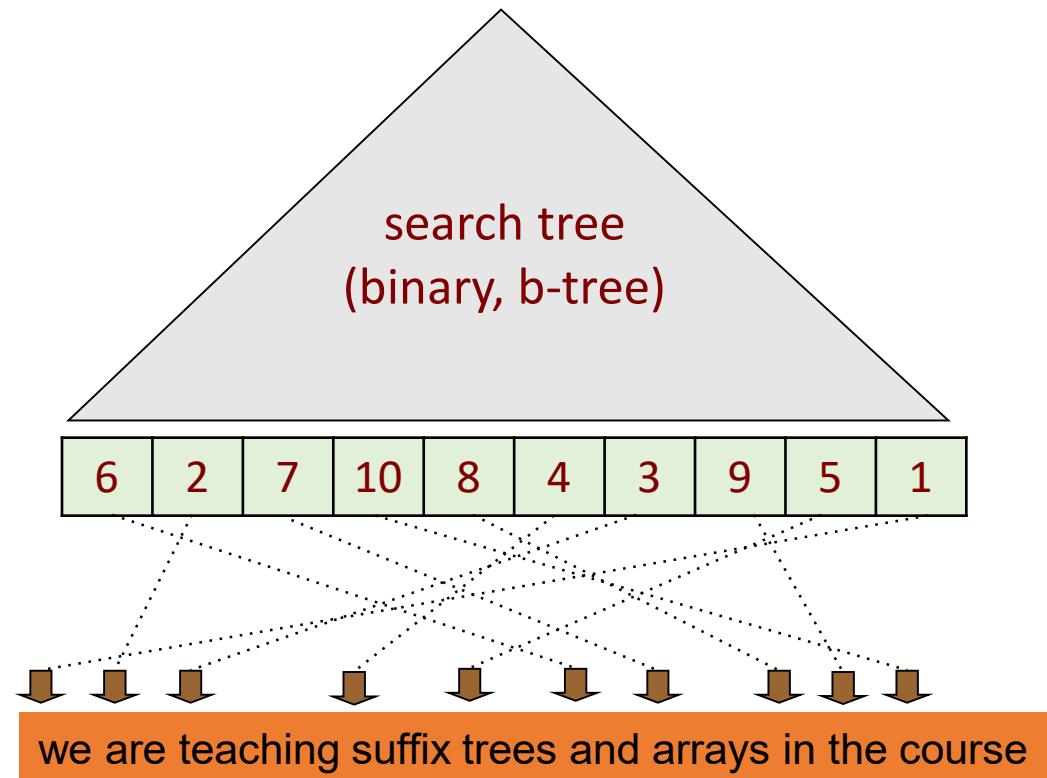
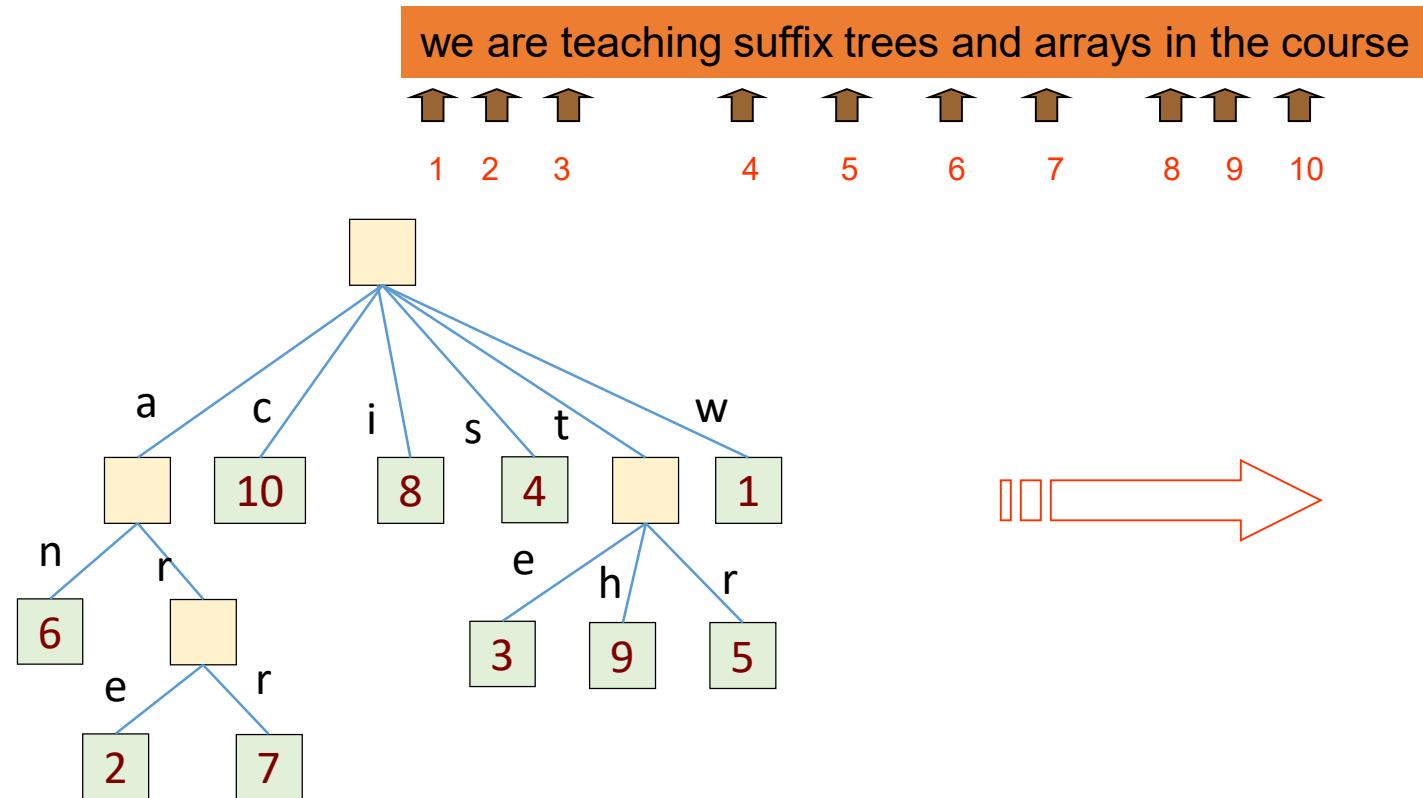
Suffix Arrays

- Suffix arrays
 - Input text: a single long string
 - each word position in the text gives a suffix



Suffix Arrays

- Suffix arrays
 - Input text: a single long string
 - each word position in the text gives a suffix



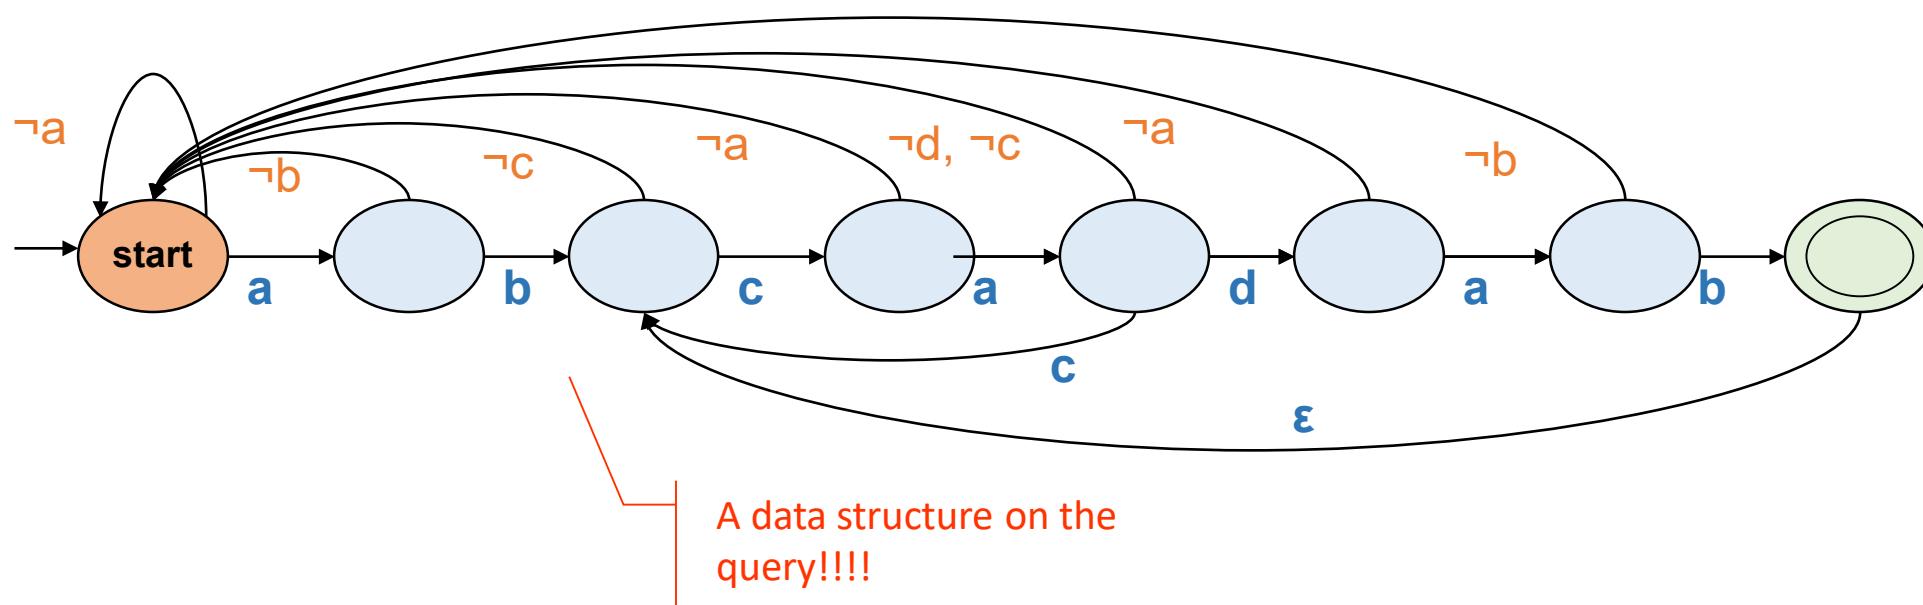
Subsequence/Pattern Search



- What if we are not given the **sequence** in advance; can we do search without a data structure on the **sequence**?
 - Yes, scan the **sequence**...
 - ..but, we have seen that this is expensive
 - Given a **sequence of length N**, and **pattern of length M**
 - Cost: $O(N \times M)$
- If we are given the **pattern** in advance, can we create a data structure on the **pattern**, instead?

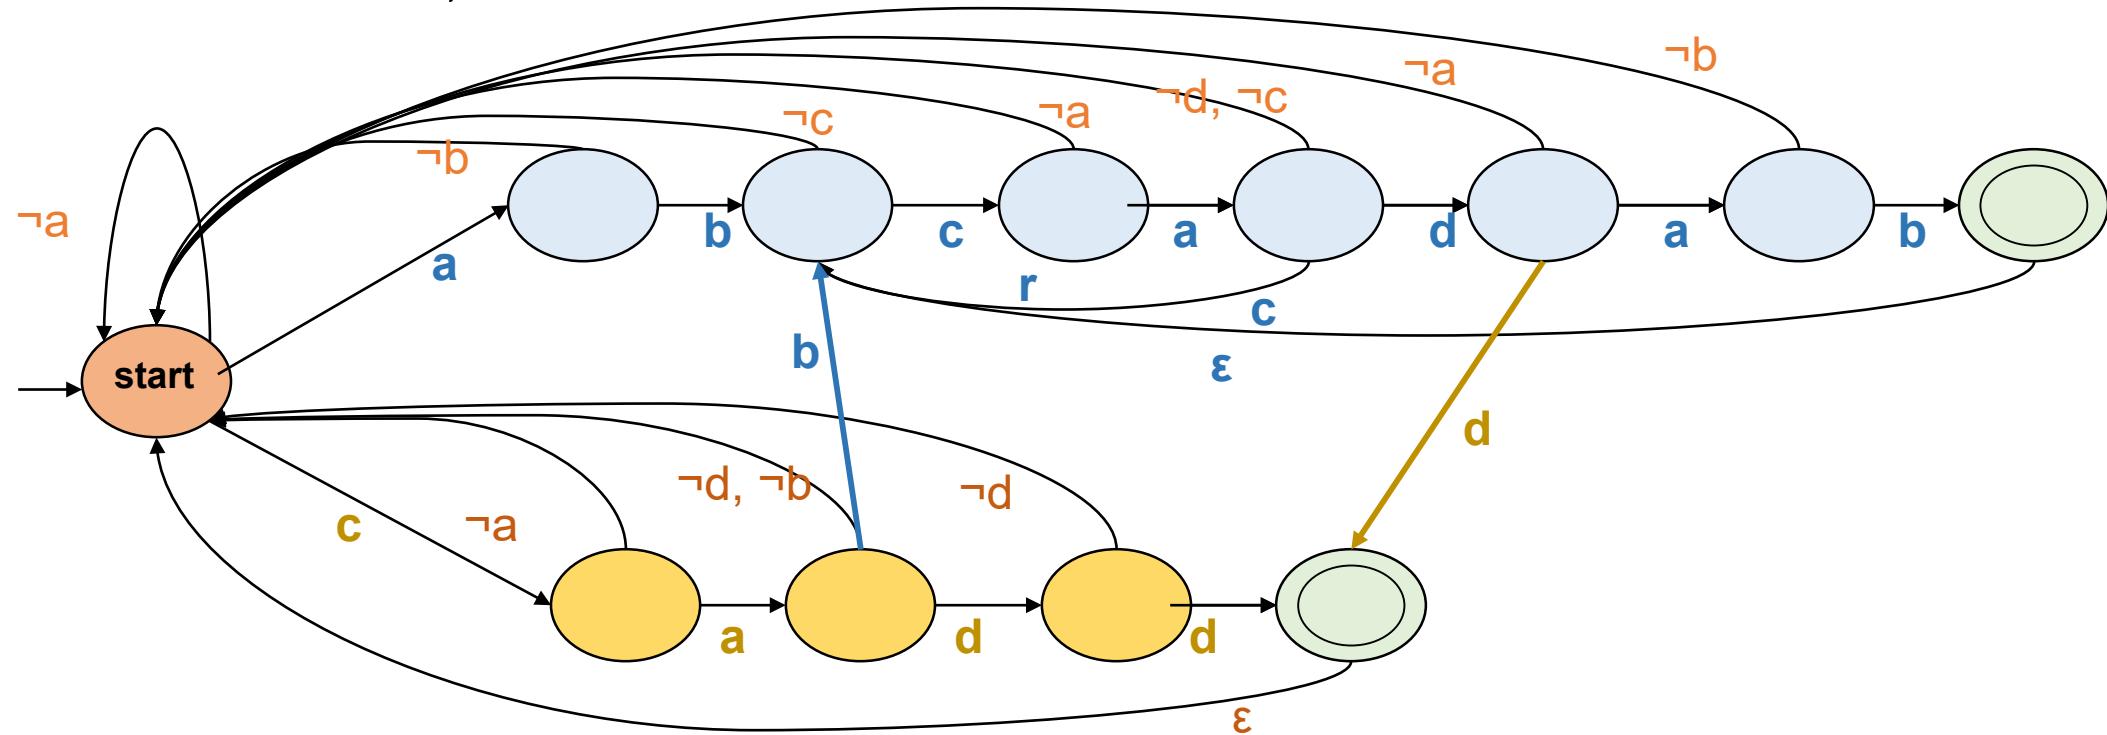
Knuth-Morris-Pratt (KMP)

- Given a sequence of length N , and pattern of length M
- Knuth-Morris-Pratt: $O(N)$
- Example
 - Pattern: $abcadab$



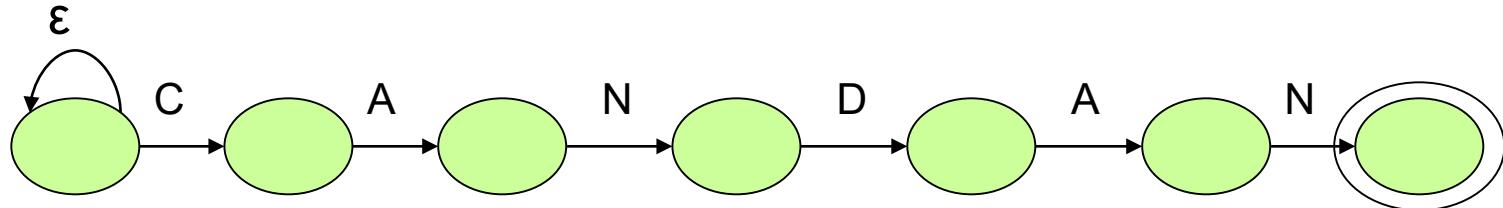
Aho-Corasick Trie

- What is we are given **multiple patterns** to search simultaneously
- Given a **sequence of length N**, and **patterns of length M₁ and M₂**
- Aho-Corasick Trie: O(N)
- Example
 - Patterns: **abcdab; cadd**

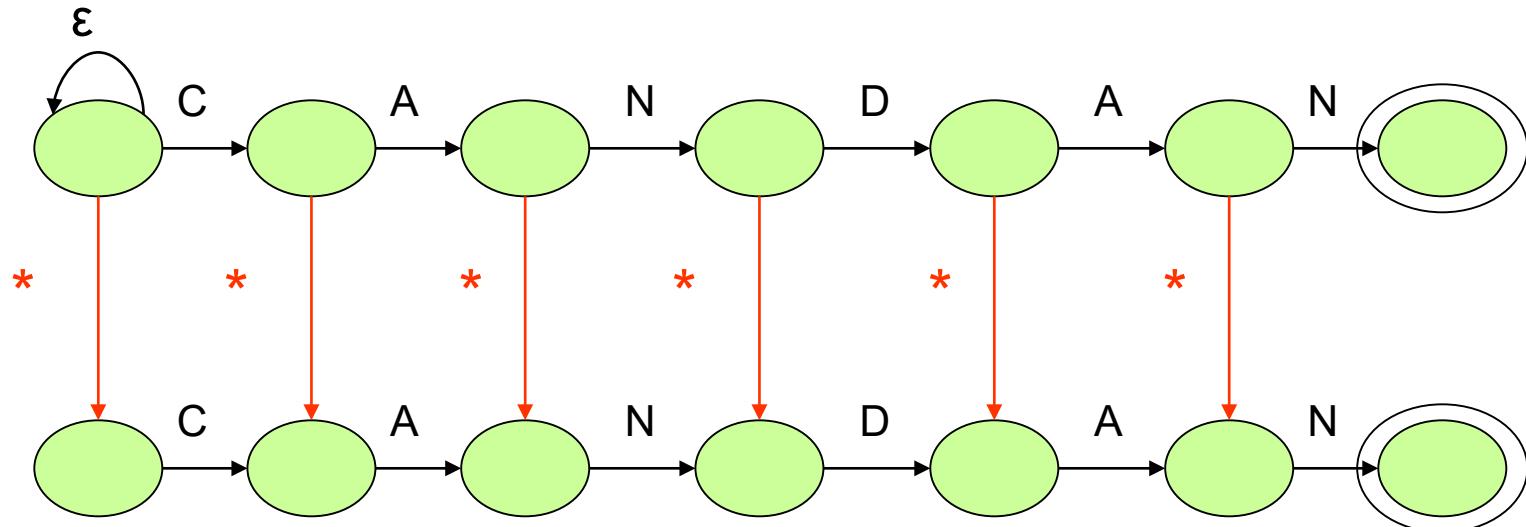


What about approximate matches?

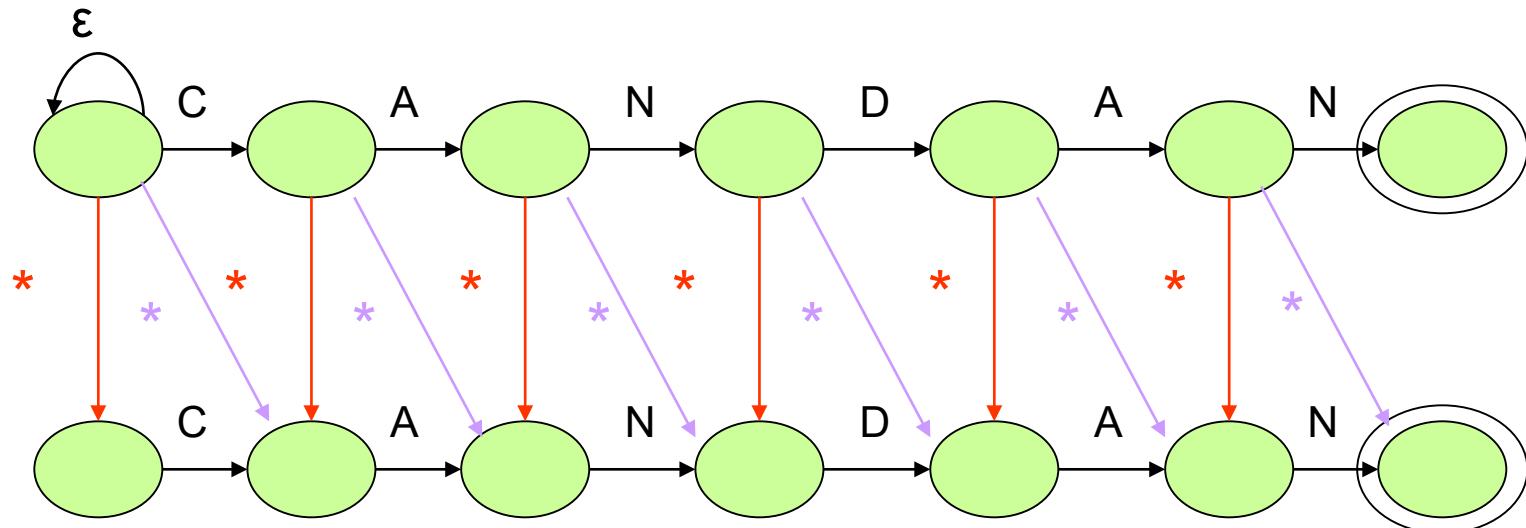
- Example
 - Pattern: CANDAN



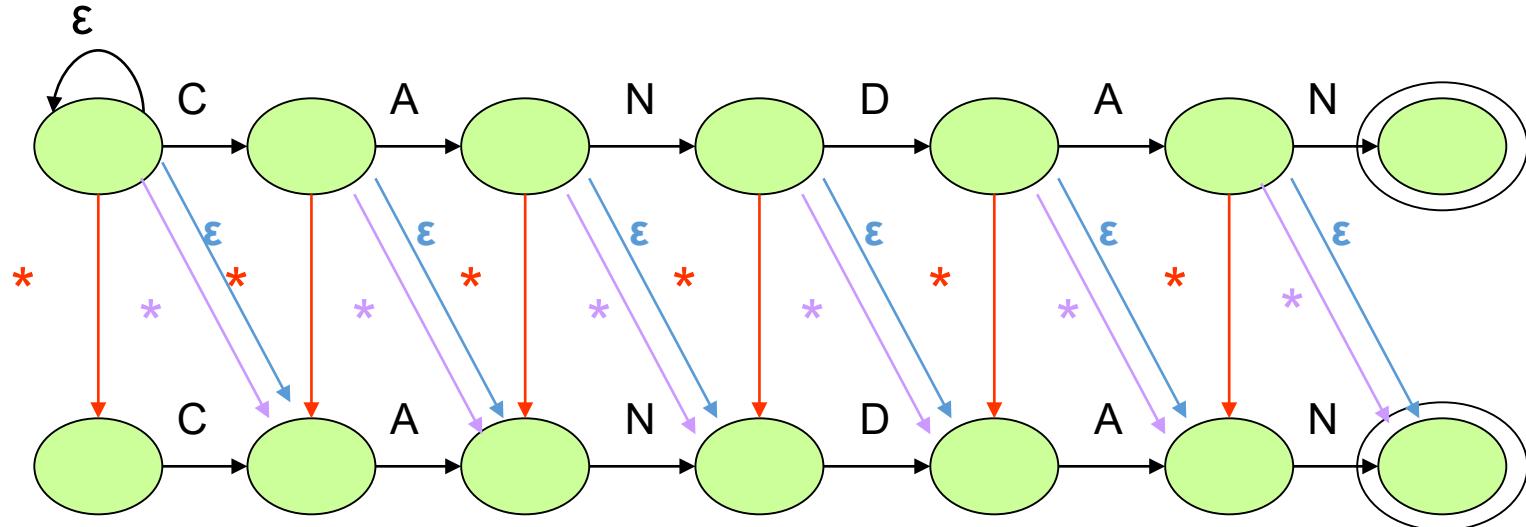
NFA...upto 1 insertion



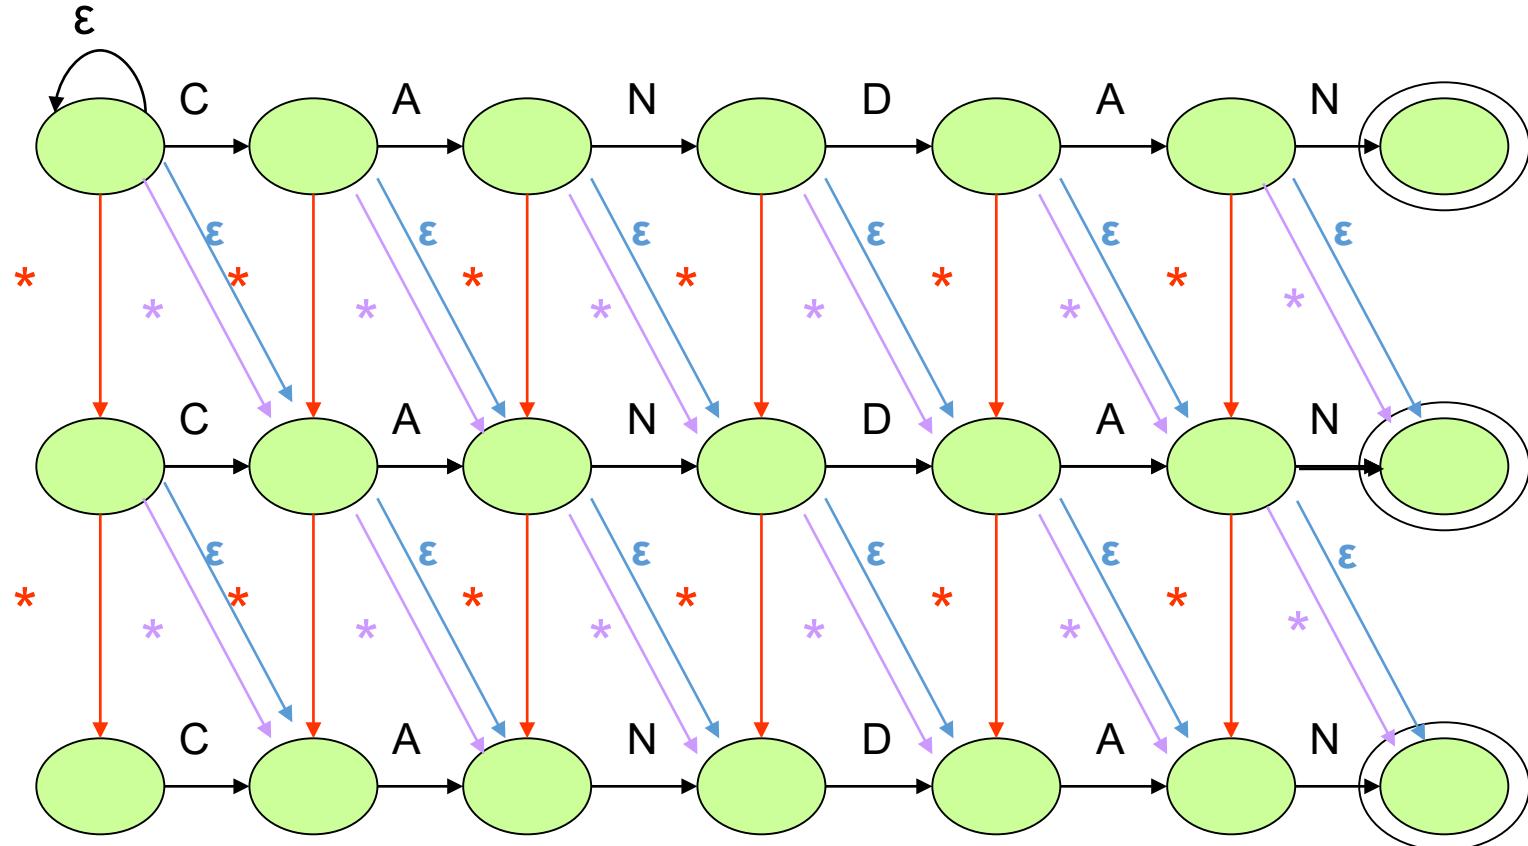
NFA...upto 1 insertion\replacement



NFA...upto 1 ins.\rep.\deletion



NFA...upto 2 ins.\rep.\deletion



Summary



- Prefix based sequence exploration:
 - Trie data structure helps prune the candidate set
- Subsequence search and exploration
 - Suffix trees and suffix arrays helps focus on the part of a long sequence
- Pattern matching
 - Non-deterministic finite automata can be used to support exact and approximate pattern matching



Sequences and Time Series

Edit Distance

K. Selçuk Candan, Professor of Computer Science and Engineering



Strings, sequences, time series



| A *string* or *sequence*, $S = (c_1, c_2, \dots, c_N)$, is a finite sequence of symbols.

abcbbbaabbaabcbbaaabbc

- Prefix search:
 - Find all strings that start with “tab”:
 - “table”; “tabular”; “tablet”; ...
- Subsequence search:
 - Find all strings that contain the subsequence “ark”:
 - “marketing”; “spark”; “quark”
 - Find all occurrences of “acd”:
 - “aabacdcdabdcababdacddcab.”
- Sequence similarity:
 - “table” vs. “cable”?
 - “table” vs. “tale”?
 - “table” vs. “tackle”?

Approximate string match



- Sequence distance/similarity:
 - “table” vs. “cable”?
 - “table” vs. “bale”?
- Edit distance:
 - “**t**able” vs. “**c**able”: 1 (replace “**t**” with “**c**”)
 - “**t**able” vs. “**b**ale”: 3 (delete “**t**”; replace “**a**” and “**b**”; replace “**b**” and “**a**”)
- Common edit operations
 - Replacement:
 - $a \rightarrow b$
 - Deletion:
 - $a \rightarrow \lambda$
 - Insertion:
 - $\lambda \rightarrow a$

Edit cost

- Let E be a sequence of edit operations to convert one string to another
- Let us associate a cost, C , to each edit operation

- Costs of edit operations can be different from each other
 - Type of the operation (replace, delete, insert)
 - Symbols involved in the operation
 - Position of the edit operation

- Given a sequence of edit operations, E

$$C(E) = \sum_{e_i \in E} C(e_i)$$

- Edit Distance:

$$D(String_1, String_2) = \min_{E \text{ takes } String_1 \text{ to } String_2} \{C(E)\}$$

Edit distance

- Let us be given two strings, P and Q, of lengths N and M
- Let us assume that all edit operations have cost = 1

$D[i,j]$ = # of edits from length-i prefix of P to length-j prefix of Q

Edit distance

- Let us be given two strings, P and Q, of lengths N and M
- Let us assume that all edit operations have cost = 1

$D[i,j]$ = # of edits from length-i prefix of P to length-j prefix of Q

● $D[0,j] = j$

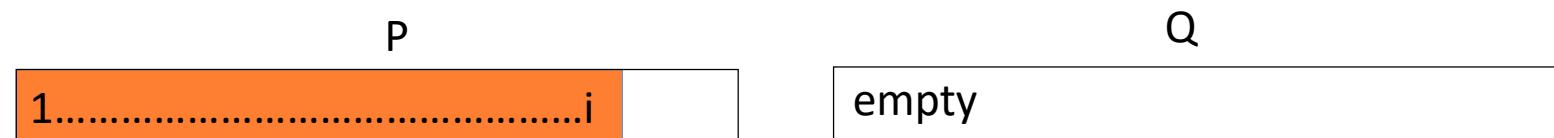


Edit distance

- Let us be given two strings, P and Q, of lengths N and M
- Let us assume that all edit operations have cost = 1

$D[i,j]$ = # of edits from length-i prefix of P to length-j prefix of Q

- $D[0,j] = j$
- $D[i,0] = i$



Edit distance

- Let us be given two strings, P and Q, of lengths N and M
- Let us assume that all edit operations have cost = 1

$D[i,j]$ = # of edits from length-i prefix of P to length-j prefix of Q

- $D[0,j] = j$
- $D[i,0] = i$

• if($P_i = Q_j$)

- $D[i,j] = D[i-1,j-1]$

for all (i<=N) and (j <= M)



Edit distance

- Let us be given two strings, P and Q, of lengths N and M
 - Let us assume that all edit operations have cost = 1

$D[i,j] = \# \text{ of edits from length-}i \text{ prefix of } P \text{ to length-}j \text{ prefix of } Q$

- $D[0,j] = j$
 - $D[i,0] = i$

else $D[i,j] = 1 + \min\{$

insert Qj

$D[i-1,j]$,

8

G

delete Pi

$D[i, j-1]$,

1

.....j-1

replace Pi with Qj

$D[i-1, j-1]$ }

1 i-1

1

1 j-1

Edit distance

- Let us be given two strings, P and Q, of lengths N and M
 - Let us assume that all edit operations have cost = 1

$D[i,j]$ = Cost of edits from length- i prefix of P to length- j prefix of Q

- $D[-1,j] = \text{infinity}$; $D[i,-1] = \text{infinity}$
 - $D[0,0] = 0$

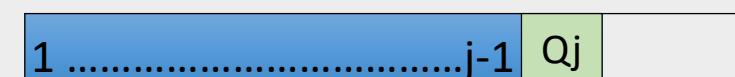
else $D[i,j] = \min\{$

P

Q

insert Gi

$$C_{ins}(Q_j) + D[i-1,j] ,$$



delete Pi

$$C_{del}(P_i) + D[i,j-1],$$



replace
Pi with Gi

$$C_{rep}(P_i, Q_j) + D[i-1, j-1]$$



}

Edit distance

- Let us be given two strings, P and Q, of lengths N and M
 - Let us assume that all edit operations have cost = 1

O(N*M)

$D[i,j]$ = Cost of edits from length- i prefix of P to length- j prefix of Q

- $D[-1,j] = \text{infinity}$; $D[i,-1] = \text{infinity}$
 - $D[0,0] = 0$

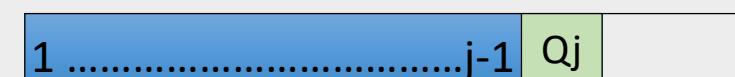
else $D[i,j] = \min\{$

P

Q

insert Gi

$$C_{ins}(Q_j) + D[i-1,j] ,$$



delete Pi

$$C_{del}(P_i) + D[i,j-1],$$



replace
Pi with Gi

$$C_{rep}(P_i, Q_j) + D[i-1, j-1]$$



}

Summary



- Edit distance can be used to assess how similar or different two strings are
- Problem: Edit distance can be costly for matching long strings.

Sequences and Time Series

W-Grams and Other Approaches to Filtering

K. Selçuk Candan, Professor of Computer Science and Engineering



Strings, sequences, time series



| A *string* or *sequence*, $S = (c_1, c_2, \dots, c_N)$, is a finite sequence of symbols.

abcbbbaabbaabcbbaaabbc

- Prefix search:
 - Find all strings that start with “tab”:
 - “table”; “tabular”; “tablet”; ...
- Subsequence search:
 - Find all strings that contain the subsequence “ark”:
 - “marketing”; “spark”; “quark”
 - Find all occurrences of “acd”:
 - “aabacdcdabdcababdacddcab.”
- Sequence similarity:
 - “table” vs. “cable”?
 - “table” vs. “tale”?
 - “table” vs. “tackle”?

Reminder: Edit Distance

- Let P (of size N) and Q (of size M) be two sequences.
- Given a sequence of edit operations, E

$$C(E) = \sum_{e_i \in E} C(e_i)$$

- Edit distance

$$D(String_1, String_2) = \min_{E \text{ takes } String_1 \text{ to } String_2} \{C(E)\}$$

- Cost: $O(N*M)$

Cross-parsing Distance

- Let P (of size N) and Q (of size M) be two sequences
- $c(P|Q)$: cost of parsing P with respect to Q
 1. Find the longest (possibly empty) prefix of P that appears as a string somewhere in Q
 2. Restart from the very next position in P and continue until P is completely parsed.

Cross-parsing Distance

- Let P (of size N) and Q (of size M) be two sequences
- $c(P|Q)$: cost of parsing P with respect to Q
 1. Find the longest (possibly empty) prefix of P that appears as a string somewhere in Q
 2. Restart from the very next position in P and continue until P is completely parsed.

$$D_{crossparse}(P, Q) = \frac{(c(P|Q)-1) + (c(Q|P)-1)}{2}$$

- Linear time if the strings are indexed using a suffix tree

Compression Distance



- $C(P)$ is the compressed size of P ,
- $C(Q)$ is the compressed size of Q , and
- $C(P:Q)$ is the compressed size of the sequence obtained by concatenating P and Q .

$$D_{NCD}(P,Q) = \frac{C(P:Q) - \min\{C(P), C(Q)\}}{\max\{C(P), C(Q)\}}$$

Filtering



- Given a string P of (length N) and a pattern Q (of length M), determine whether the string P may contain an approximate match to Q

Filtering



- Given a **string P of (length N)** and a **pattern Q (of length M)**, determine whether the **string P may contain an approximate match to Q with at most k errors**
- Approach 1: Given a maximum error rate, **k**,
 - cut the **pattern Q** into **k + 1** pieces
 - verify that **at least one piece of Q exists in P exactly**This is because **k errors** cannot affect more than **k pieces**.

Filtering



- Given a **string P of (length N)** and a **pattern Q (of length M)**, determine whether the **string P may contain an approximate match to Q with at most k errors**
- Approach 2: Given a maximum error rate, **k**,
 - slide a **window of length M** over **the string P** and count the number of symbols that are included in **the pattern Q**
 - only windows that have at least $M - k$ matching symbols need to be considered

Fingerprinting with w-grams



- **w-grams:** Given a sequence P , its w -grams are obtained by sliding a window of length w over P .

Common w-gram counting

- Given a **string P** of (length **N**) and a **string Q** (of length **M**), determine whether the two strings may match each other with at most **k errors**
 - w-grams:** Given a **sequence P**, its **w-grams** are obtained by sliding a window of length **w** over **P**.
- Approach (common w-gram counting):
 - Identify $(M - w + 1)$ **w-grams** of the **query string Q**
 - Each mismatch between **Q** and **P** can affect **w** many **w-grams**
 - given an upper bound of **k errors**, at least $(M - w + 1 - kw)$ **w-grams** must match
 - Search for these matches using a suffix tree **(in linear time)**

String kernels

- Given a **string P** of (length **N**) and a **string Q** (of length **M**), determine whether the strings may approximately match each other
 - **w-grams:** Given a **sequence P**, its **w-grams** are obtained by sliding a window of length **w** over **P**.
- Approach (string kernels):
 - Identify all **w-grams** of the **query sequence Q**; create a counting vector, **q**
 - Identify all **w-grams** of the **data sequence P**; create a counting vector, **p**
 - Measure the (dot product) similarity of the two counting vectors
 - If the dot product similarity is low, then **P** is not likely to match **Q**

Min-sampling similarity



- Given a **string P** of (length **N**) and a **string Q** (of length **M**), determine whether the strings may approximately match each other
 - w-grams:** Given a **sequence P**, its **w-grams** are obtained by sliding a window of length **w** over **P**.
- Approach (min-sampling similarity):
 - Consider **r** random hash orders of the **w-grams** of **P** and **Q**
 - For each order o_i
 - Find the **smallest B w-grams** of string **P** based on the chosen order
 - Find the **smallest B w-grams** of string **Q** based on the chosen order
 - If they agree, $\text{match}_i(P, Q) = 1$
 - After computing the match for all **r** orders

$$sim(P, Q) = \frac{\sum_{i=1}^r \text{match}_i(P, Q)}{r}$$

Summary



- Edit distance can be costly for matching long strings
 - Cross-Parsing and Compression-Distance can be used to approximate the edits distance comparison
- W-grams (more commonly known as n-grams) can be used to help filter unpromising candidates before more costly distance computations

Sequences and Time Series

Time Series Modeling

K. Selçuk Candan, Professor of Computer Science and Engineering

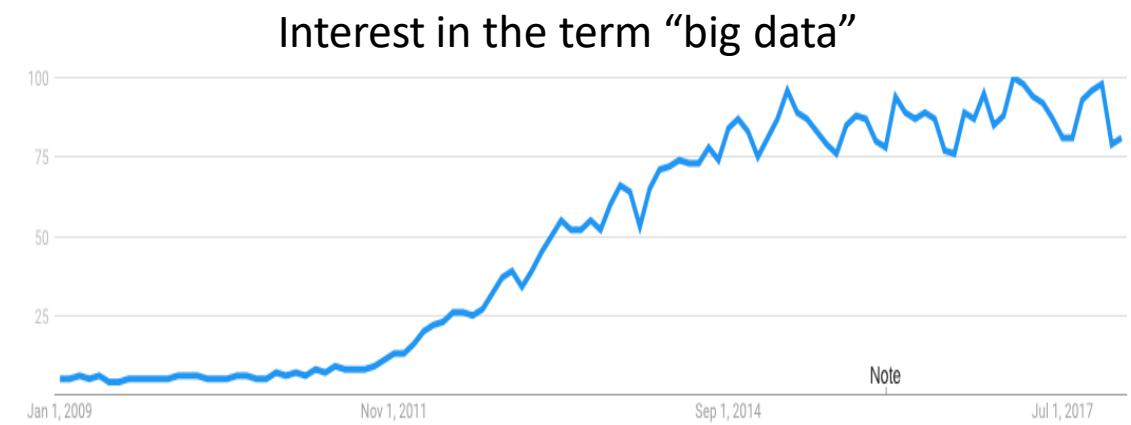


Strings, sequences, time series

| A *string* or *sequence*, $S = (c_1, c_2, \dots, c_N)$, is a finite sequence of symbols.

abcbbbaabbaabcbbaaaabbcc

| A *time series*, $T = (d_1, d_2, \dots, d_N)$, is a finite sequence of data values.



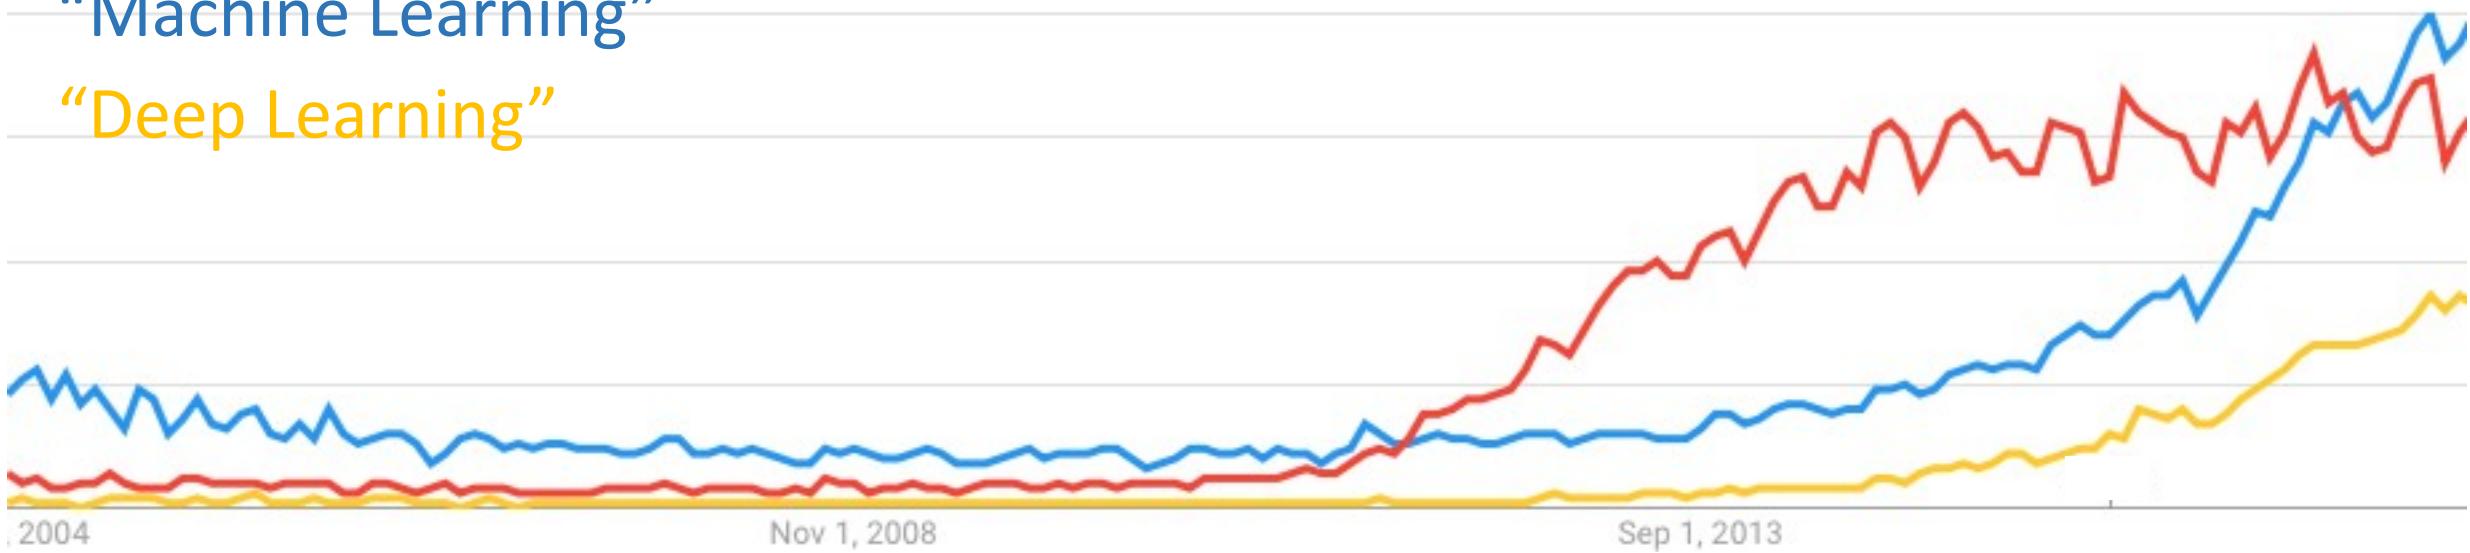
<https://trends.google.com/trends/explore?date=2008-12-19%202018-01-19&q=big%20data>

Comparing time series

“Big Data”

“Machine Learning”

“Deep Learning”

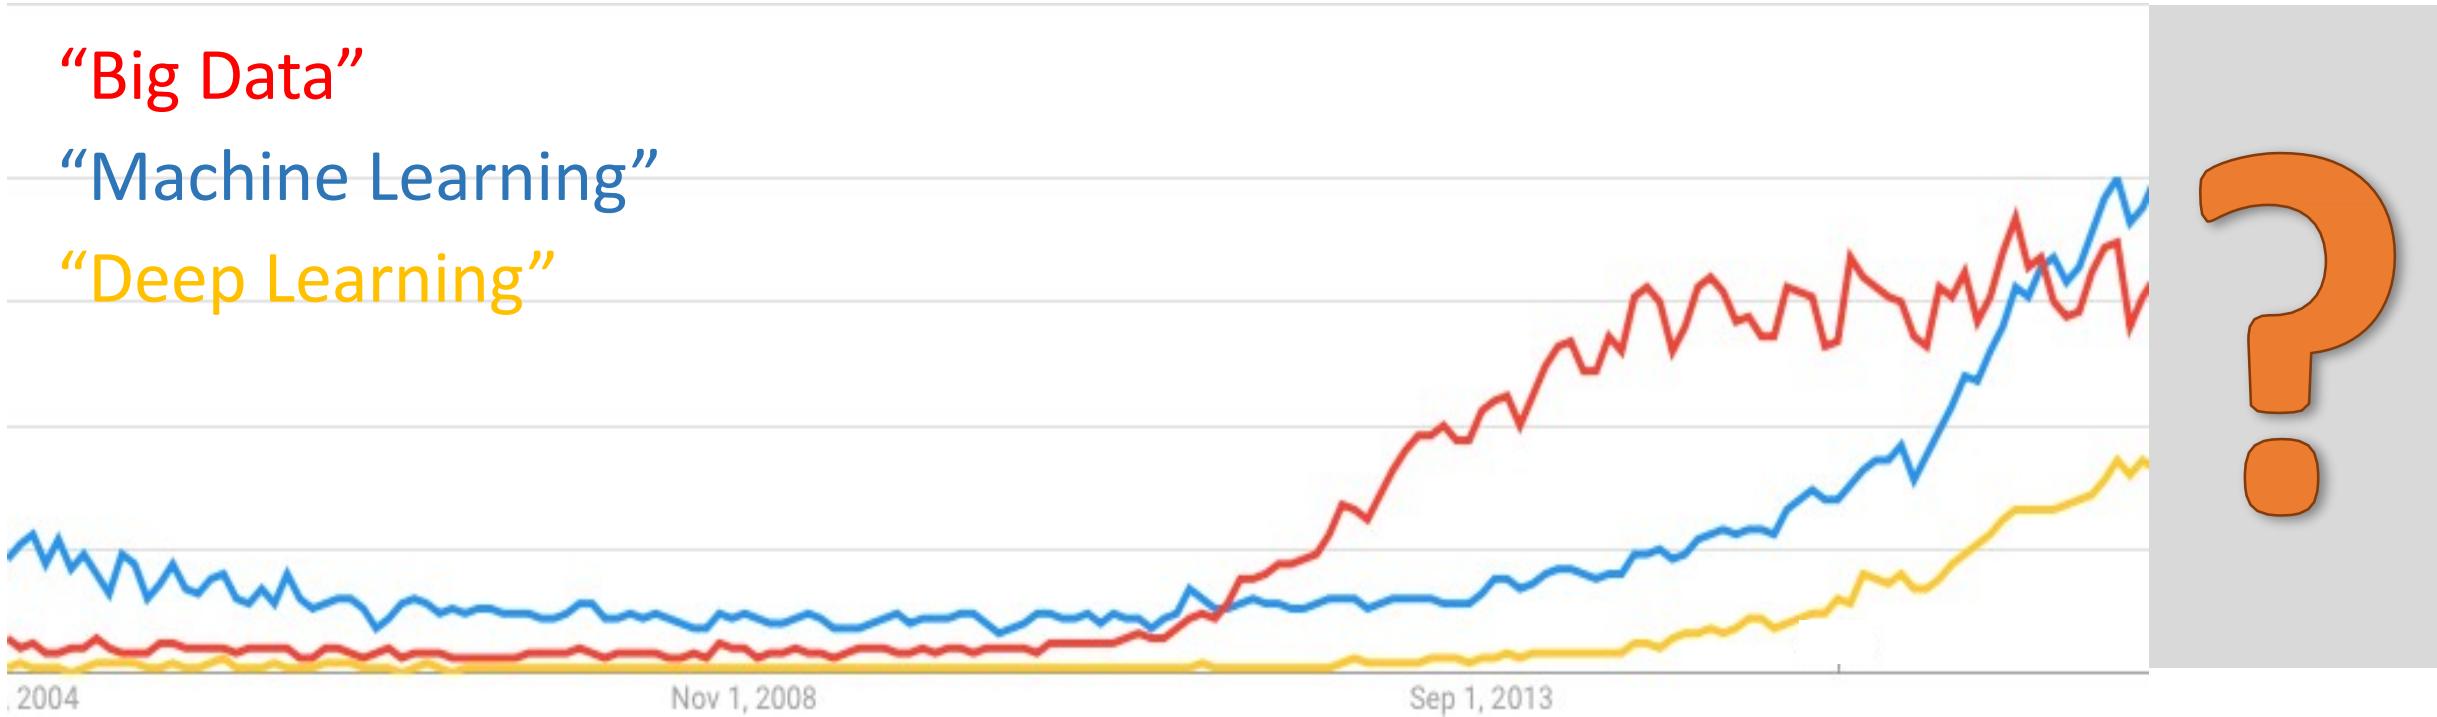


Forecasting time series

“Big Data”

“Machine Learning”

“Deep Learning”

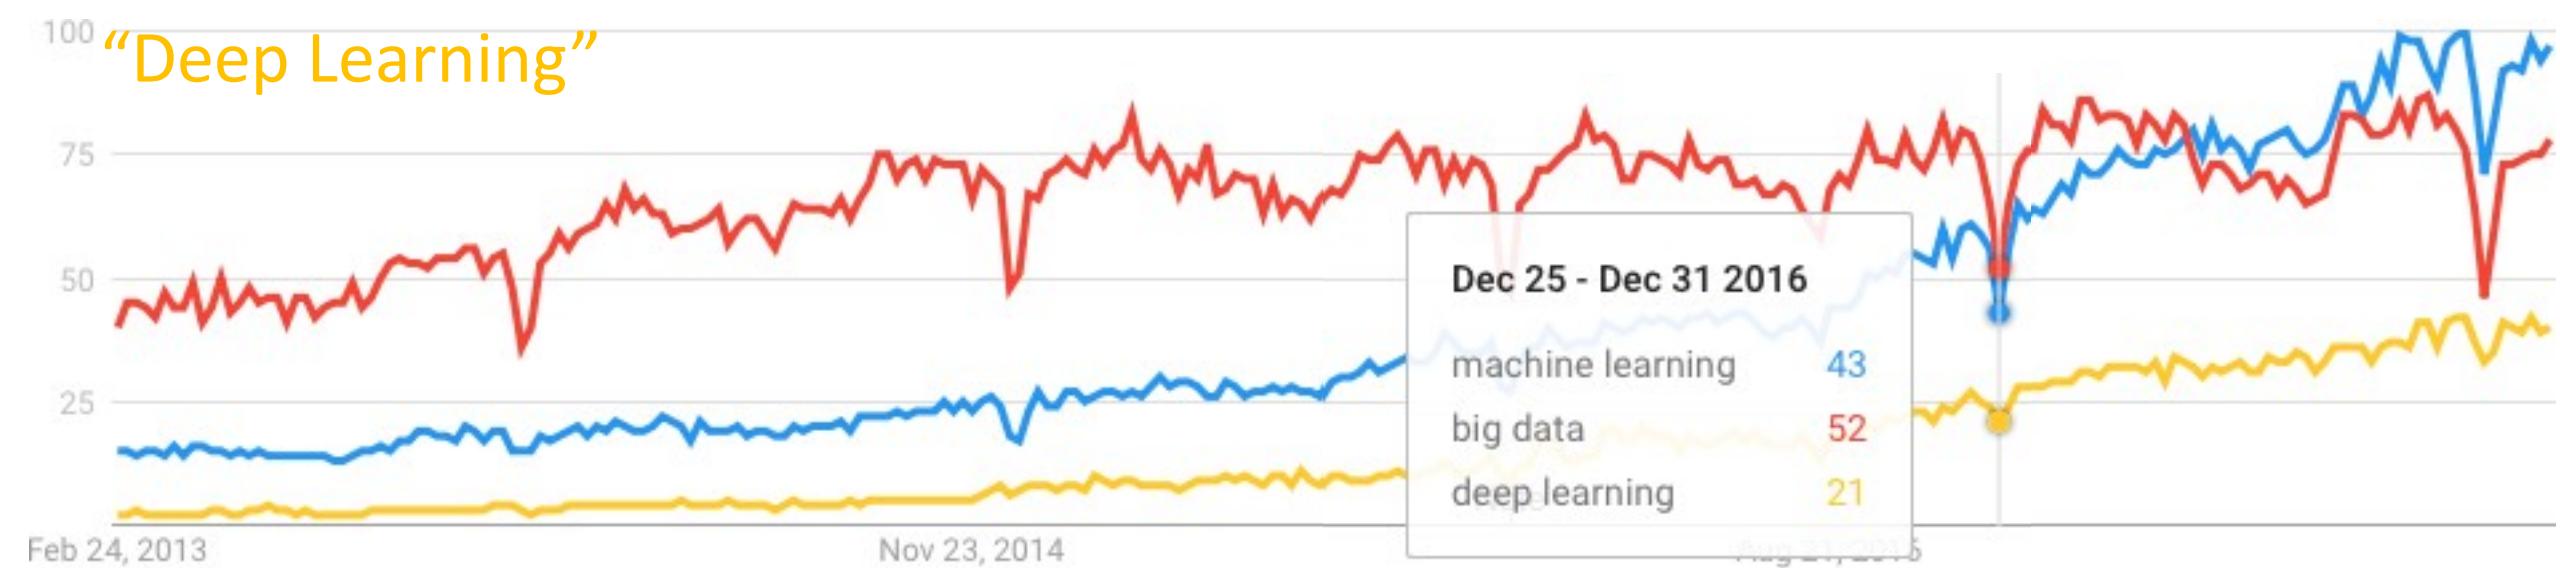


Motifs

“Big Data”

“Machine Learning”

“Deep Learning”



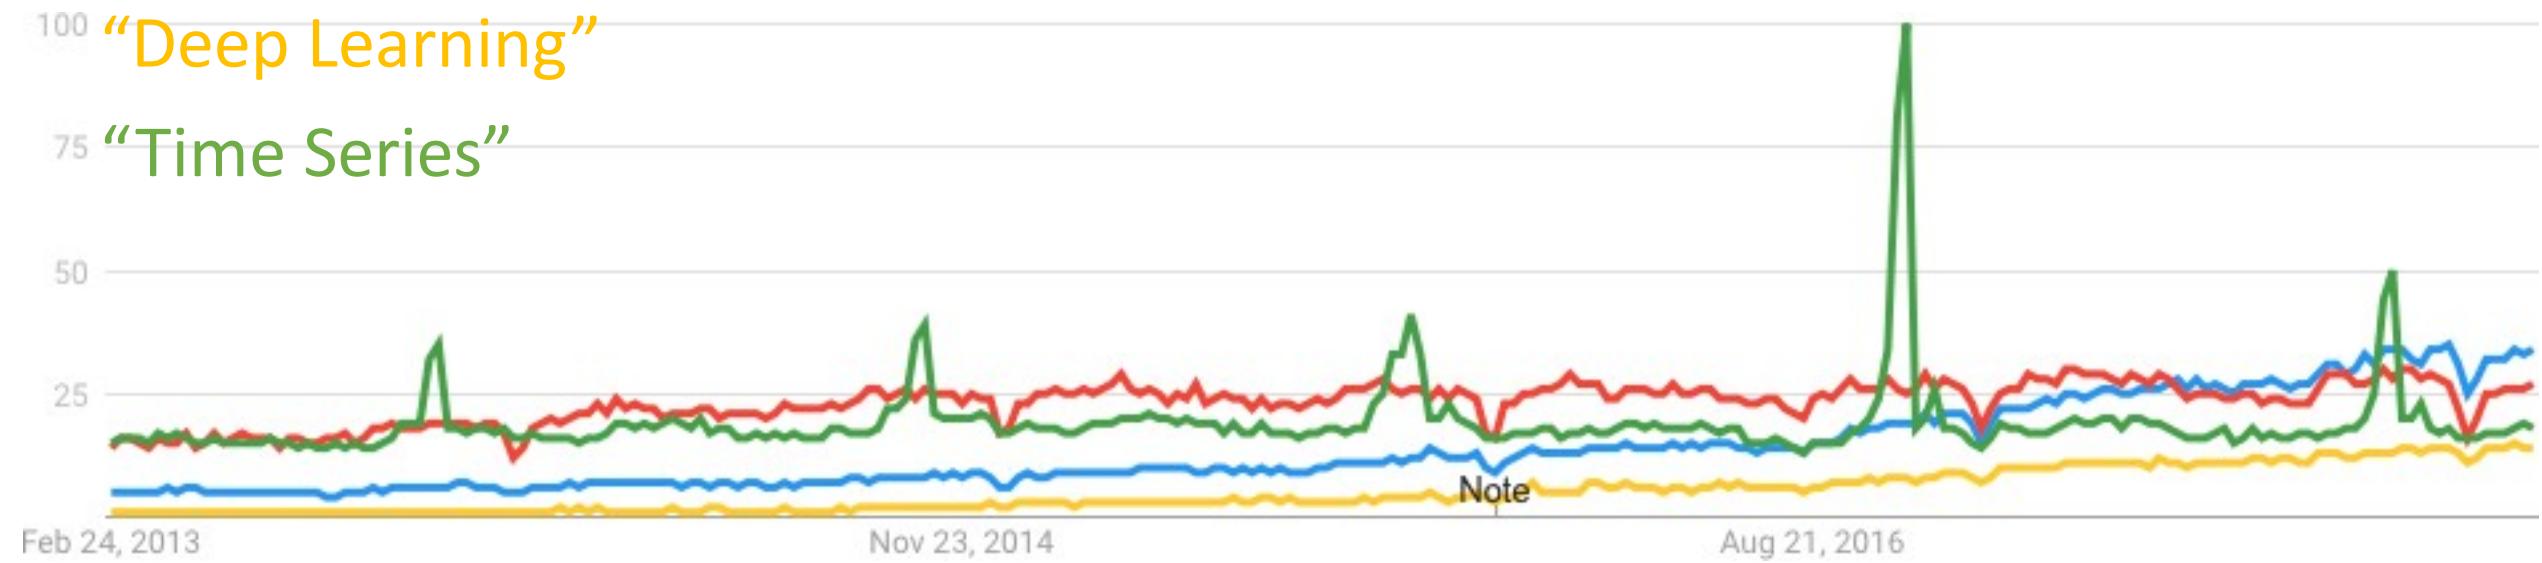
Motifs

“Big Data”

“Machine Learning”

“Deep Learning”

“Time Series”



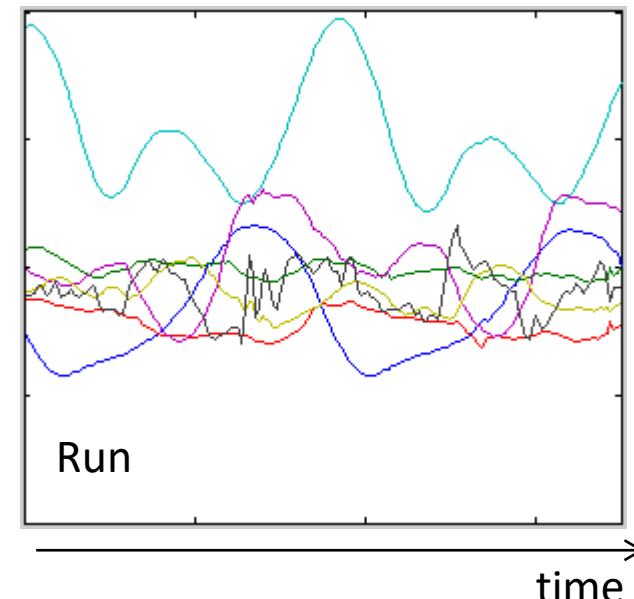
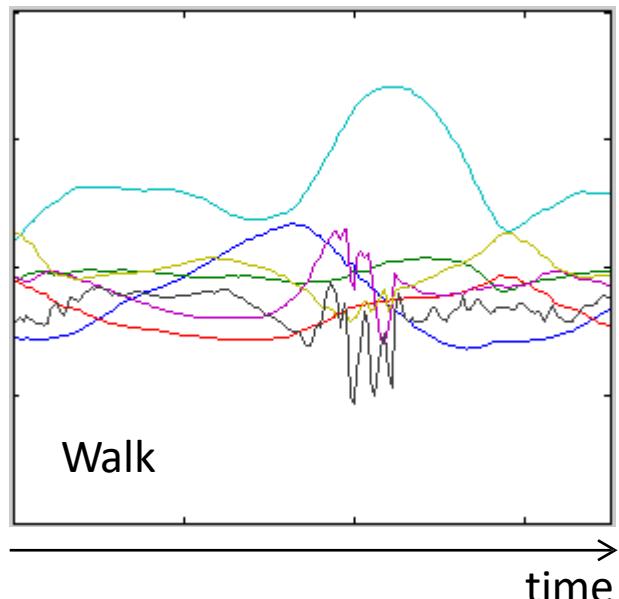
Note

Classification

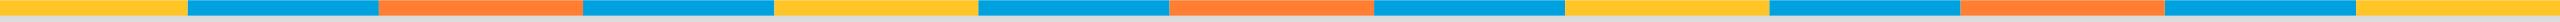


Markers on
human body [1]

Multi-variate time series tracking 7 markers on subject's left leg



Time Series Models



- **Question:** Can we discover a closed form formula (or a model) that describes the time series?
- **Simpler Question:** Can we characterize high-level properties of a given time series?

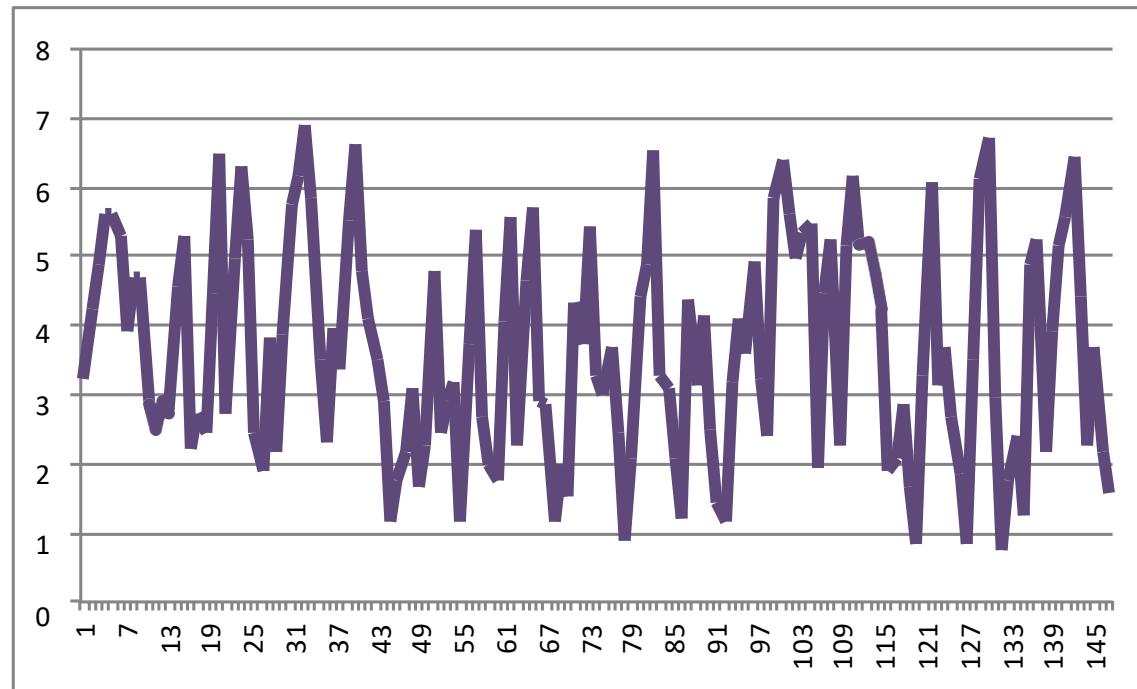
Time Series



Types of Time Series

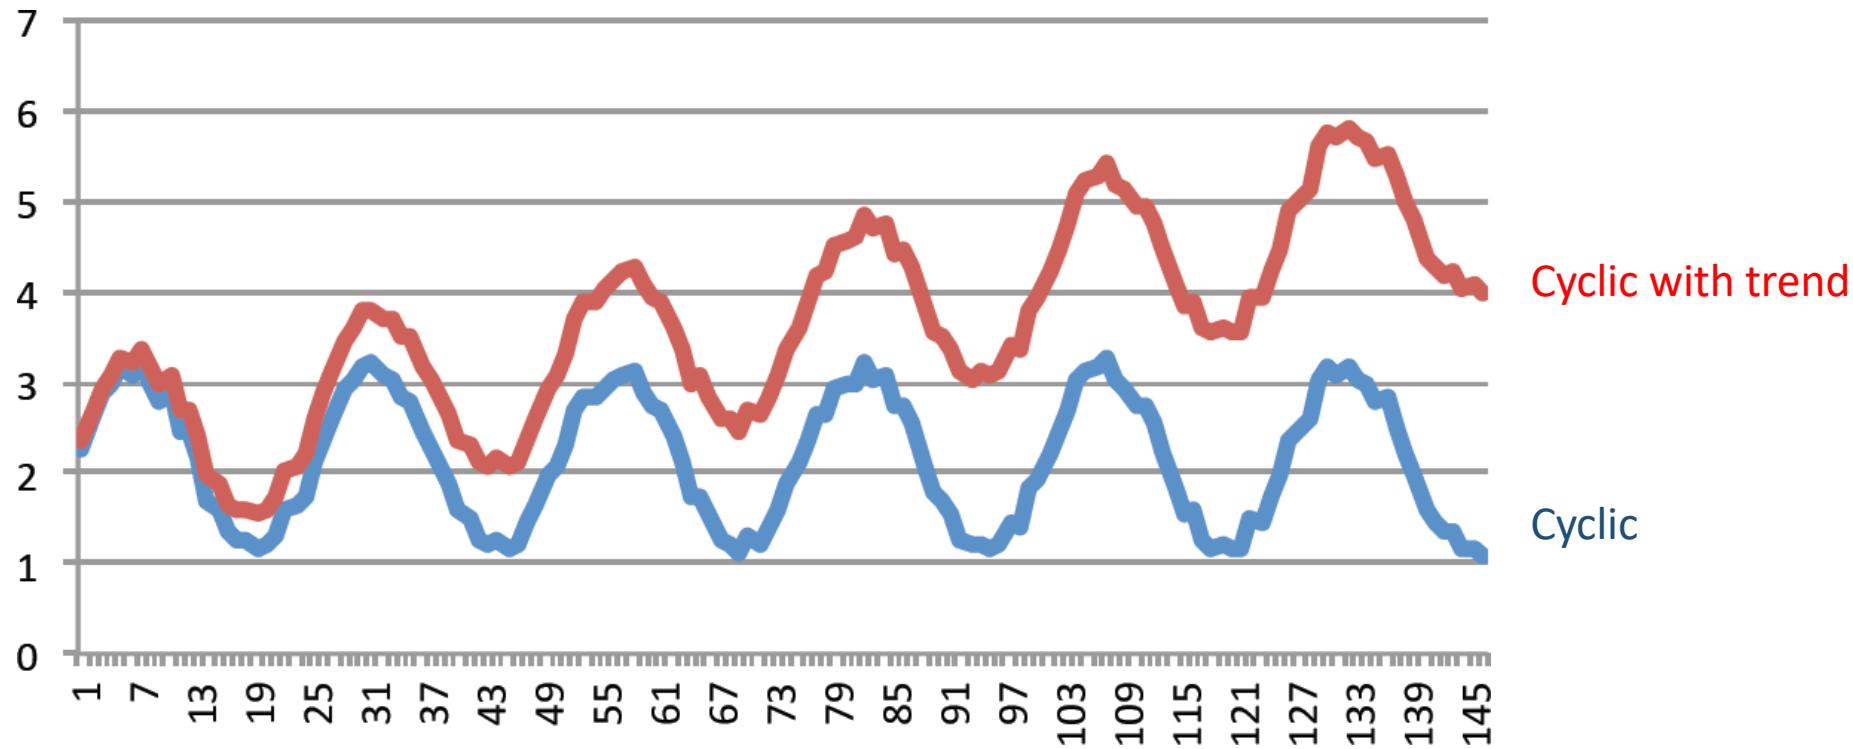
Stationary Series

- Stationary series
 - statistical properties, such as mean and variance, are constant over time.



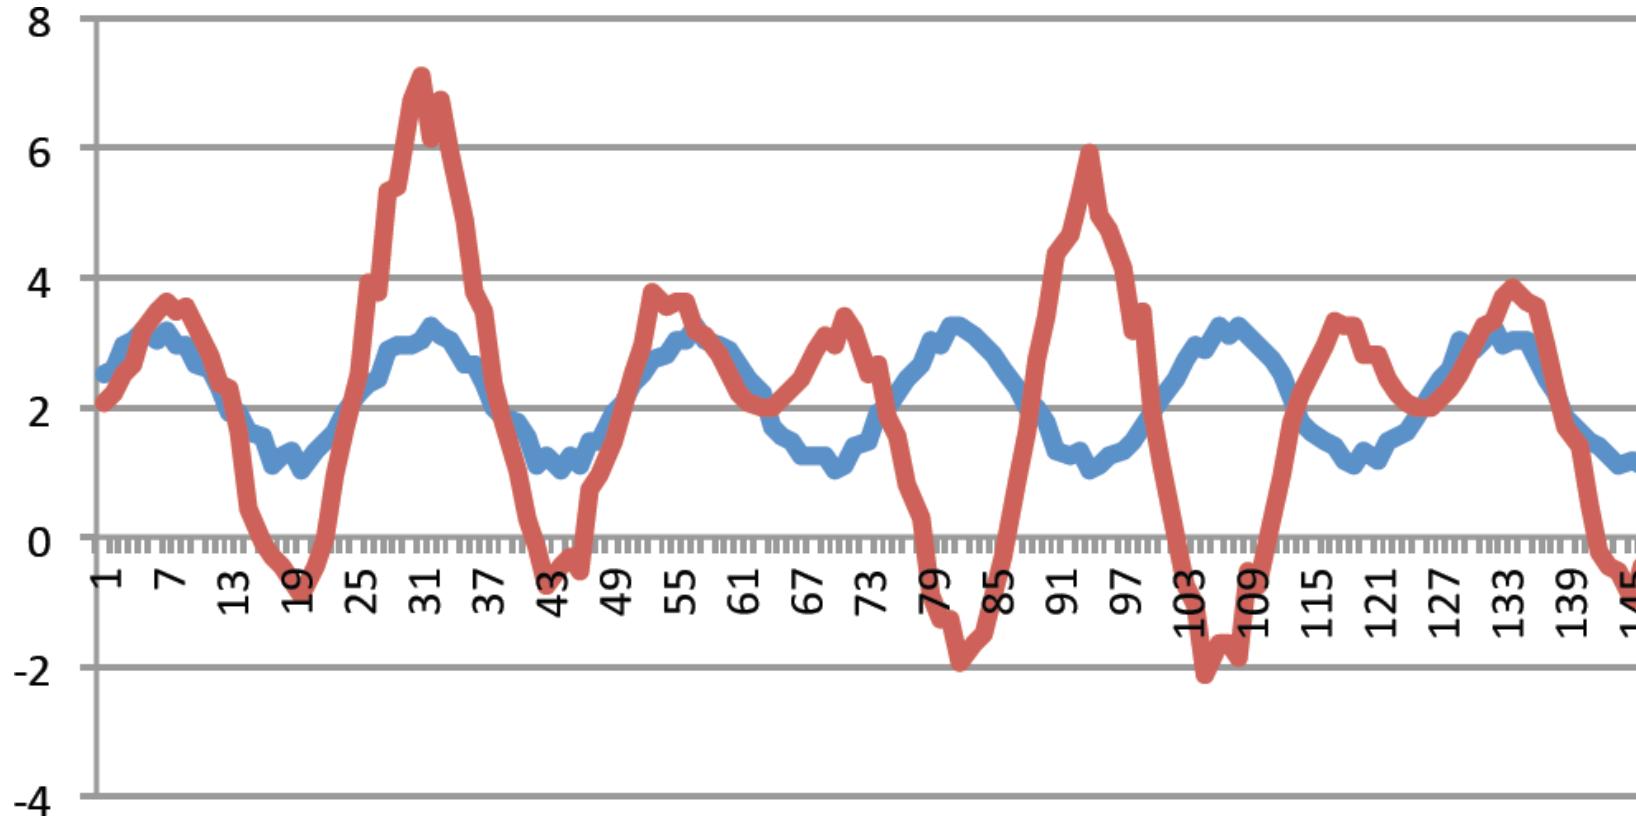
Non-Stationary Series

- Non-stationary series
 - statistical properties **change over time**



Non-Stationary Series

- Non-stationary series
 - statistical properties **change over time**

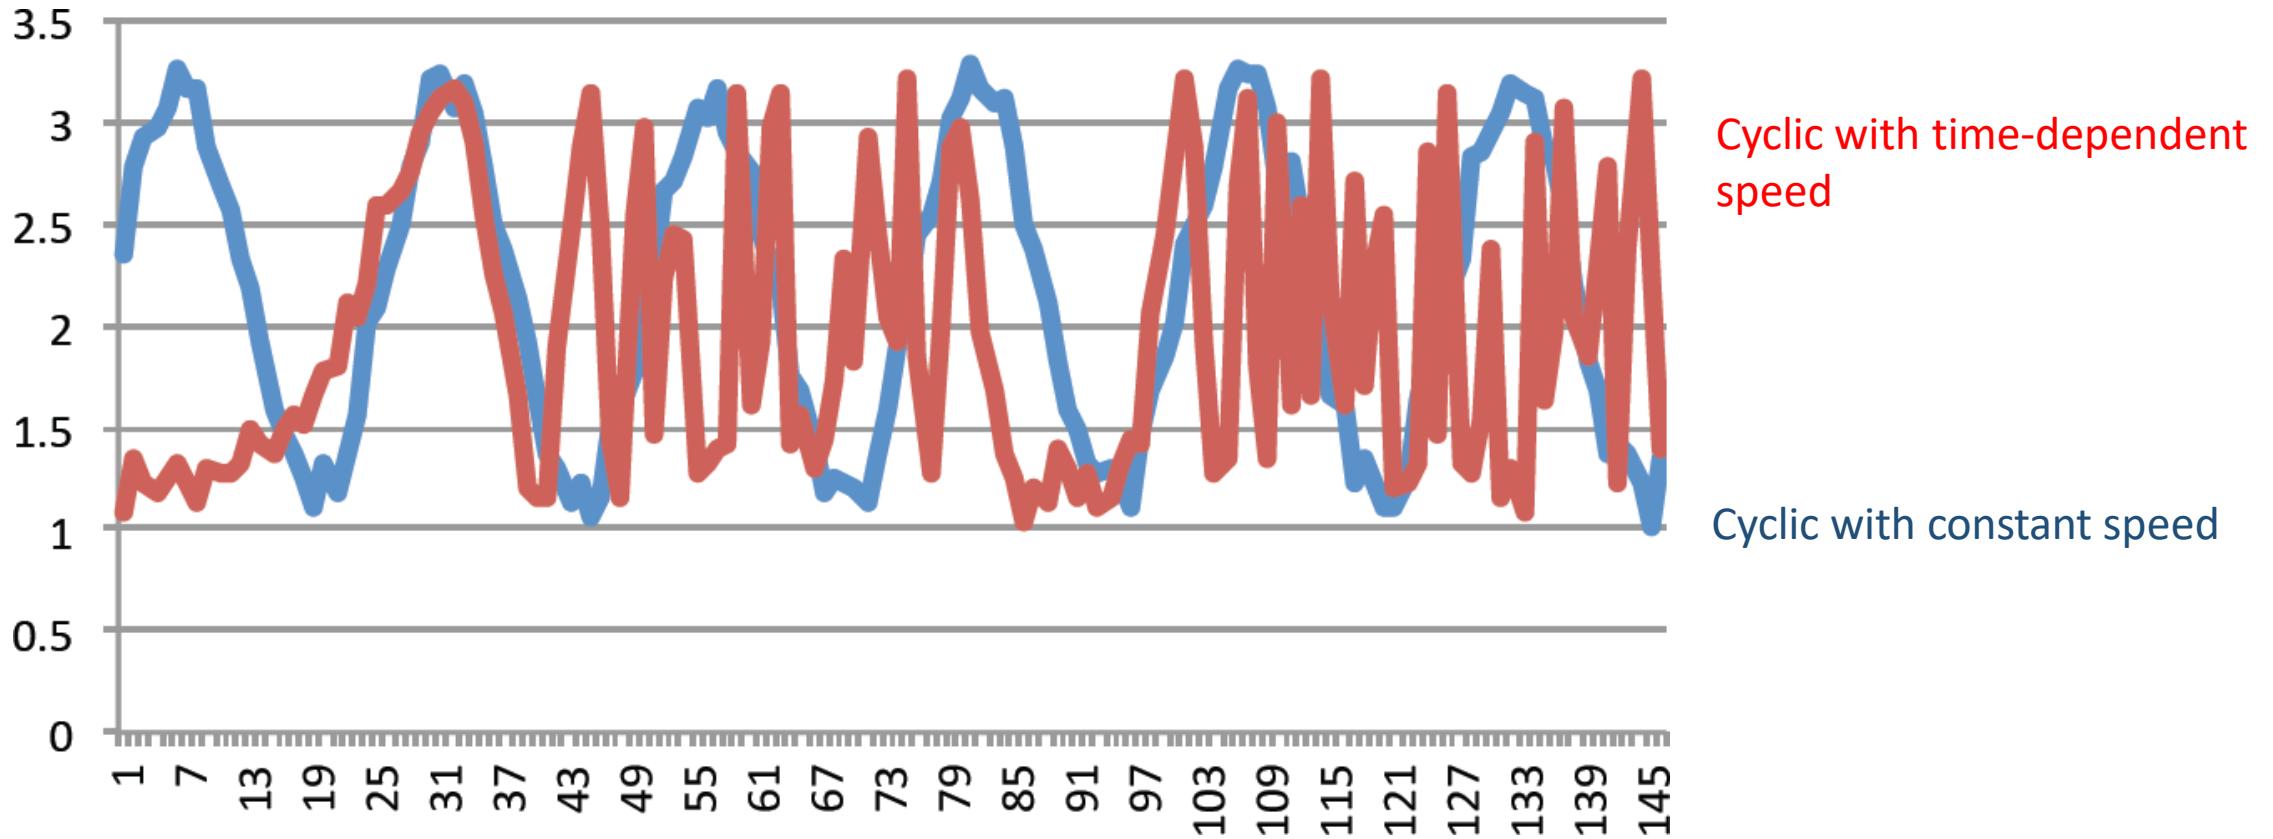


Cyclic with constant variance

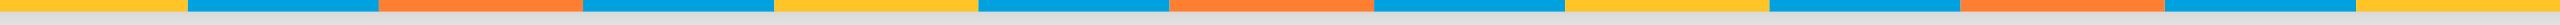
Cyclic with time-dependent variance

Non-Stationary Series

- Non-stationary series
 - statistical properties **change over time**



Why do these matter?



- Stationary series
 - statistical properties, such as mean and variance, are constant over time.
- Non-stationary series
 - statistical properties change over time
- Why important?
 - most statistical forecasting methods assume that the series **can be rendered (approximately) stationary** through mathematical **transformations**

Time Series



Models of Time Series

Random Time Series Model



- Current observations only reflect current (random) error:

$$X_t = E_t \quad E_t \sim N(0, \sigma^2)$$

- Intuitively, the “random error” represents a stochastic process that does not depend on the past

Autoregressive Time Series Model

- Current observations only reflect current (random) error and has no dependence on the past observations:

$$X_t = E_t \quad E_t \sim N(0, \sigma^2)$$

- AR(1): the current observation depends only the previous time instance and the current error

$$X_t = \alpha X_{t-1} + E_t + \lambda$$

Autoregressive Time Series Model

- Current observations only reflect current (random) error and has no dependence on the past observations:

$$X_t = E_t \quad E_t \sim N(0, \sigma^2)$$

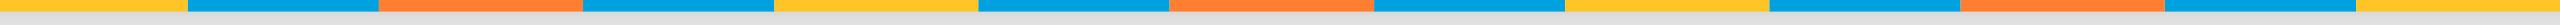
- AR(1): the current observation depends only the previous time instance and the current error

$$X_t = \alpha X_{t-1} + E_t + \lambda$$

- AR(2): the current observation depends only the previous 2 time instances and the current error

$$X_t = \alpha_1 X_{t-1} + \alpha_2 X_{t-2} + E_t + \lambda$$

Moving Average Time Series Model



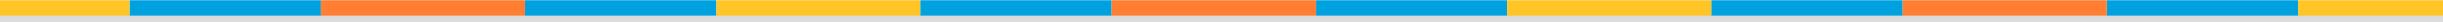
- MA(1): The current observation depends only to the error in the previous time instance and the current error

$$X_t = \beta E_{t-1} + E_t + \lambda$$

- MA(2): The current observation depends only to the error in the previous 2 time instances and the current error

$$X_t = \beta_1 E_{t-1} + \beta_2 E_{t-2} + E_t + \lambda$$

ARMA model



- ARMA(a,m):
 - AR(a): The model has **a** auto-regressive terms
 - MA(m): The model has **m** moving-average terms

$$X_t = (\alpha_1 X_{t-1} + \dots + \alpha_a X_{t-a}) + (\beta_1 E_{t-1} + \dots + \beta_m E_{t-m}) + E_t + \lambda$$

- **Shortcoming of ARMA models**
 - These models cannot model where the current value is determined by taking into account the speed or degree of acceleration observed in the past.

Models with Differencing



- “differencing” enables models to also consider **speed of change** and **degree of acceleration** in determining the current value:

- Order-1 differencing:

$$X_t^{(1)} = X_t - X_{t-1}$$

speed of change

- Order-2 differencing:

$$X_t^{(2)} = X_t^{(1)} - X_{(1)t-1}^{(1)}$$

degree of acceleration

- ...

- Order-d differencing:

$$X_t^{(d)} = X_t^{(d-1)} - X_{t-1}^{(d-1)}$$

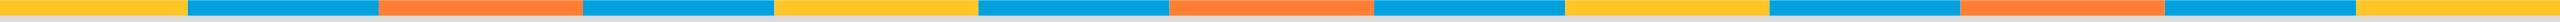
ARIMA model



- ARIMA(a,d,m)
 - AR(a): The model has a auto-regressive terms
 - MA(m): The model has m moving-average terms
 - Integrated with **order-d differencing**

$$X_t = (\alpha_1 X_{t-1} + \dots + \alpha_a X_{t-a}) + (\beta_1 E_{t-1} + \dots + \beta_m E_{t-m}) + E_t \\ + \lambda + (\theta_1 X_{t-1}^{(d)} + \dots + \theta_m X_{t-m}^{(d)})$$

Models with Seasonal Differencing



- Seasonal “differencing” enables models to also consider speed of changes of values across a lag:
 - Lag “s” seasonal differencing:

$$s_t^{(s)} = x_t - x_{t-5}$$

Models with Seasonal Differencing

- Seasonal “differencing” enables models to also consider speed of changes of values across a lag:

- lag “s” seasonal differencing:

$$S_t^{(s)} = x_t - x_{t-s}$$

- Seasonal ARIMA models also incorporate seasonal terms

- additive seasonal models

$$X_t = \text{ARIMA}(a, d, m) + \text{SEASONAL}_s(A, D, M)$$

- Multiplicative seasonal models

$$X_t = \text{ARIMA}(a, d, m) \times \text{SEASONAL}_s(A, D, M)$$

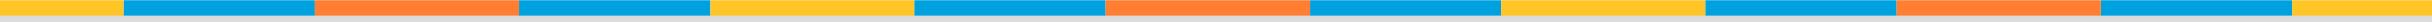
- A= # seasonal autoregressive terms, D= # seasonal differences, M = # seasonal moving average terms

Time Series



Model Discovery
(Box-Jenkins Procedure)

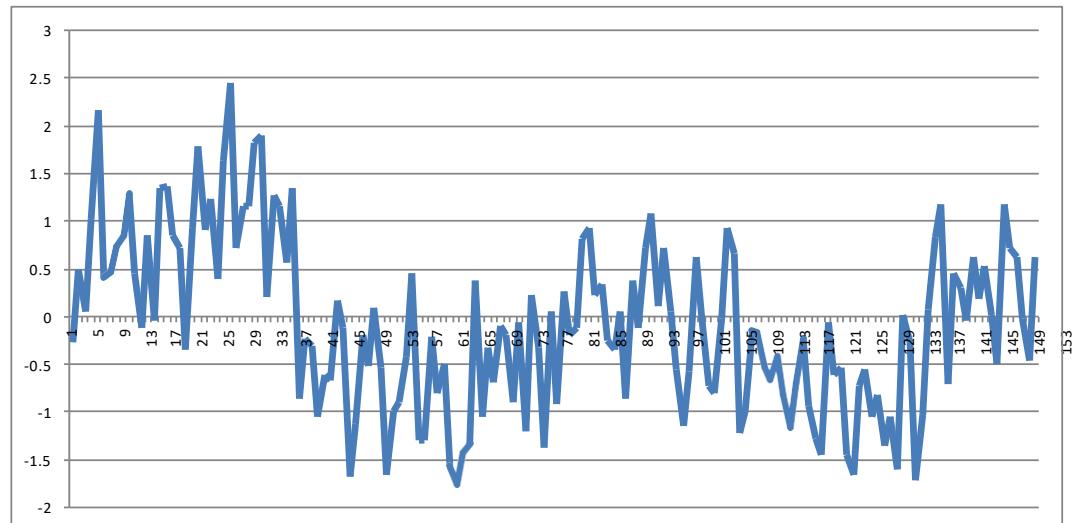
Model based Time Series Analysis



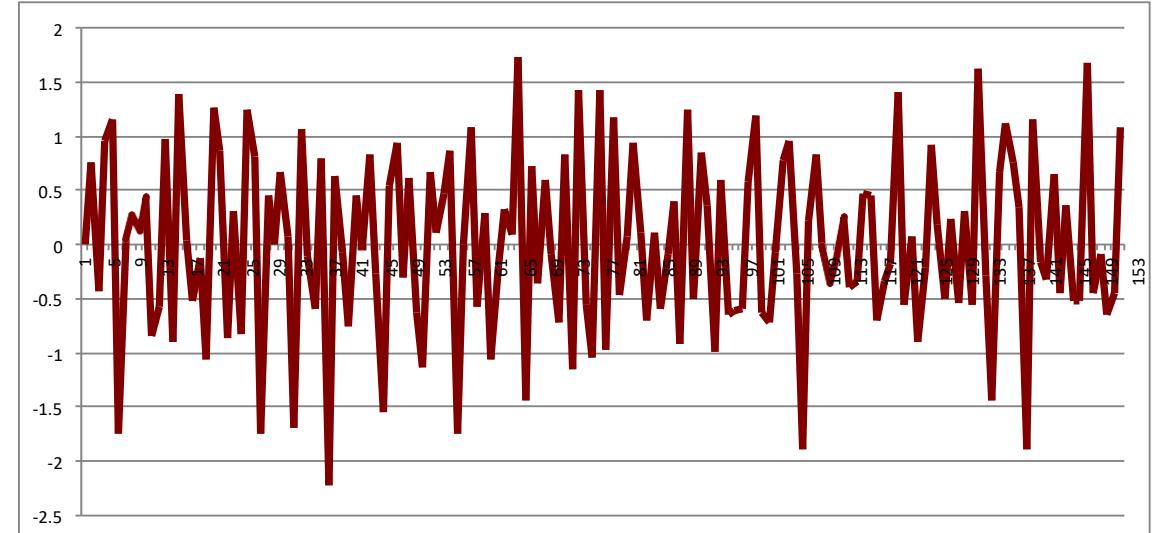
- Find the parameters of the model
 - Model fitting
 - the model should be as simple as possible (contain as few terms as possible)
 - the fit to historic data should be as good as possible

Box-Jenkins procedure

- Remove any **seasonal patterns** and **deterministic trends** that may hide valuable information and patterns **through differencing**
- When the mean trend is stochastic, **differencing the series may yield a stationary stochastic process** and, thus, may help convert a non-stationary series to a stationary one

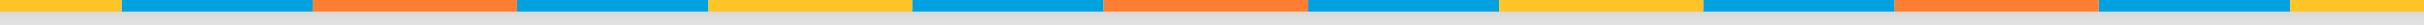


Original Series



Differenced Series

Box-Jenkins procedure



- Remove any **seasonal patterns** and **deterministic trends** that may hide valuable information and patterns **through differencing**
- When the mean trend is stochastic, **differencing the series may yield a stationary stochastic process** and, thus, may help convert a non-stationary series to a stationary one
- Series that can be rendered stationary by D time differencing are called “**integrated processes of order D**”

Model based Time Series Analysis

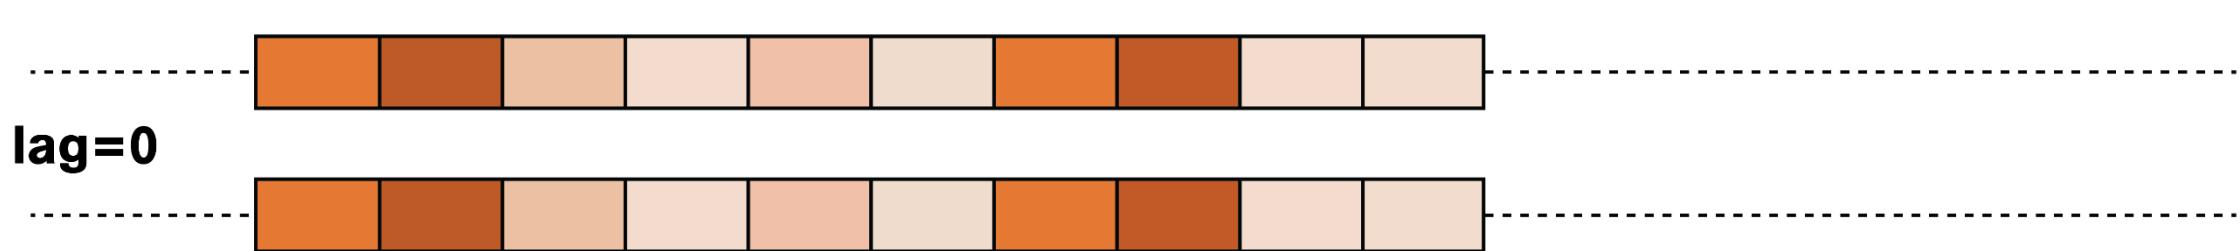


- Find the parameters of the model
 - Model fitting
 - the model should be as simple as possible (contain as few terms as possible)
 - the fit to historic data should be as good as possible

Model based Time Series Analysis

- Find the parameters of the model
 - Model fitting
 - Plot analysis
 - Autocorrelation Function (ACF) helps observe linear relationships between lagged values of a time series

$$ACF(X, lag) = \frac{E[(X_t - \mu)(X_{t+lag} - \mu)]}{\sigma^2} = \frac{\text{Covar}(X_t, X_{t+lag})}{\sigma^2}$$

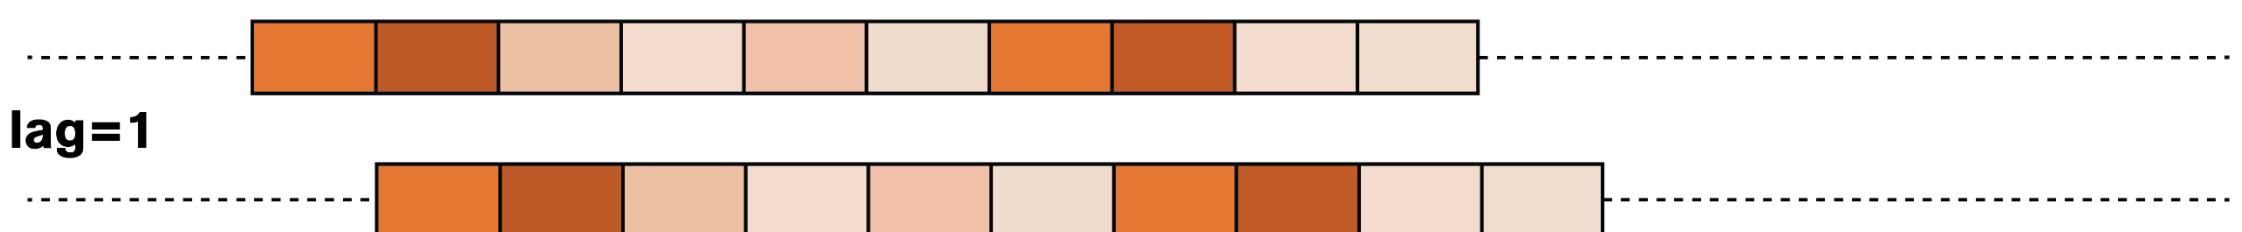


Model based Time Series Analysis

- Find the parameters of the model
 - Model fitting
 - Plot analysis
 - Autocorrelation Function (ACF) helps observe linear relationships between lagged values of a time series

$$\text{ACF}(X, \text{lag}) = \frac{E[(X_t - \mu)(X_{t+lag} - \mu)]}{\sigma^2} = \frac{\text{Covar}(X_t, X_{t+lag})}{\sigma^2}$$

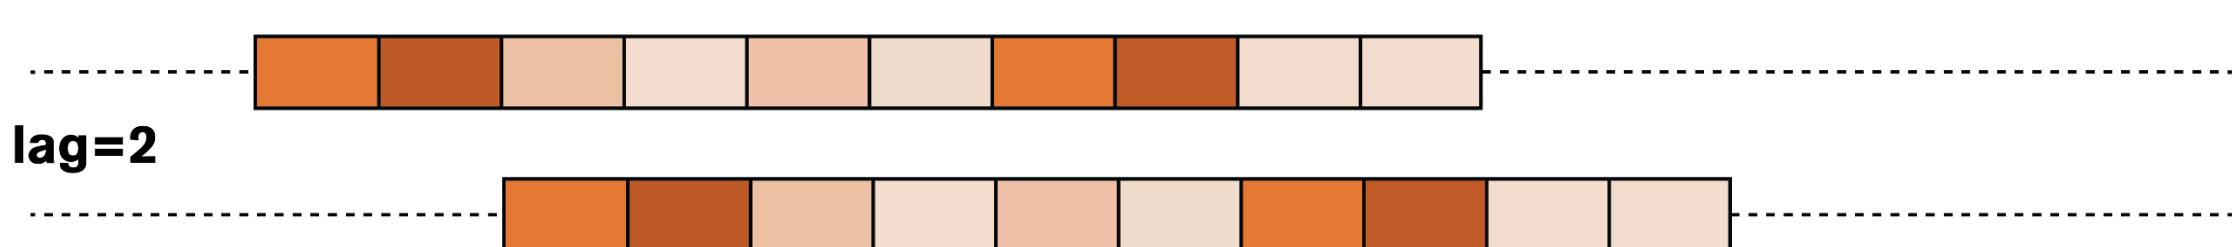
=



Model based Time Series Analysis

- Find the parameters of the model
 - Model fitting
 - Plot analysis
 - Autocorrelation Function (ACF) helps observe linear relationships between lagged values of a time series

$$ACF(X, lag) = \frac{E[(X_t - \mu)(X_{t+lag} - \mu)]}{\sigma^2} = \frac{\text{Covar}(X_t, X_{t+lag})}{\sigma^2}$$

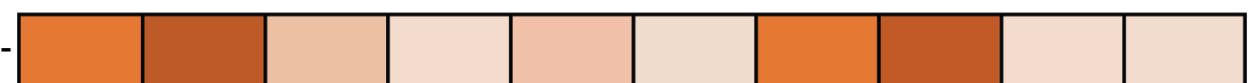
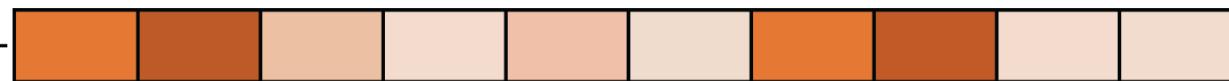


Model based Time Series Analysis

- Find the parameters of the model
 - Model fitting
 - Plot analysis
 - Autocorrelation Function (ACF) helps observe linear relationships between lagged values of a time series

$$\text{ACF}(X, \text{lag}) = \frac{E[(X_t - \mu)(X_{t+\text{lag}} - \mu)]}{\sigma^2} = \frac{\text{Covar}(X_t, X_{t+\text{lag}})}{\sigma^2}$$

lag=7



Model based Time Series Analysis

- Find the parameters of the model
 - Model fitting
 - Plot analysis
 - Autocorrelation Function (ACF) helps observe linear relationships between lagged values of a time series

$$\text{ACF}(X, \text{lag}) = \frac{E[(X_t - \mu)(X_{t+\text{lag}} - \mu)]}{\sigma^2} = \frac{\text{Covar}(X_t, X_{t+\text{lag}})}{\sigma^2}$$

- Partial Autocorrelation Function (PACF) discount the impact of intermediate values

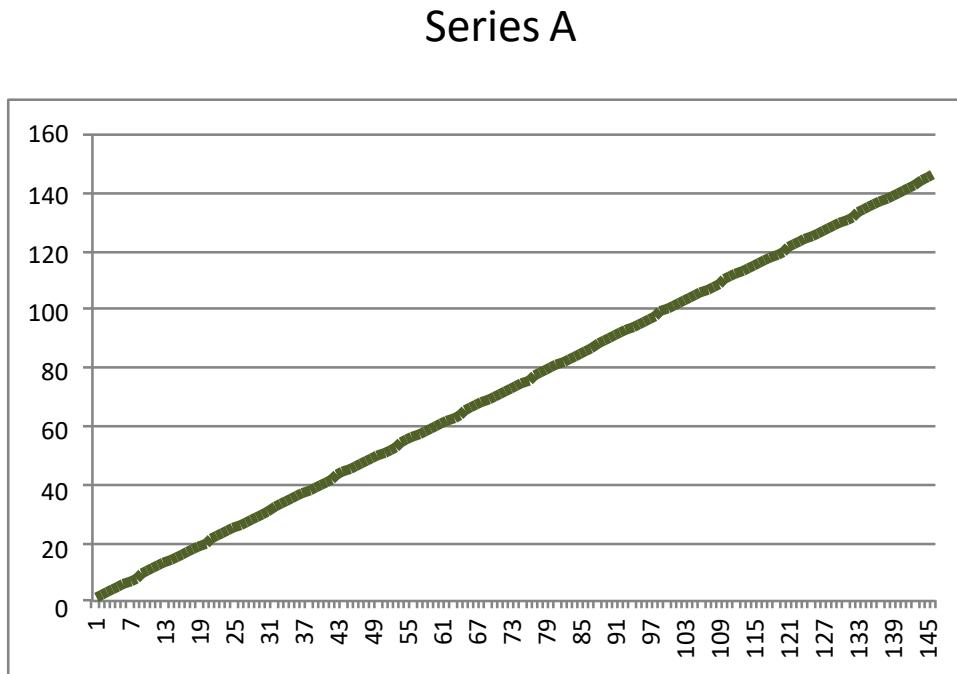
$$\text{PACF}(X, \text{lag}) = \frac{\text{Covar}(X_t, X_{t+\text{lag}} \mid X_{t+1}, \dots, X_{t+\text{lag}-1})}{\sigma^2}$$

Key Properties of ACF and PACF



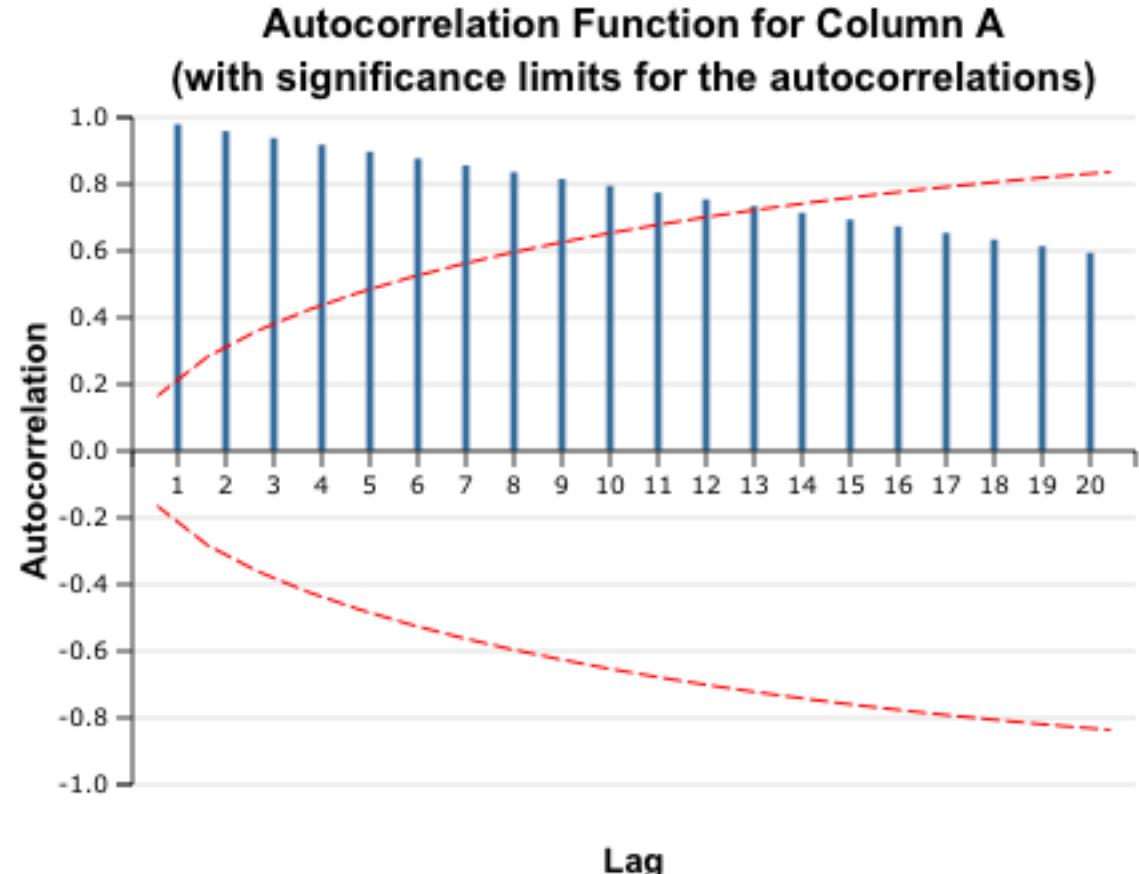
- For trend series autocorrelation function (ACF) slowly decays.

Autoregressive model



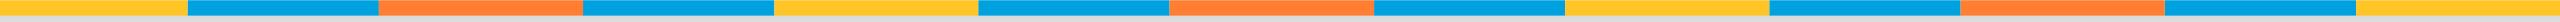
$$X_t = X_{t-1} + E_t$$

An AR(1) model



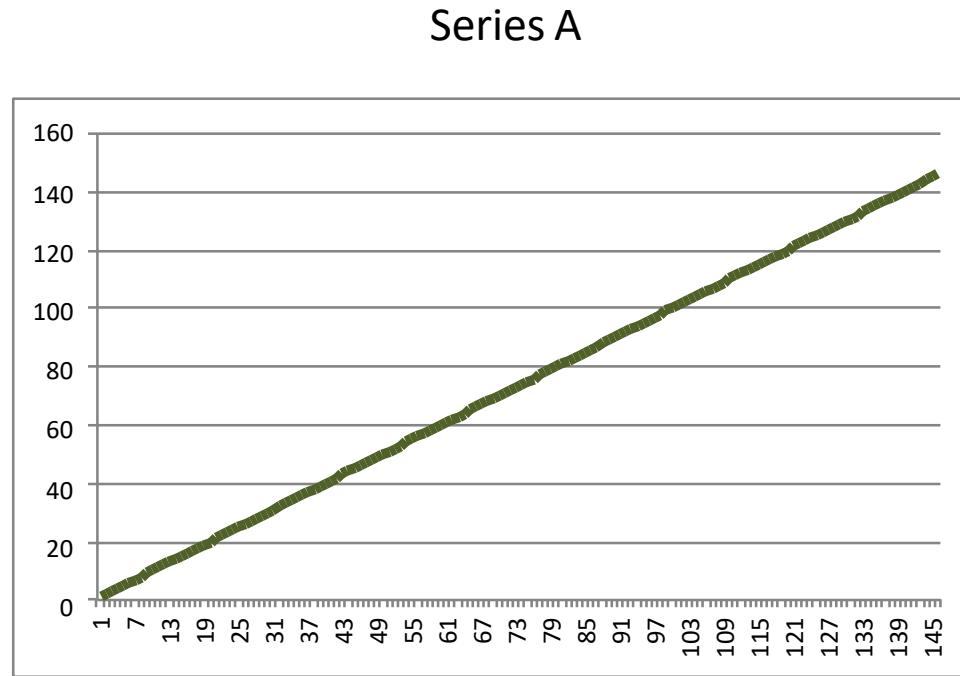
- Slowly decaying ACF indicates a trend
- If the series has positive autocorrelations for high number of lags, then it may need differencing to remove the trend or seasonal patterns

Key Properties of ACF and PACF



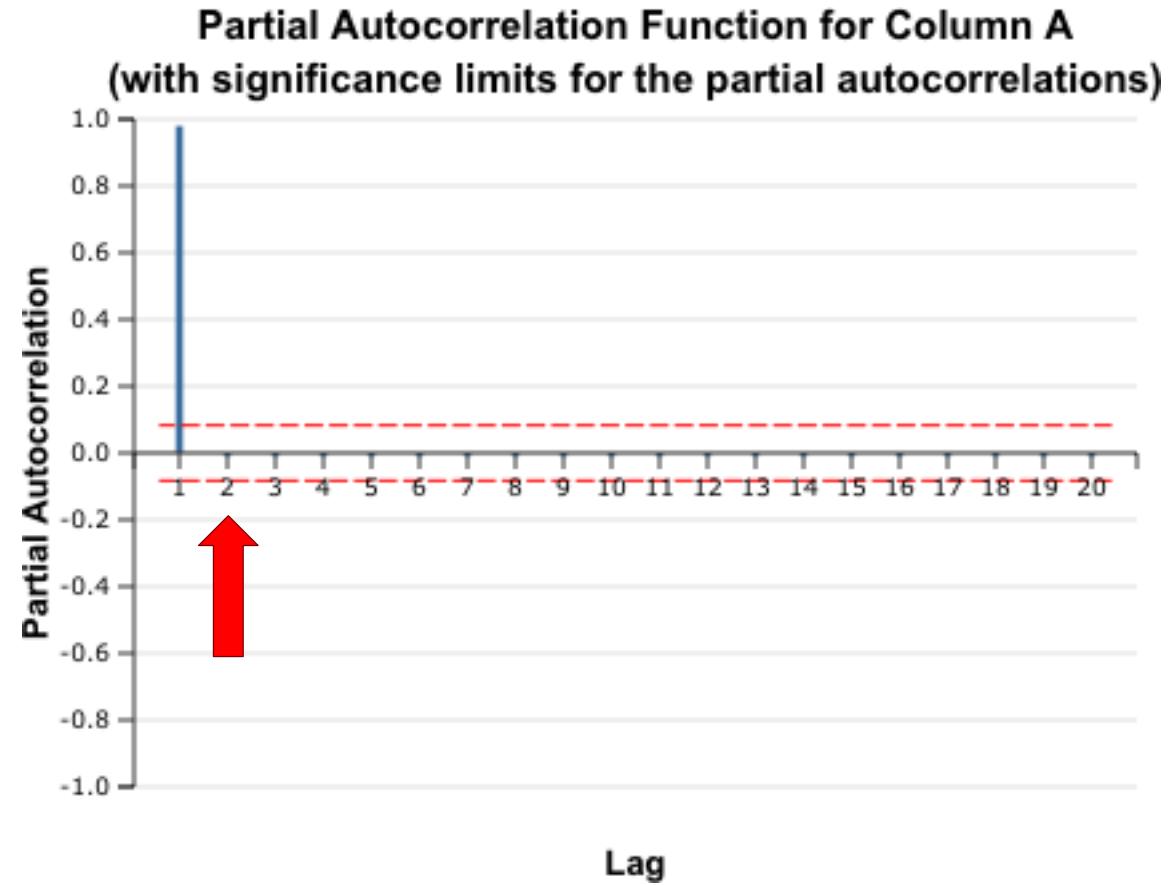
- For trend series autocorrelation function (ACF) slowly decays.
- For an autoregressive, AR(a), series, the partial autocorrelation function (PACF) gives 0 at lag $\geq a + 1$

Autoregressive model



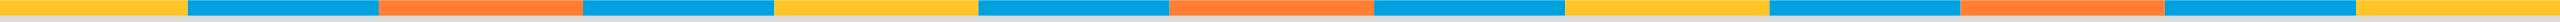
$$X_t = X_{t-1} + E_t$$

An AR(1) model



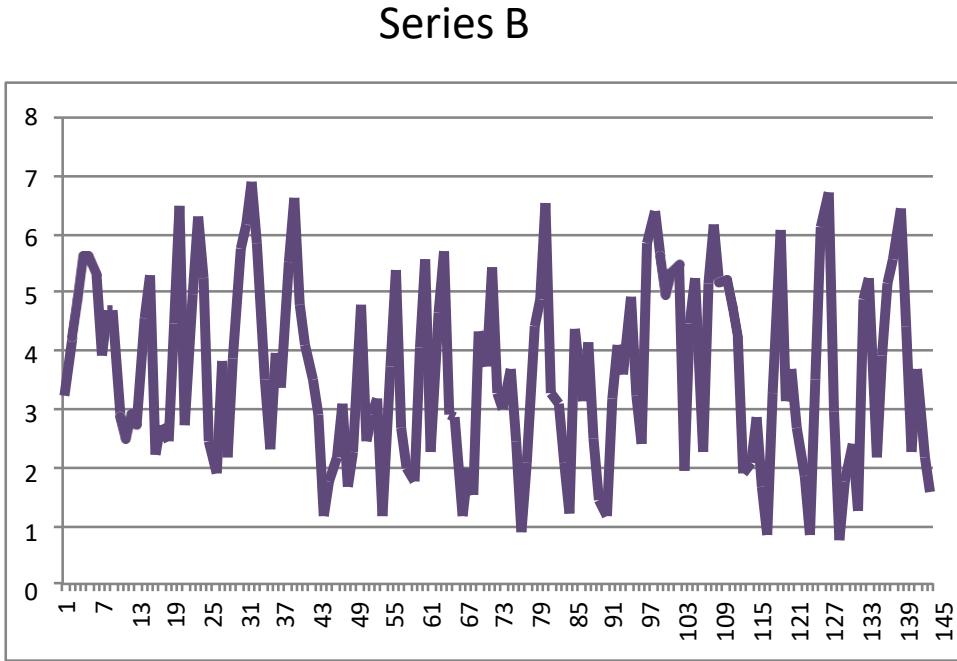
- Partial autocorrelation of an AR(a) process is zero at lag $\geq a + 1$

Key Properties of ACF and PACF



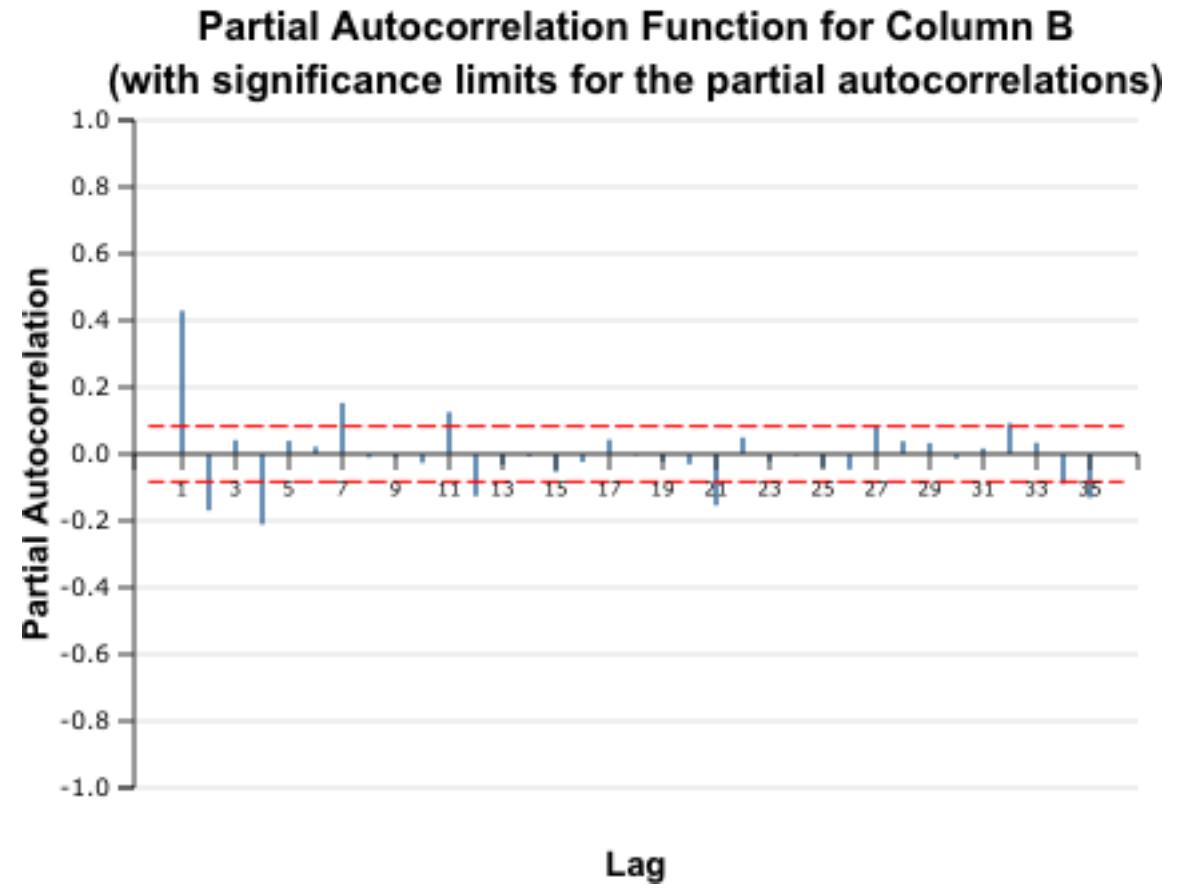
- For an autoregressive, AR(a), series, the partial autocorrelation function (PACF) gives 0 at lag $\geq a + 1$
- For a moving average, MA(m), series,
 - the partial autocorrelation function (PACF) does not “shut off” at a fixed lag, but moves toward 0.

Moving average model



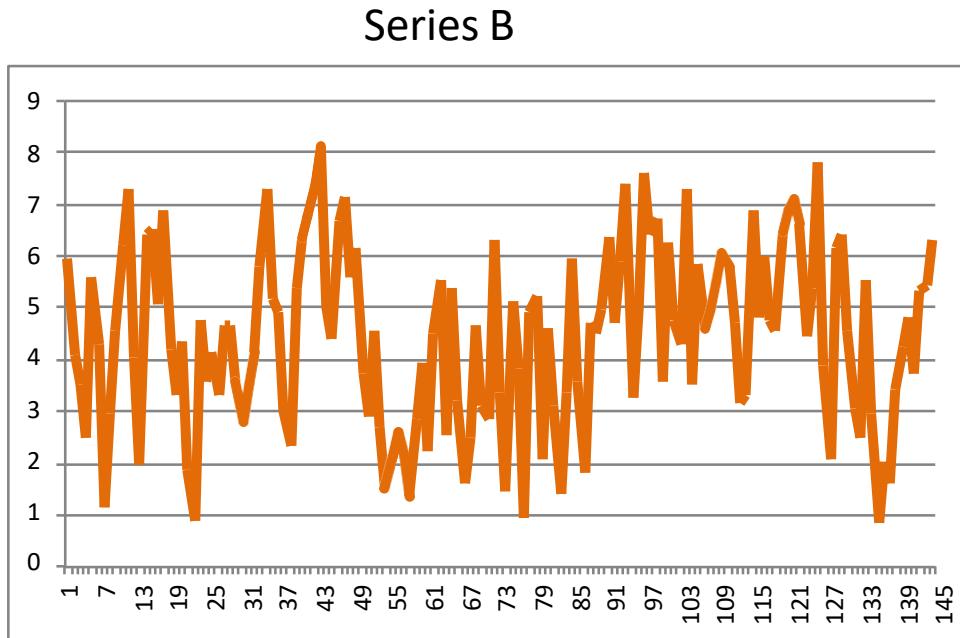
$$X_t = 0.5E_{t-1} + E_t$$

An MA(1) model



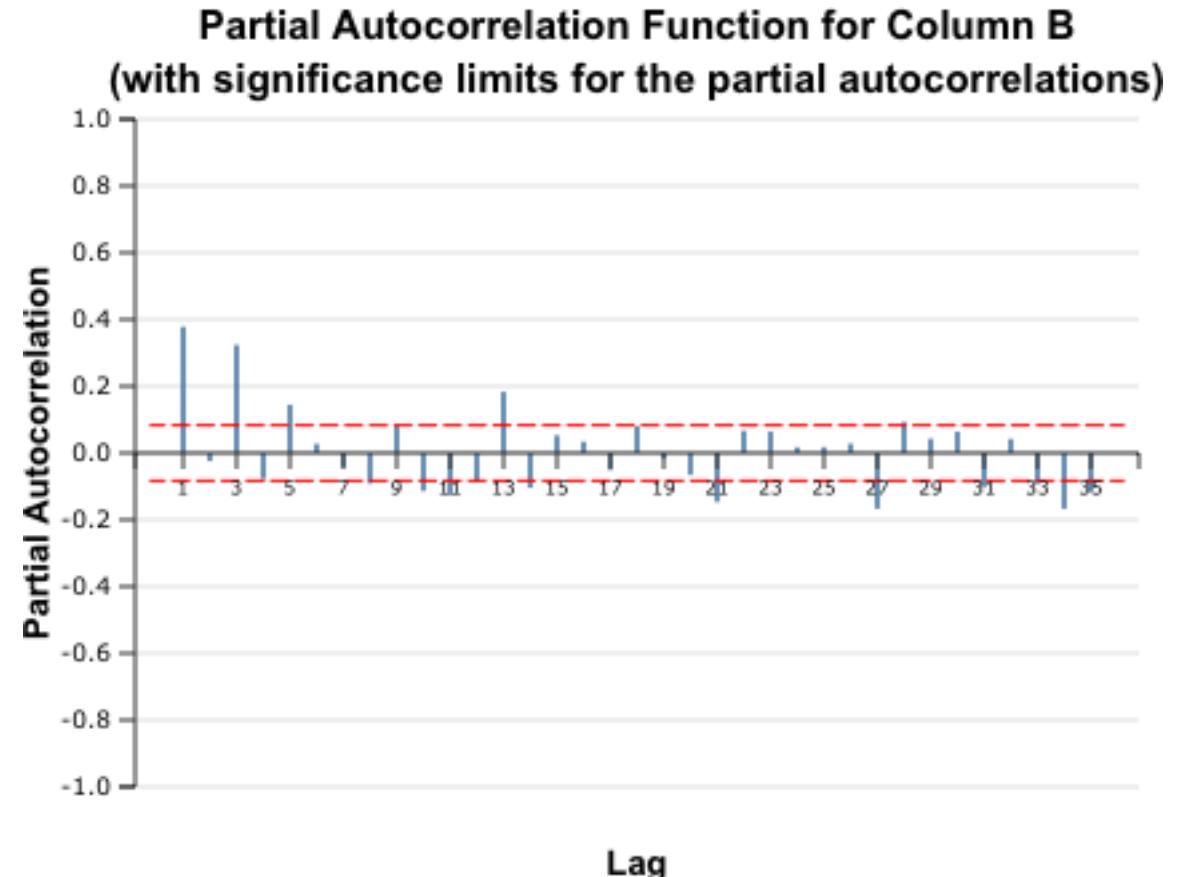
- Partial autocorrelation of an MA(m) process does not “shut off” at a fixed lag, but moves toward 0

Moving average model



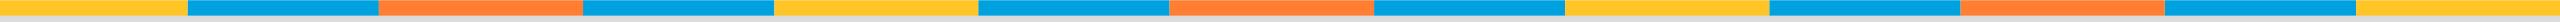
$$X_t = 0.5E_{t-1} + 0.3E_{t-3} + E_t$$

An MA(3) model



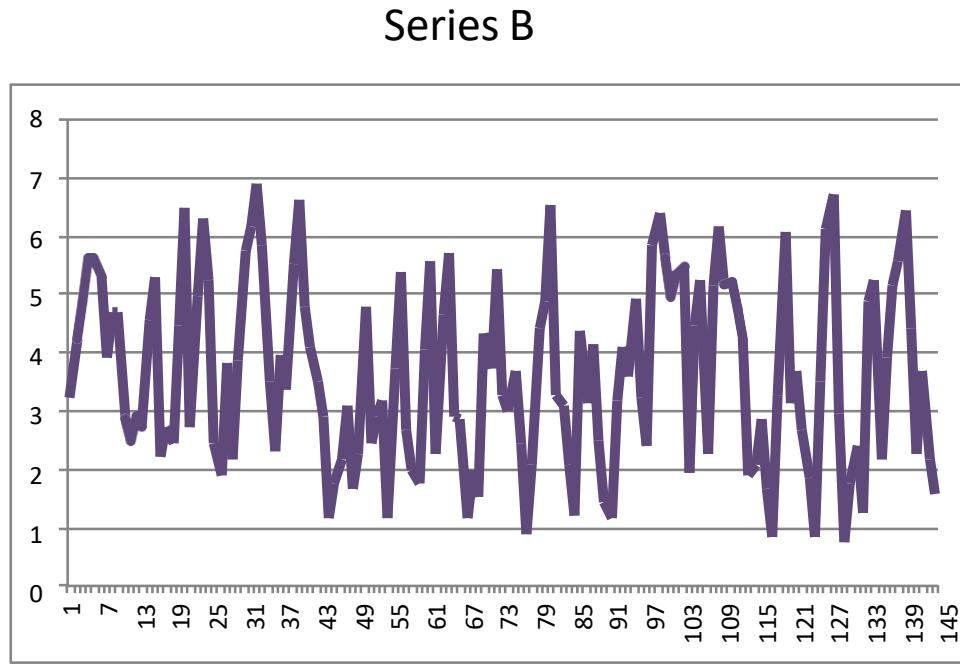
- Partial autocorrelation of an MA(m) process does not “shut off” at a fixed lag, but moves toward 0

Key Properties of ACF and PACF



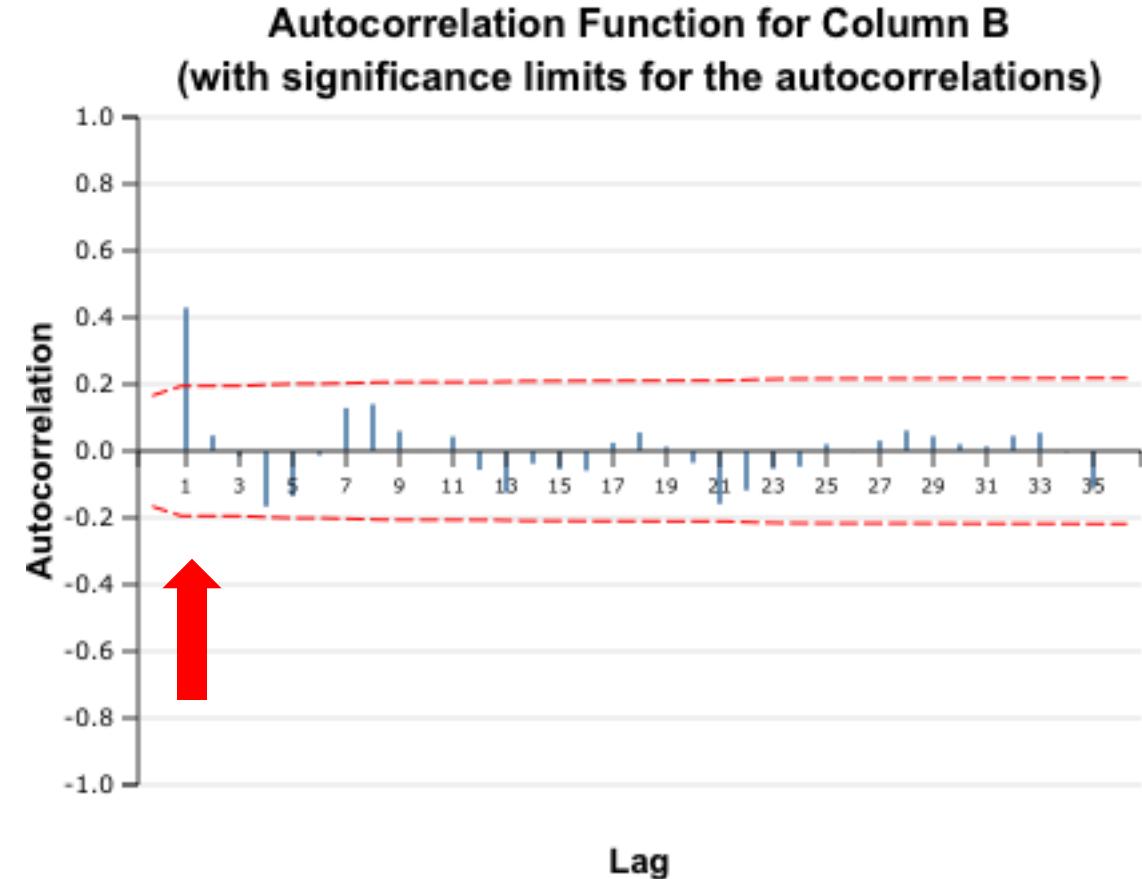
- For an autoregressive, AR(a), series, the partial autocorrelation function (PACF) gives 0 at lag $\geq a + 1$
- For a moving average, MA(m), series,
 - the partial autocorrelation function (PACF) does not “shut off” at a fixed lag, but moves toward 0.
 - the autocorrelation function (ACF) has non-zero autocorrelation only at the lags of the model.

Moving average model



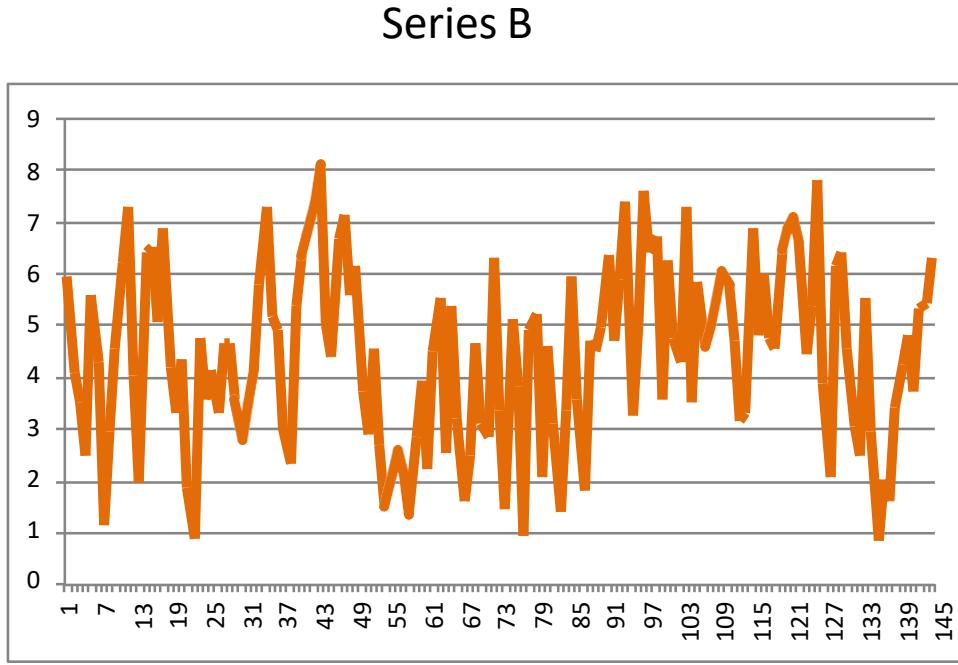
$$X_t = 0.5E_{t-1} + E_t$$

An MA(1) model



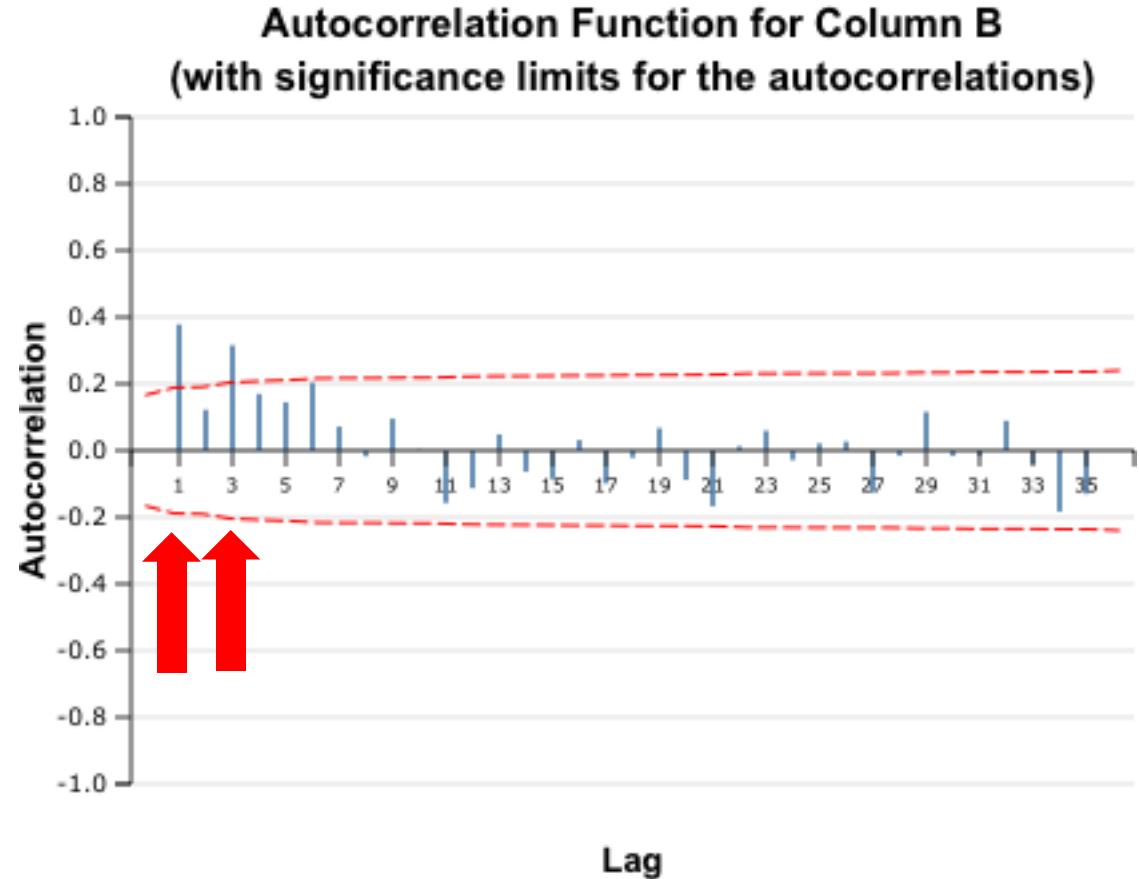
- Autocorrelation of an MA(m) process has non-zero autocorrelation only at the lags of the model

Moving average model



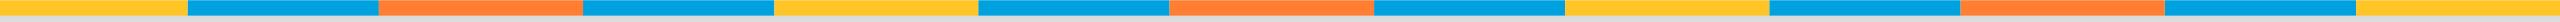
$$X_t = 0.5E_{t-1} + 0.3E_{t-3} + E_t$$

An MA(3) model



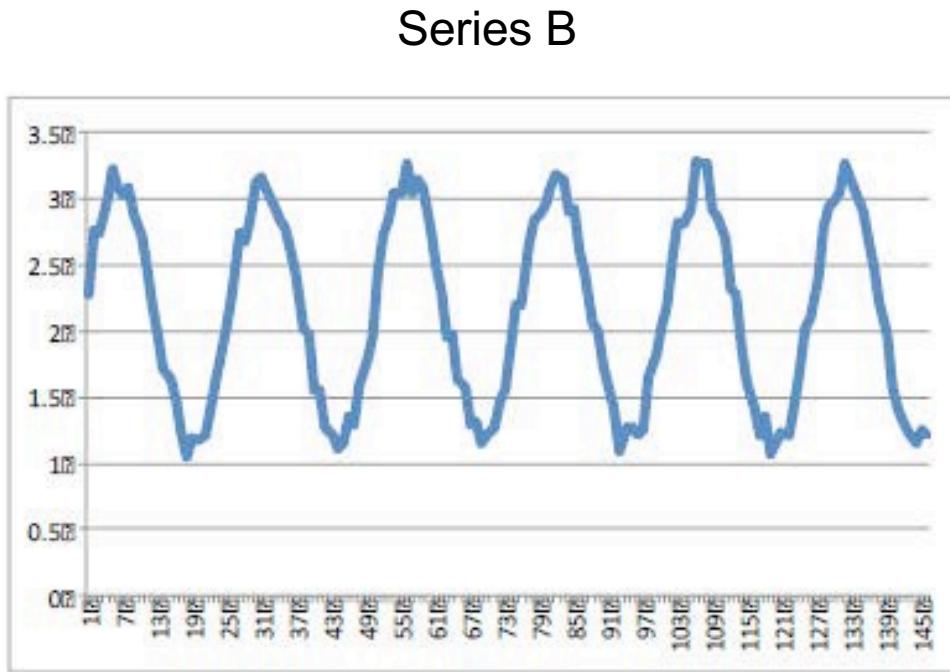
- Autocorrelation of an MA(m) process has non-zero autocorrelation only at the lags of the model

Key Properties of ACF and PACF

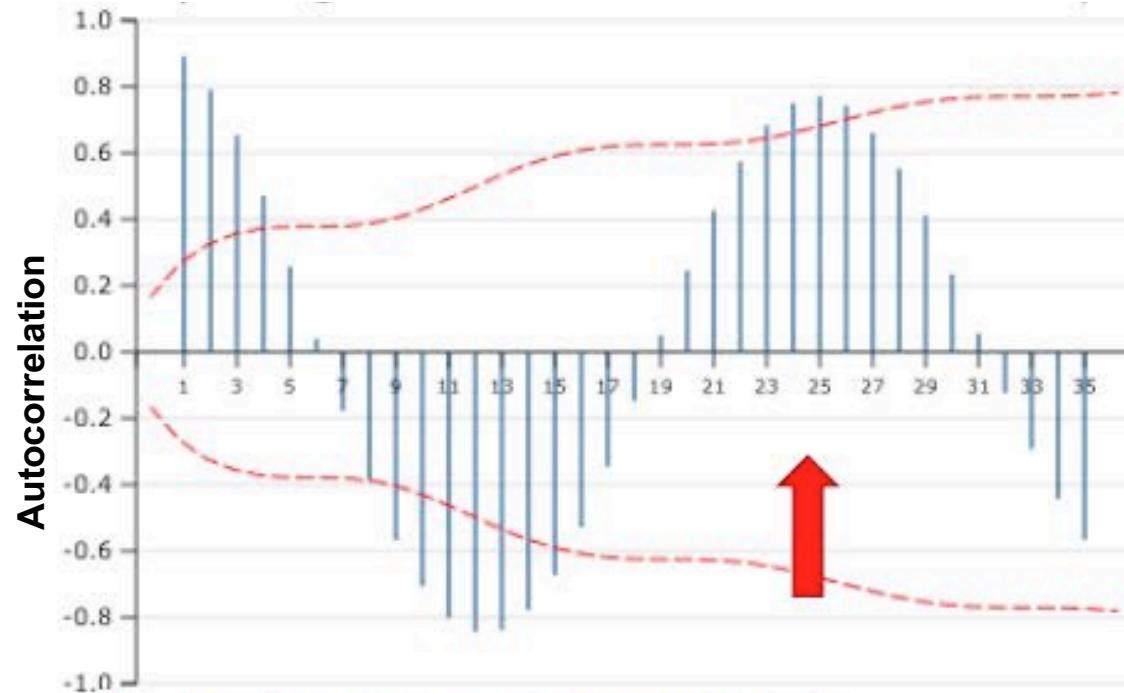


- For trend series autocorrelation function (ACF) slowly decays.
- For an autoregressive, AR(a), series, the partial autocorrelation function (PACF) gives 0 at lag $\geq a + 1$
- For a moving average, MA(m), series,
 - the partial autocorrelation function (PACF) does not “shut off” at a fixed lag, but moves toward 0.
 - the autocorrelation function (ACF) has non-zero autocorrelation only at the lags of the model.
- For **seasonal series** autocorrelation function **(ACF) at the seasonal lag will be large and positive.**

Non-stationary (cyclic) model

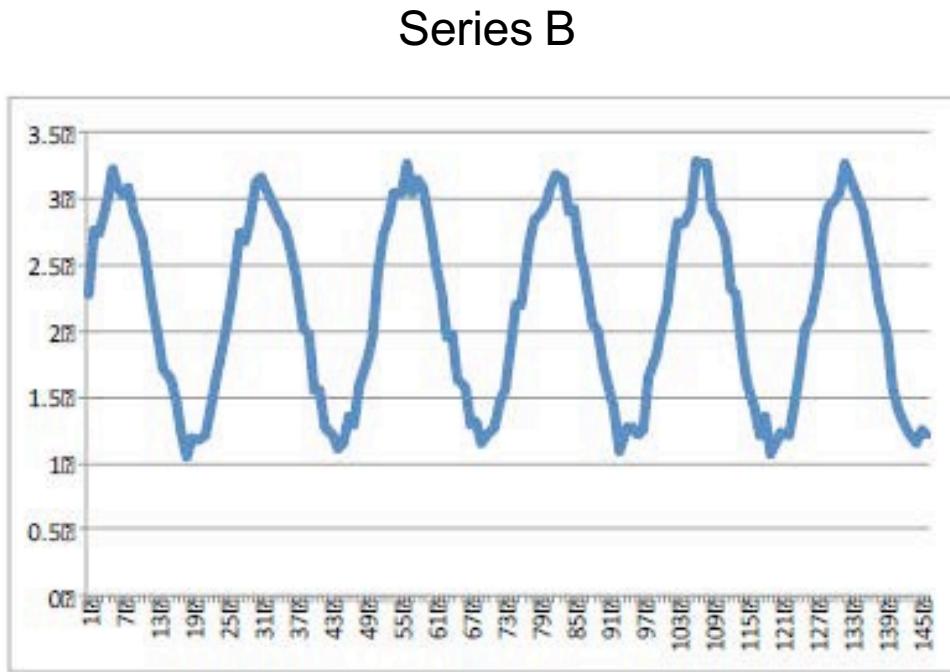


Autocorrelation Function for Column B
(with significance limits for the autocorrelations)

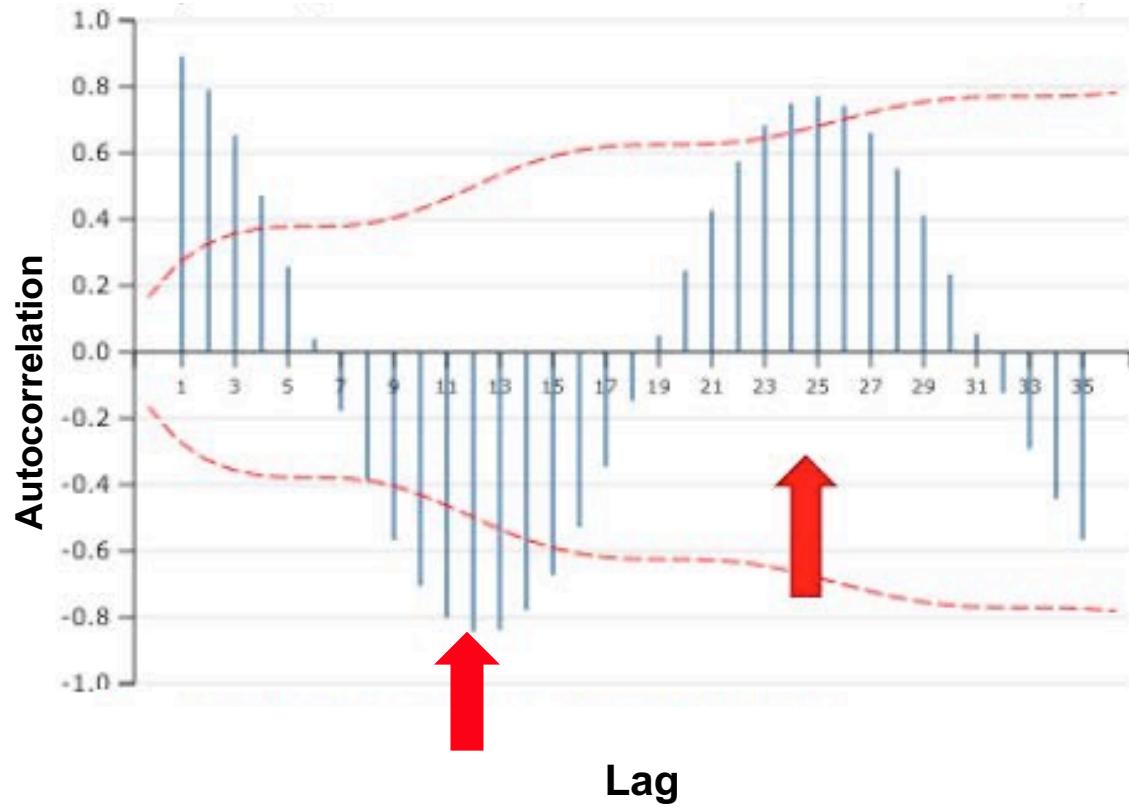


- No decay to zero – not stationary
- If the series has positive autocorrelations for high number of lags, then it may need differencing to remove the trend or seasonal patterns
- High values at fixed interval – seasonal autoregressive term
- (seasonal pattern with period of 24)

Non-stationary (cyclic) model

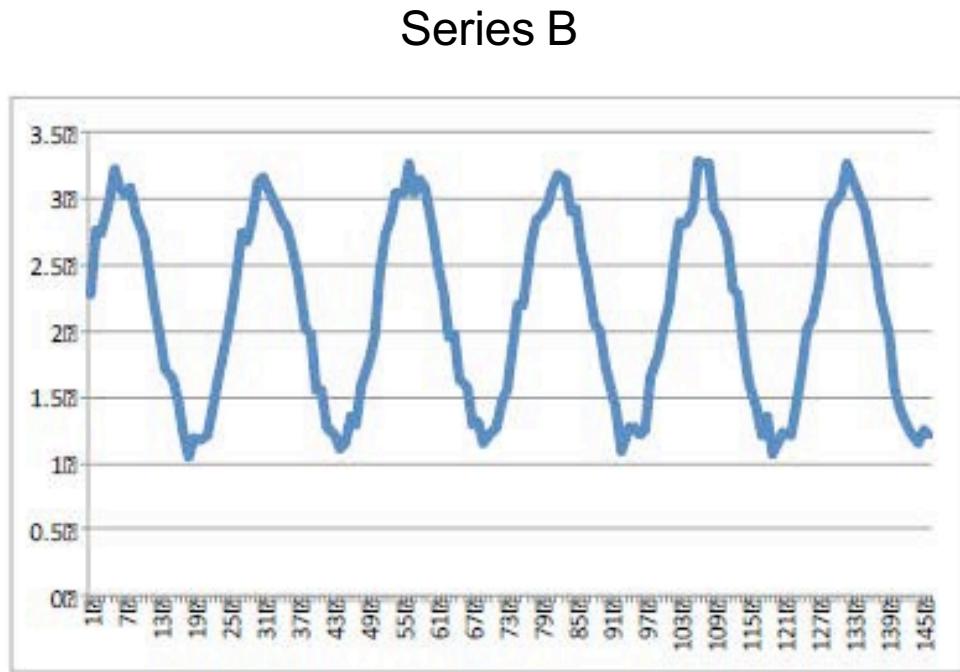


Autocorrelation Function for Column B
(with significance limits for the autocorrelations)

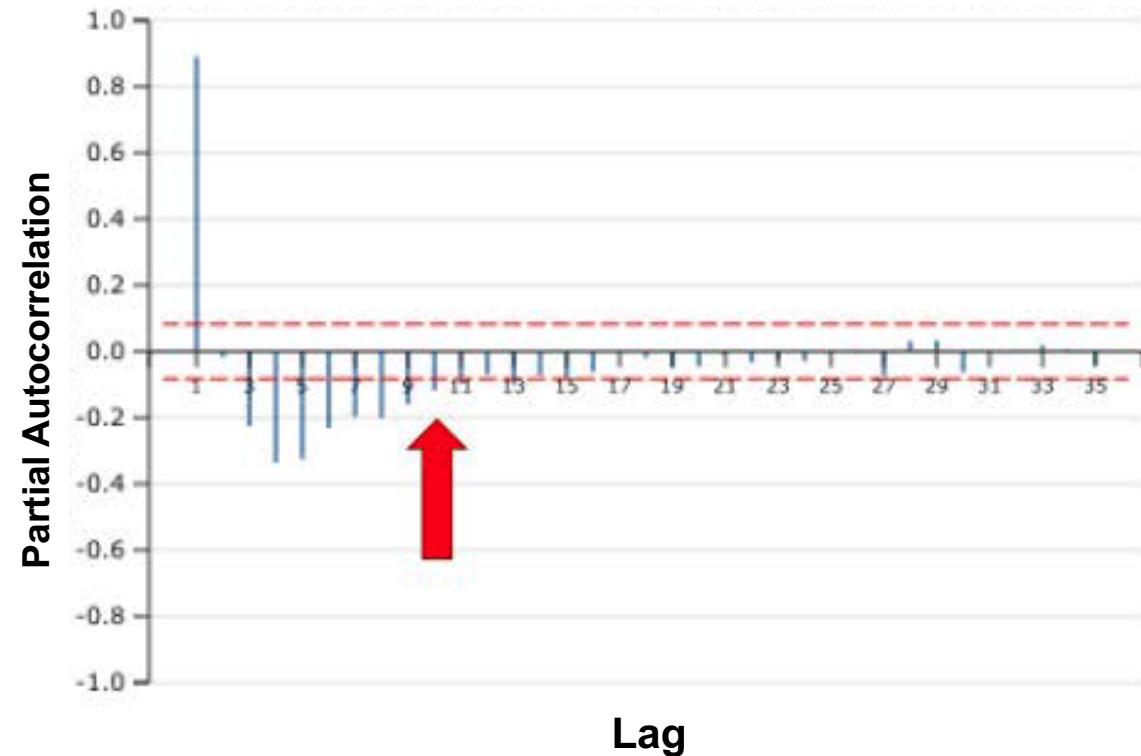


- The overall autocorrelation pattern indicates a cyclic pattern with period of 24

Non-stationary (cyclic) model

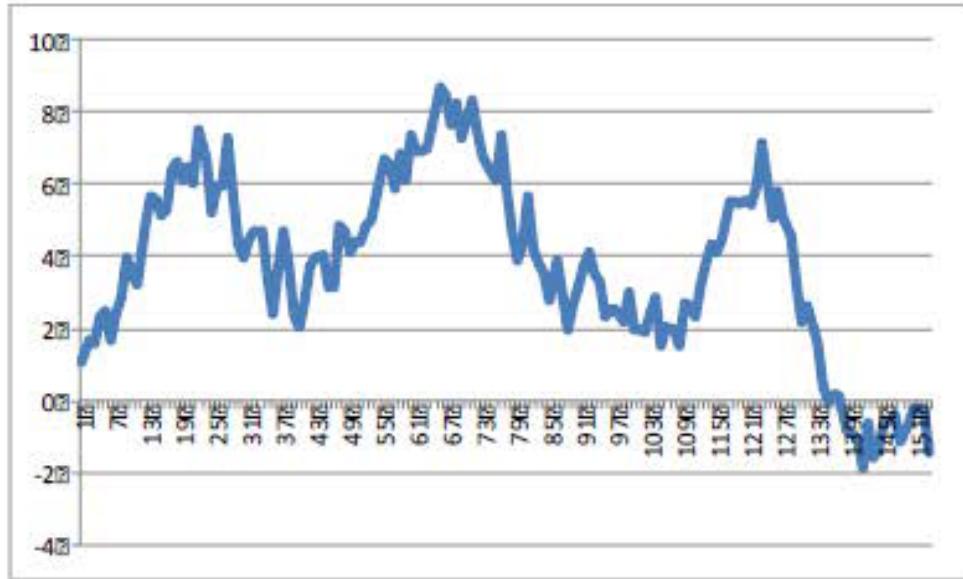


Partial Autocorrelation Function for Column B
(with significance limits for the partial autocorrelations)

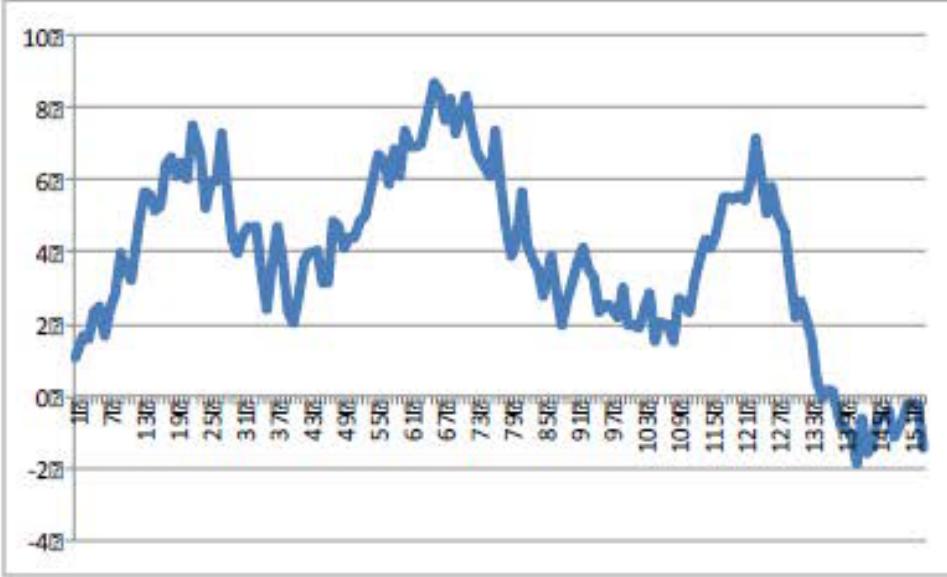


- The partial autocorrelation pattern may indicate a cyclic pattern with 12 unit half lag

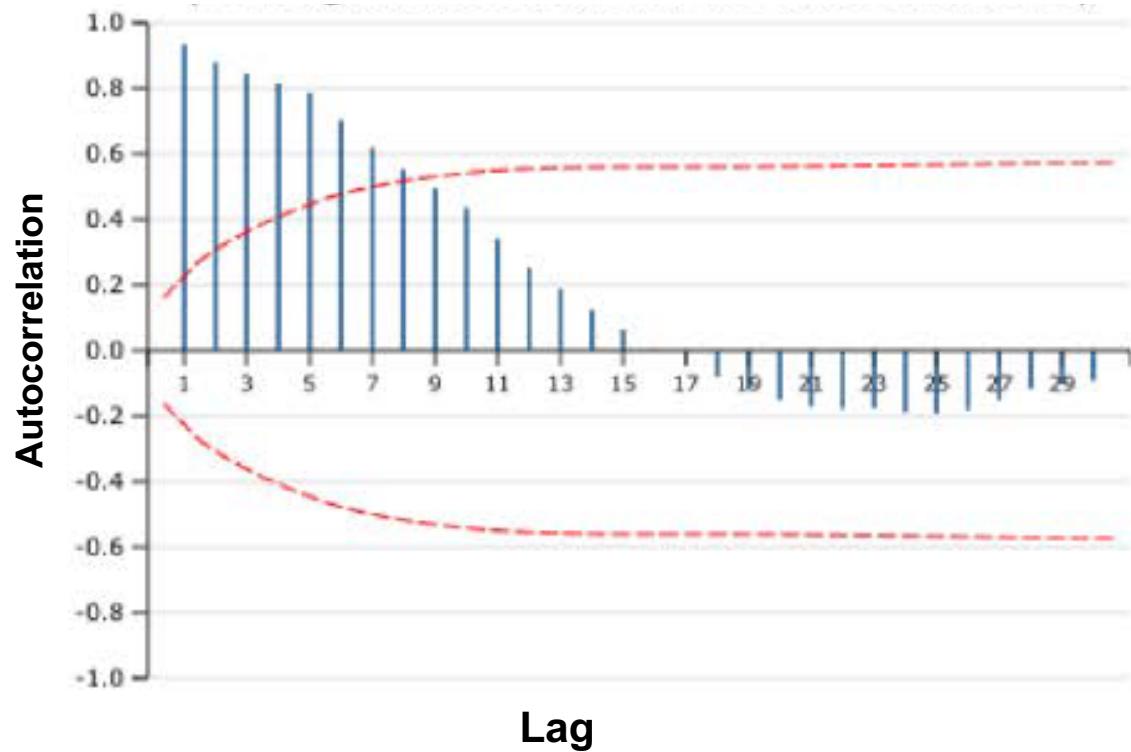
Complex model



Complex model

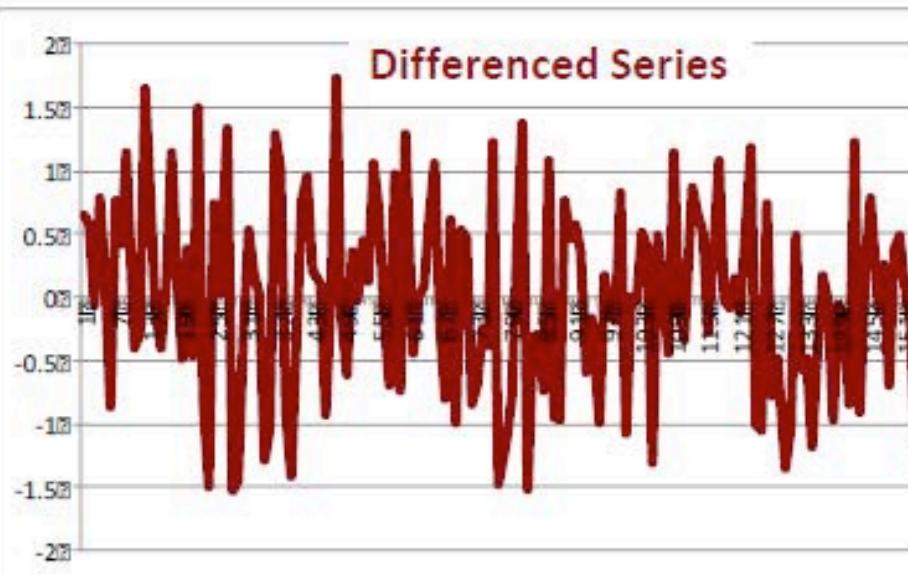
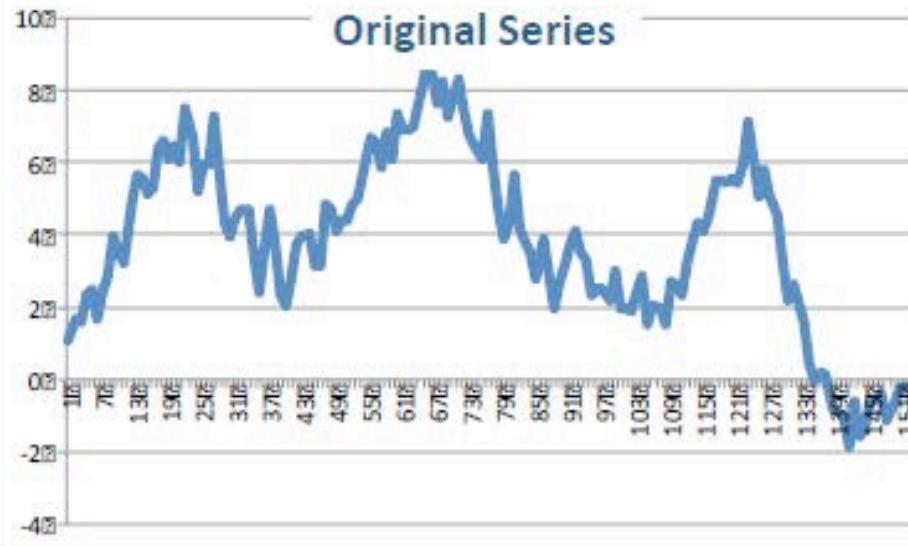


**Autocorrelation Function for Column A
(with significance limits for the autocorrelations)**

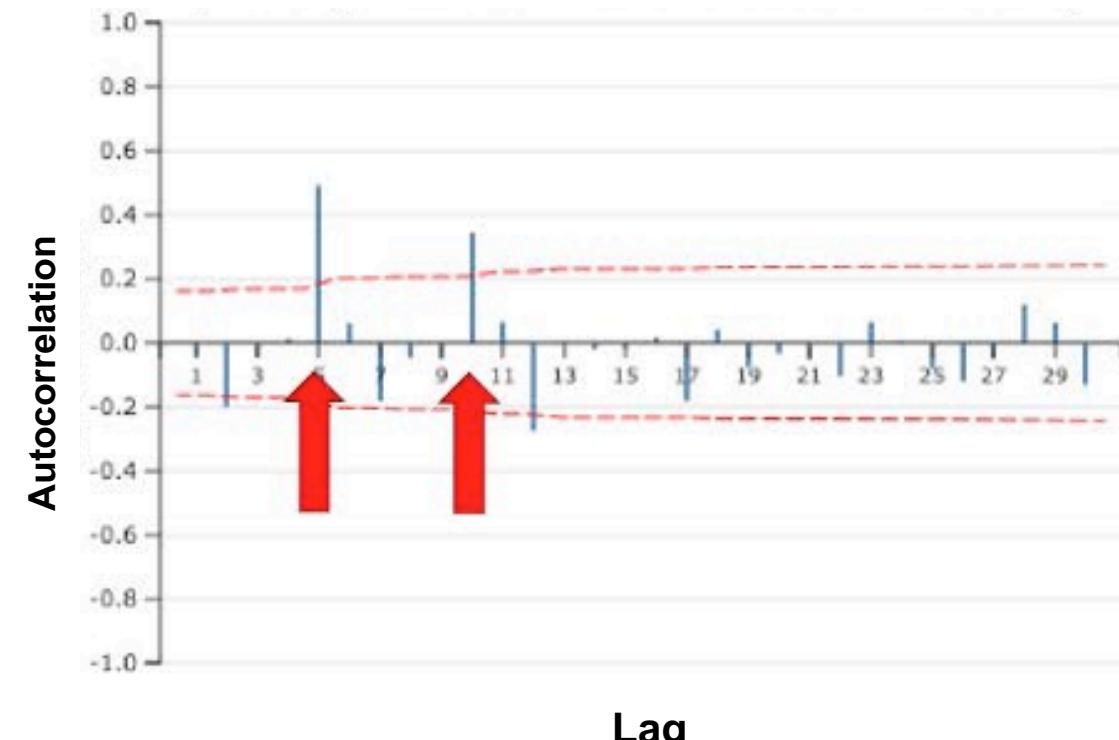


- Slowly decaying ACF indicates a trend
- if the series has positive autocorrelation for high number of lags, then it may need differencing to remove the trend or seasonal patterns

Complex model

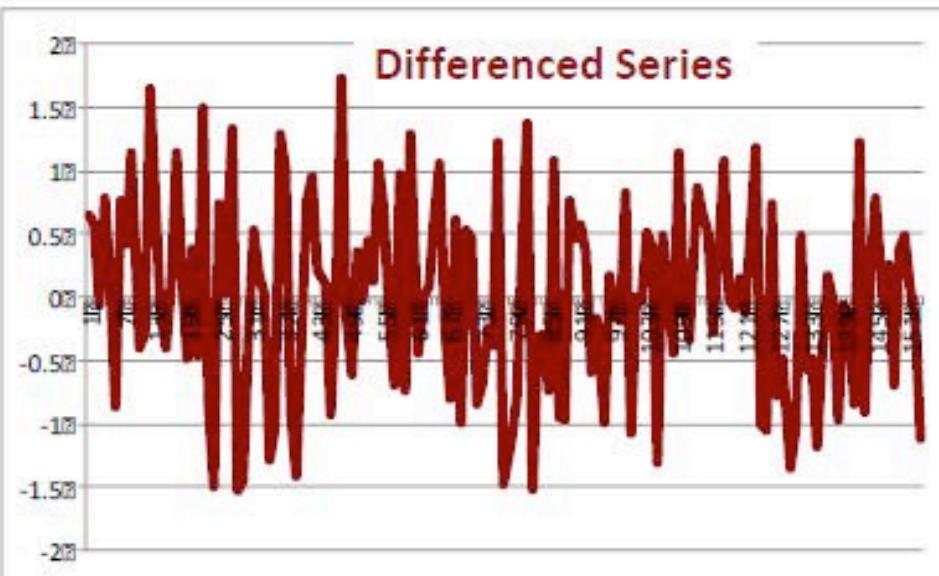
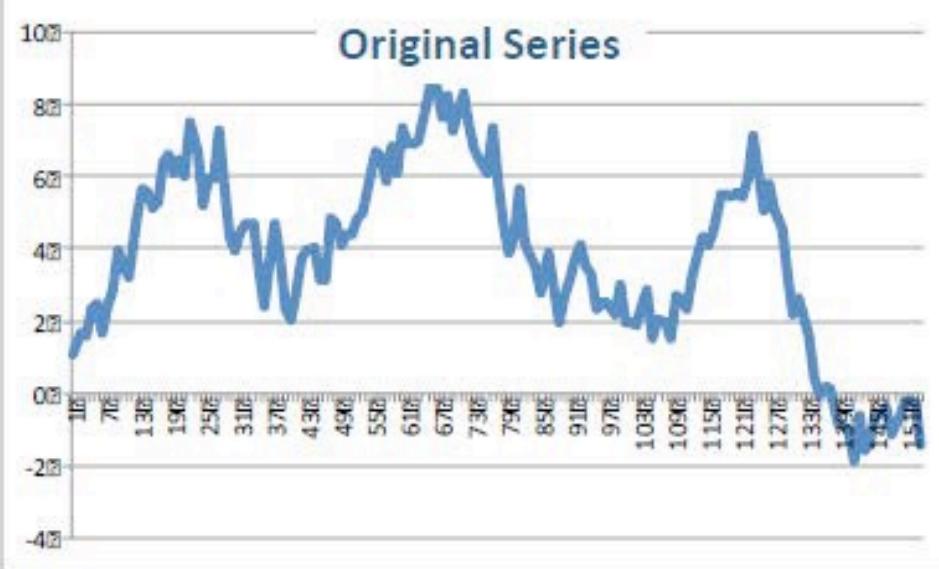


Autocorrelation Function for Column B
(with significance limits for the autocorrelations)

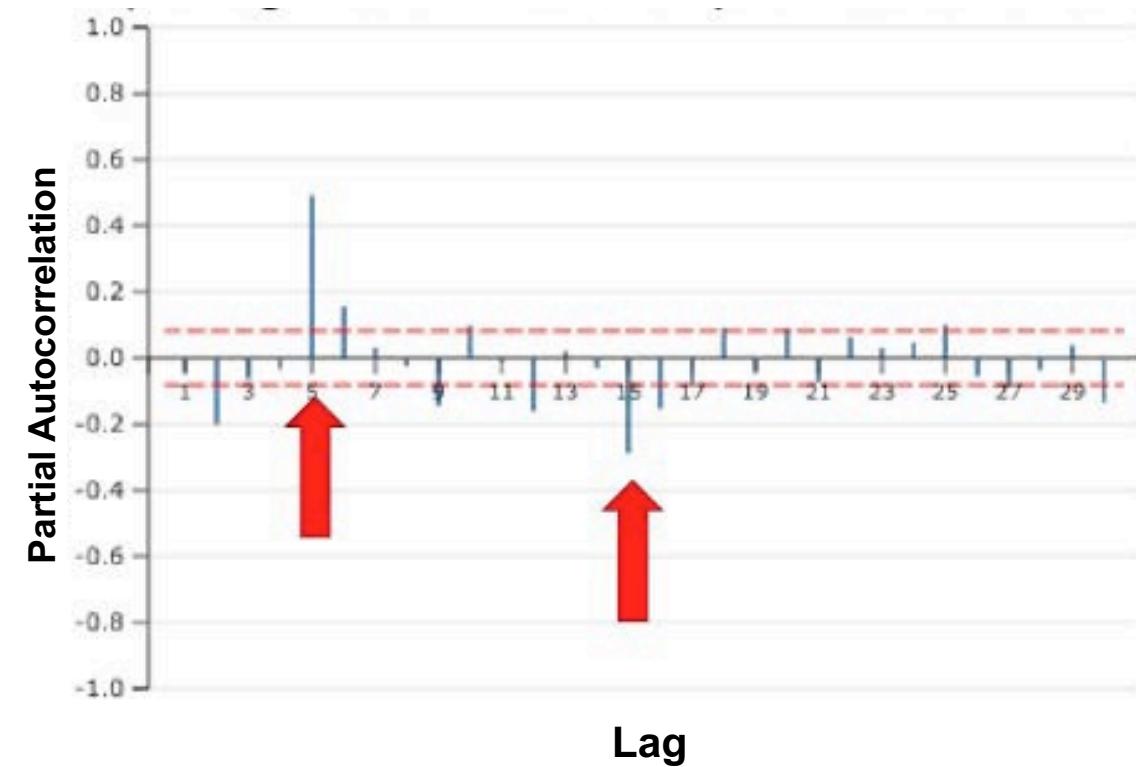


- Autocorrelation of an MA(m) process has non-zero autocorrelation only at the lags of the model
- Indicates MA(5) and MA(10) components.

Complex model

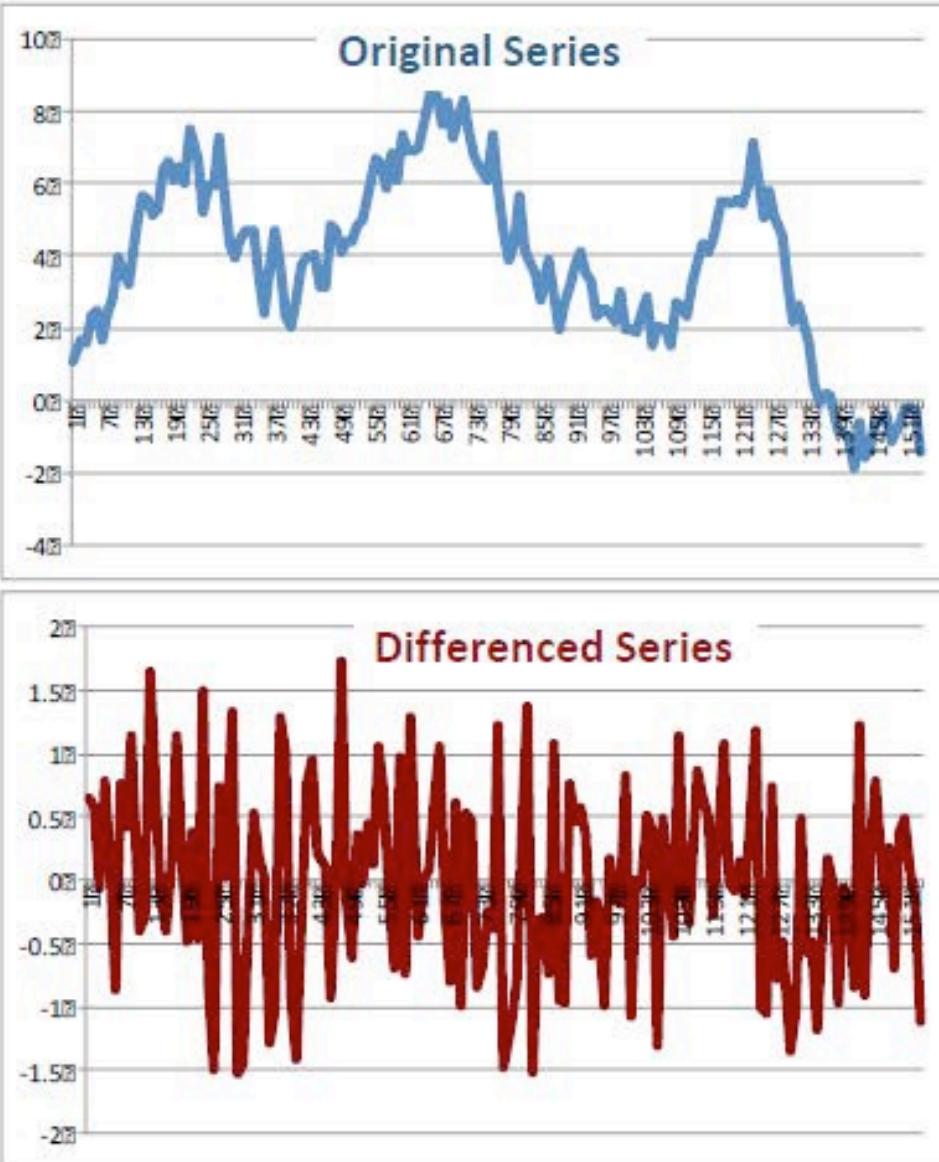


**Partial Autocorrelation Function for Column B
(with significance limits for the partial autocorrelations)**



- Partial autocorrelation of an MA(m) process does not “shut off” at a fixed lag, but moves toward 0

Complex model



- The differenced series appears to be a moving average series with MA(5) and MA(10) components.

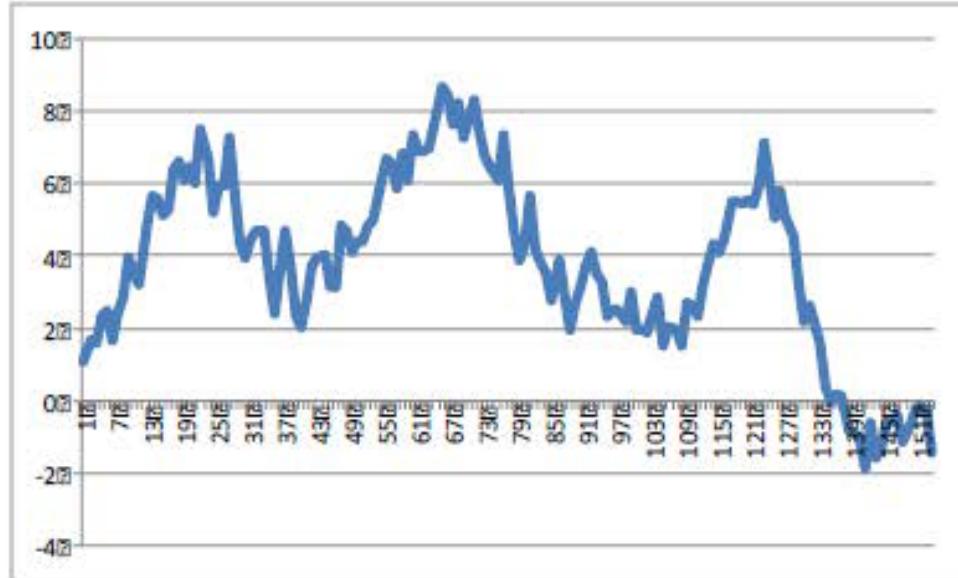
$$D_t = X_t - X_{t-1} = \alpha E_{t-5} + \beta E_{t-10} + E_t$$

$$X_t = X_{t-1} + \alpha E_{t-5} + \beta E_{t-10} + E_t$$

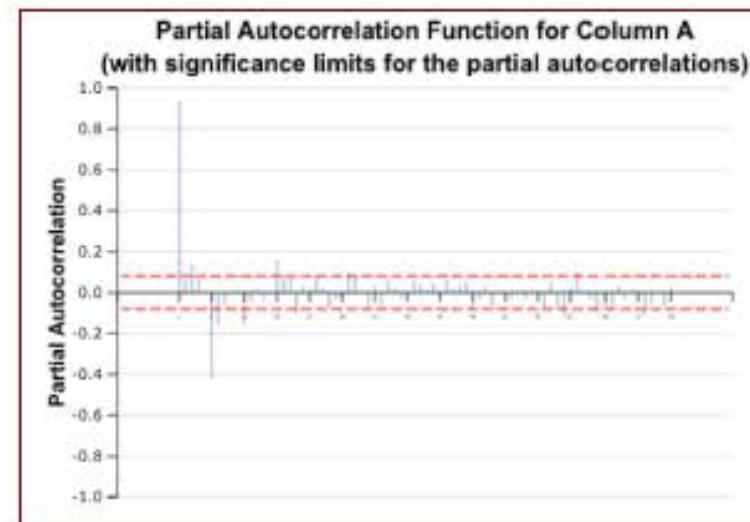
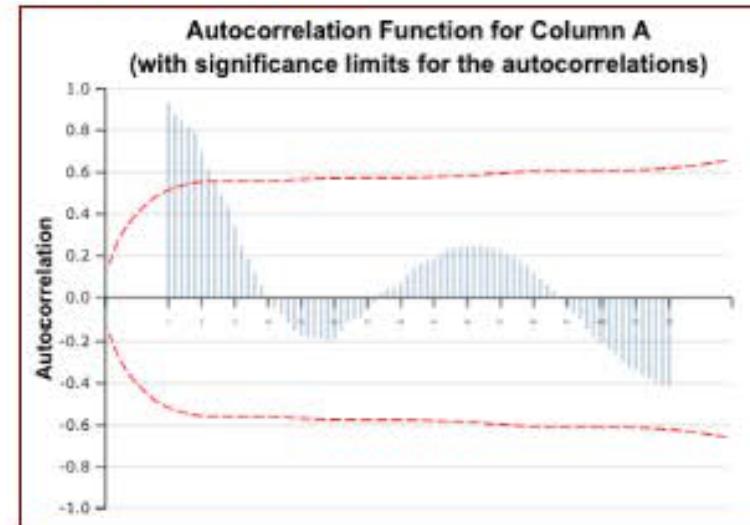
$$X_t = X_{t-1} + 0.5 E_{t-5} + 0.5 E_{t-10} + E_t$$

Moving average model

Original Series



- But isn't this series cyclic?



$$X_t = X_{t-1} + 0.5E_{t-5} + 0.5E_{t-10} + E_t$$



Sequences and Time Series

Time Series Matching

K. Selçuk Candan, Professor of Computer Science and Engineering

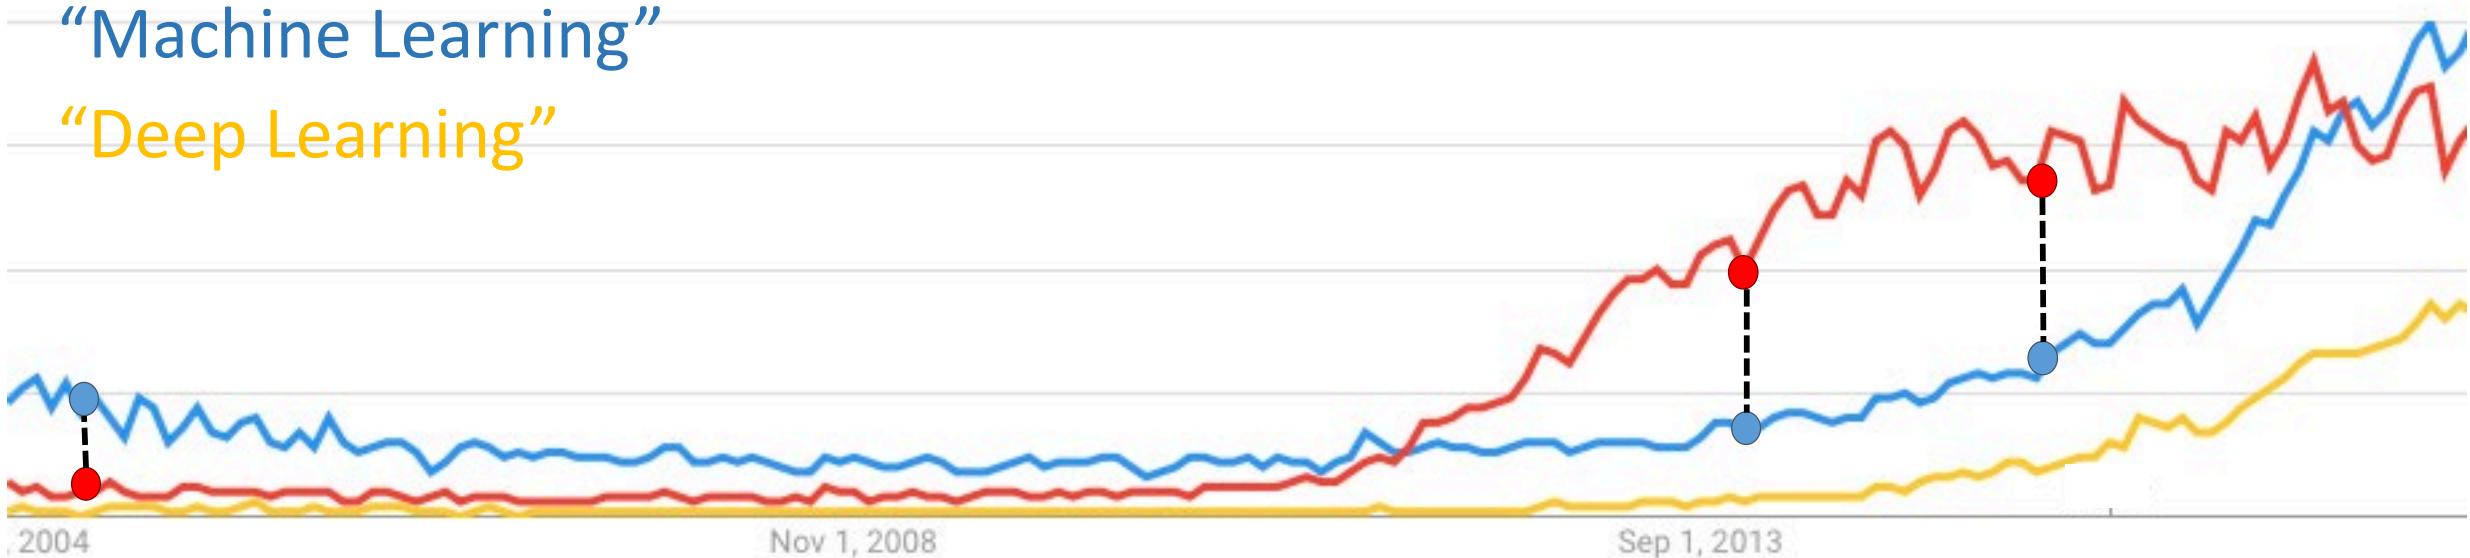


Euclidean Distance

“Big Data”

“Machine Learning”

“Deep Learning”



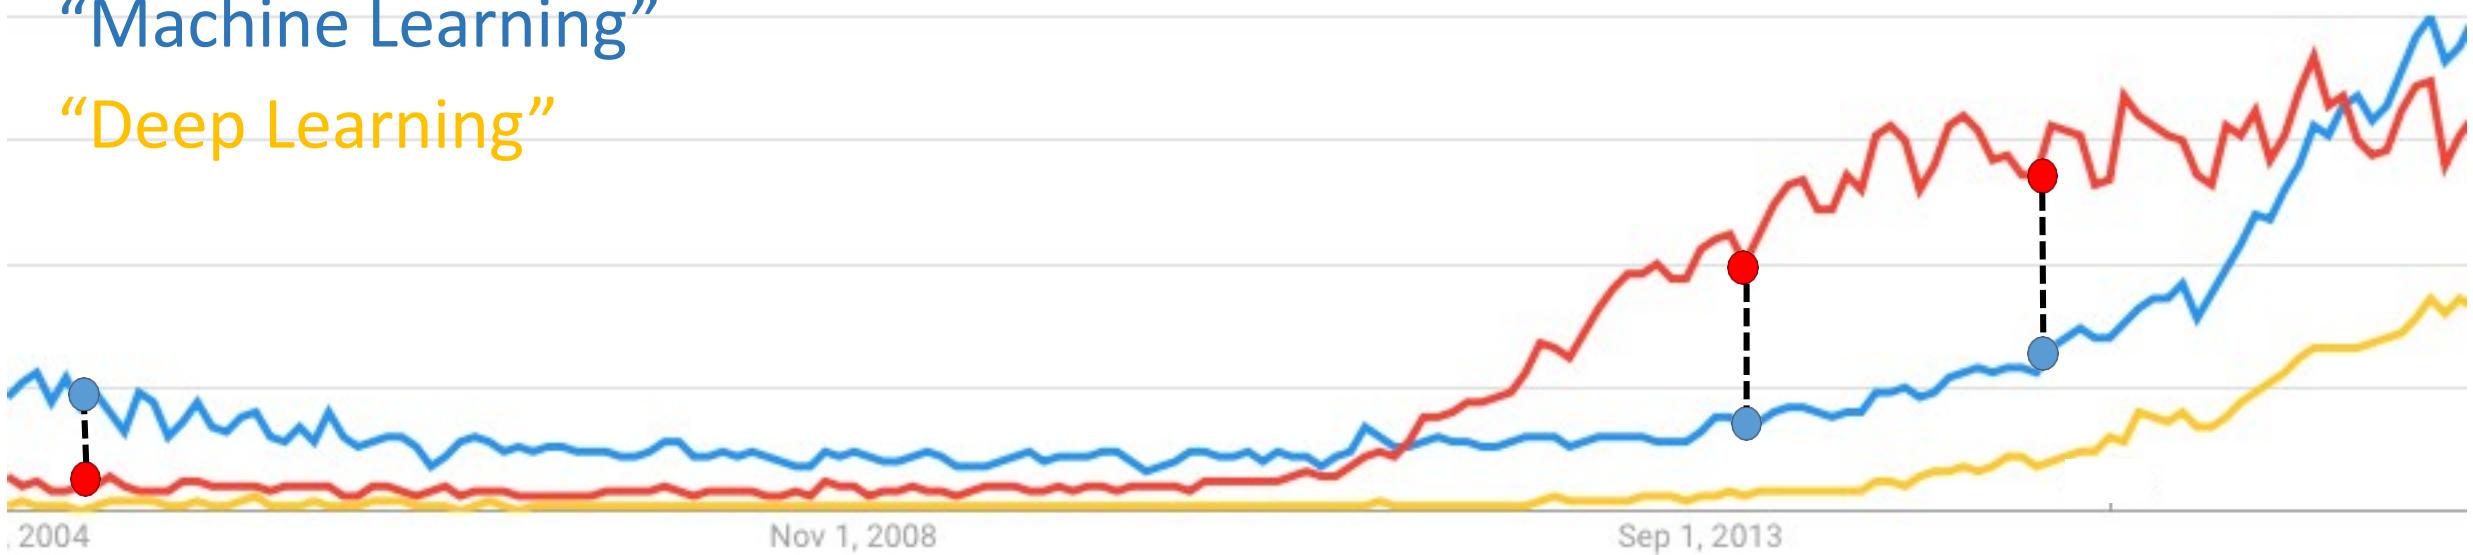
$$\Delta_{Euc}(B, M) = \sqrt{\sum_{i=1..N} (B_i - M_i)^2}$$

Euclidean Distance

“Big Data”

“Machine Learning”

“Deep Learning”



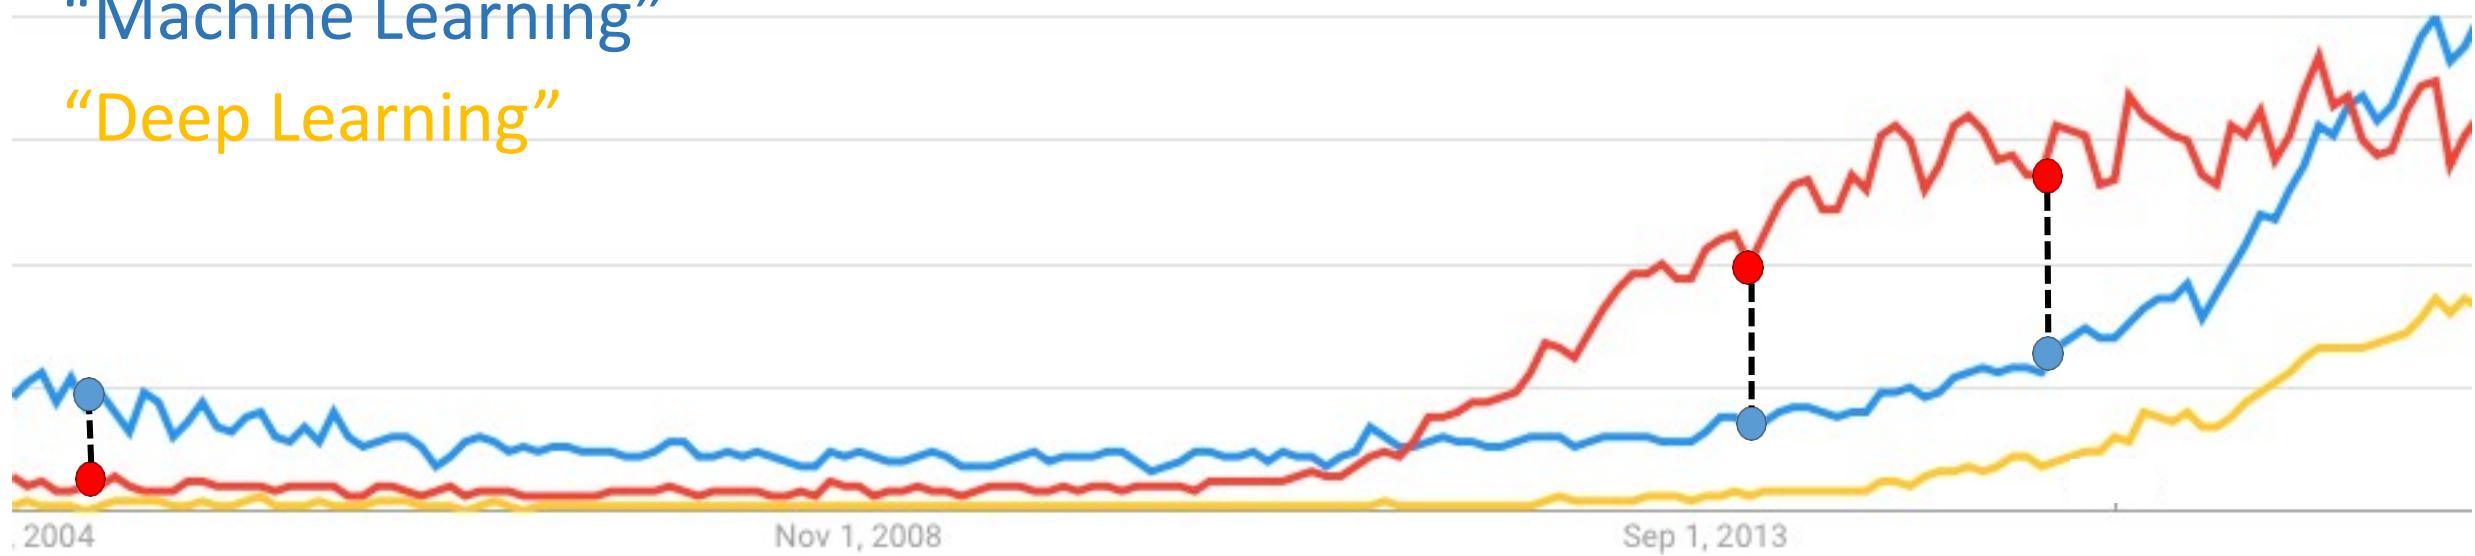
$$\Delta_{Euc}(B, M) = \sqrt{\sum_{i=1..N} (B_i - M_i)^2}$$

Correlation Similarity

“Big Data”

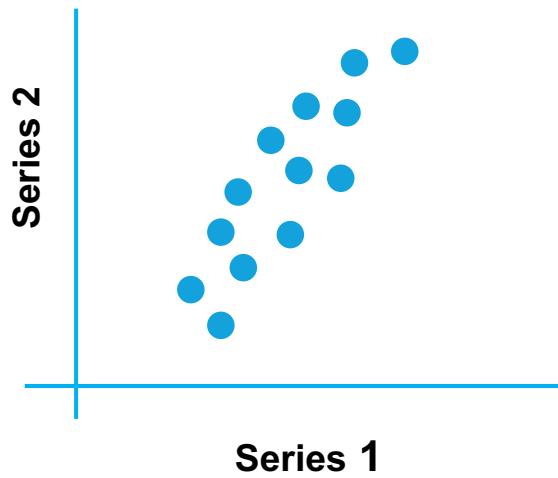
“Machine Learning”

“Deep Learning”

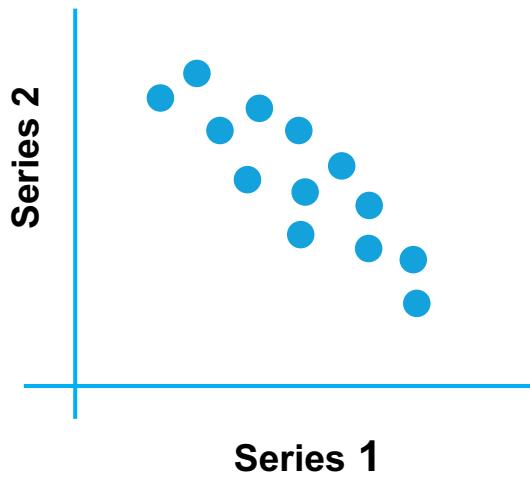


$$\text{Sim}_{correl}(B, M) = \frac{E[(B - \mu_B)(M - \mu_M)]}{\sigma_B \sigma_M}$$

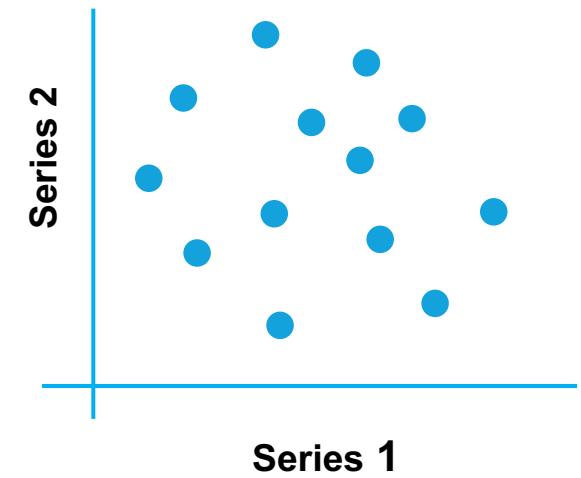
Correlation



positively correlated



negatively correlated



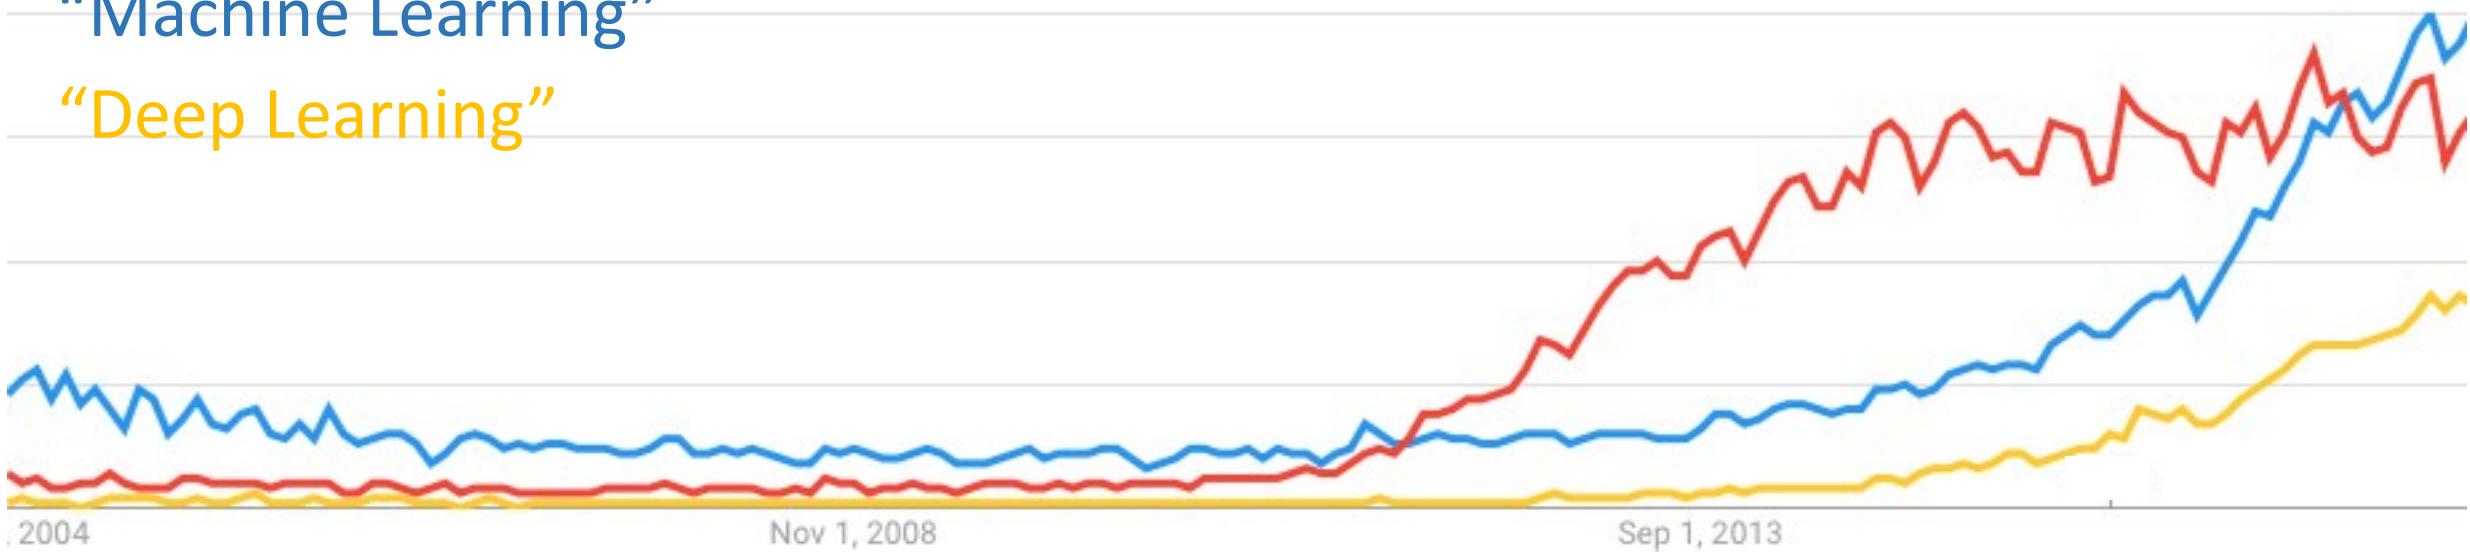
uncorrelated

Comparing time series

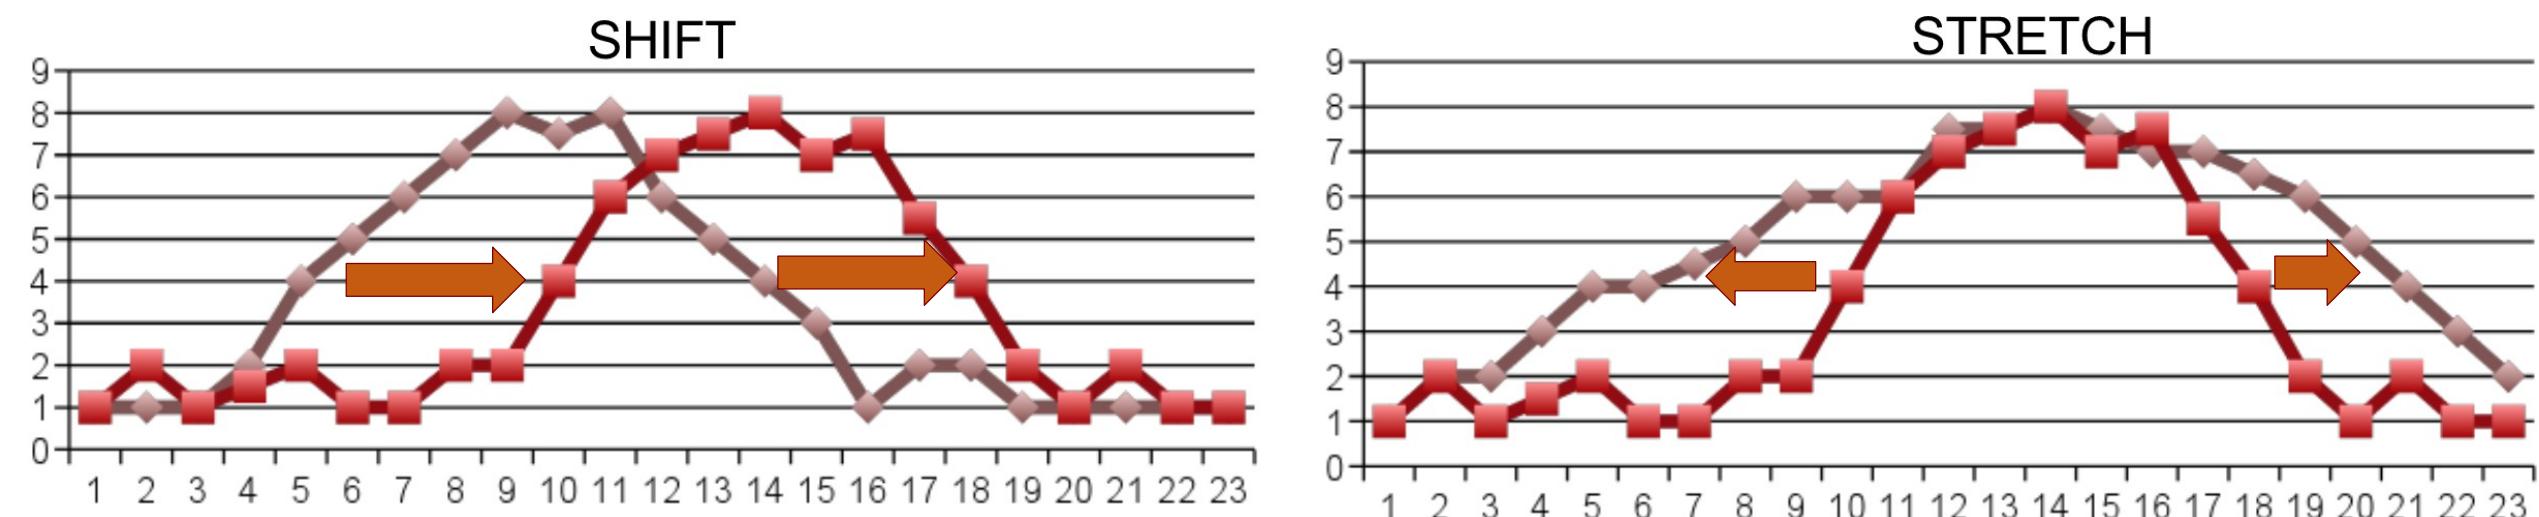
“Big Data”

“Machine Learning”

“Deep Learning”



Issues with Synchronized Measures

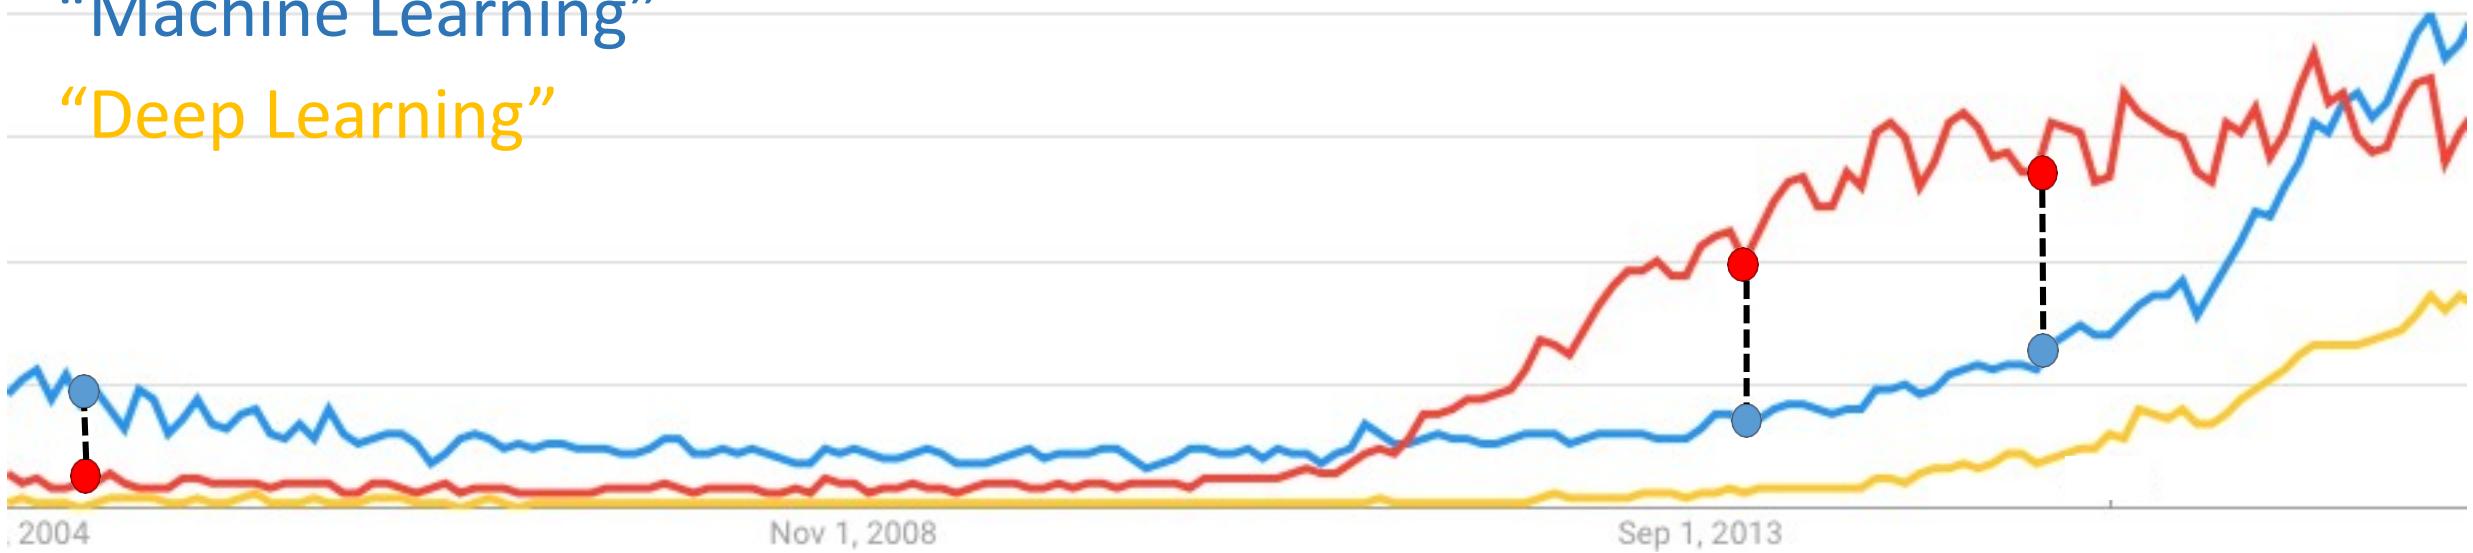


Euclidean Distance

“Big Data”

“Machine Learning”

“Deep Learning”



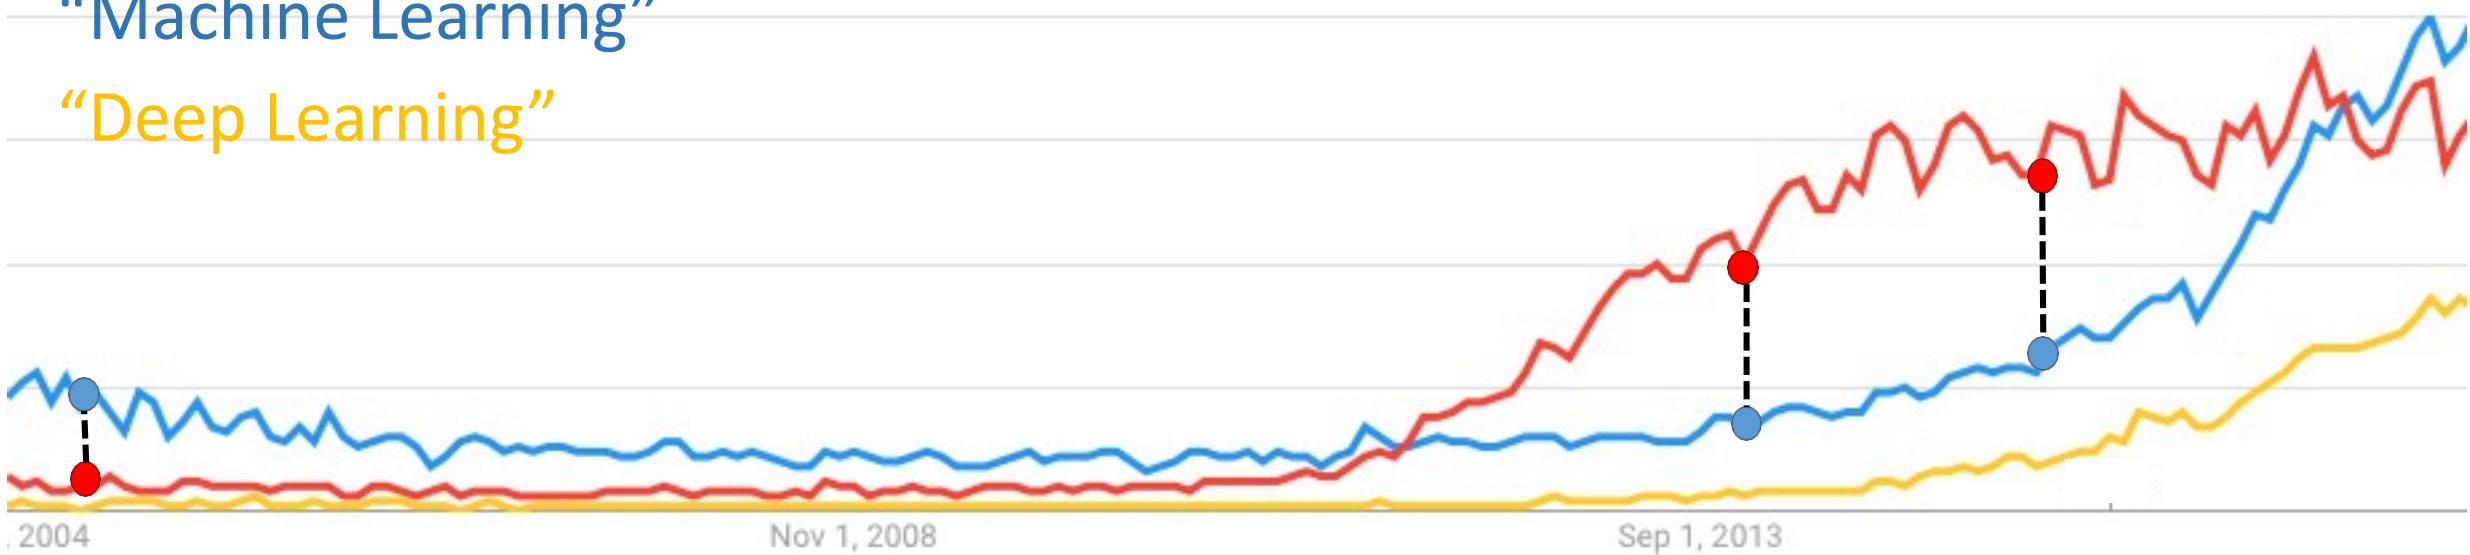
$$\Delta_{Euc}(B, M) = \sqrt{\sum_{i=1..N} (B_i - M_i)^2}$$

Correlation Similarity

“Big Data”

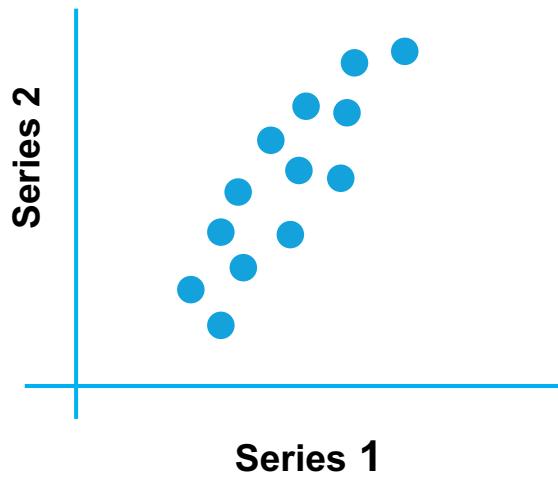
“Machine Learning”

“Deep Learning”

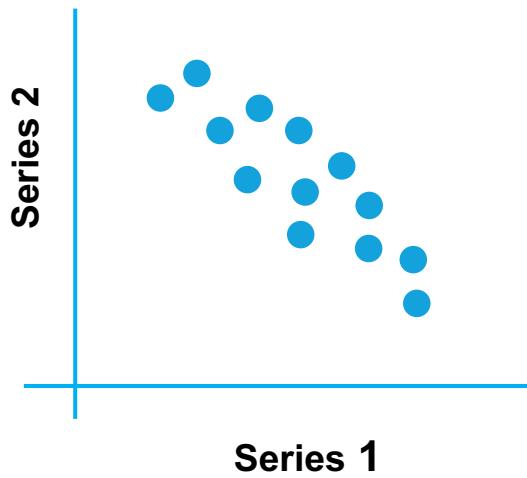


$$\text{Sim}_{correl}(B, M) = \frac{E[(B - \mu_B)(M - \mu_M)]}{\sigma_B \sigma_M}$$

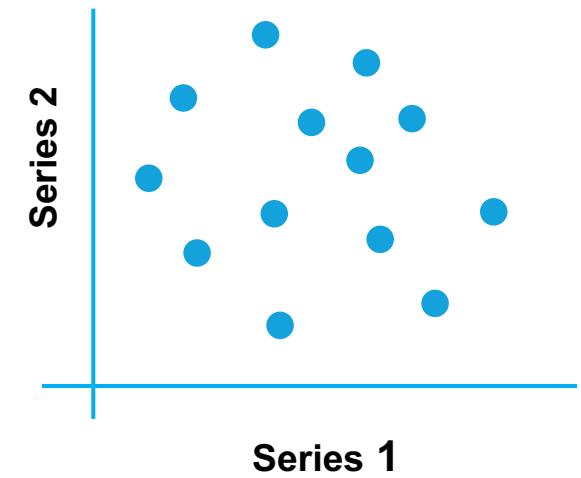
Correlation



positively correlated

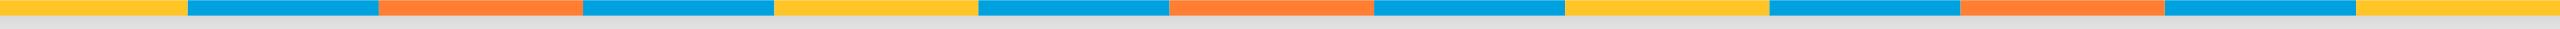


negatively correlated



uncorrelated

Reminder: Edit cost



| Let E be a sequence of edit operations to convert one string to another

| Let us associate a cost, C , to each edit operation

- Costs of edit operations can be different from each other
 - Type of the operation (replace, delete, insert)
 - Symbols involved in the operation
 - Position of the edit operation

| Given a sequence of edit operations, E

$$C(E) = \sum_{e_i \in E} C(e_i)$$

Dynamic Time Warping

- Let us be given two time series, P and Q, of lengths N and M

| $D[i,j] = \# \text{ of edits from length-}i \text{ prefix of } P \text{ to length-}j \text{ prefix of } Q$

- $D[0,j] = \infty; D[i,0] = \infty$
- $D[0,0] = 0$

else $D[i,j] = \text{abs}(P_i - Q_j) + \min \{$

$D[i-1,j],$



G



$D[i,j-1],$



$D[i-1,j-1]$

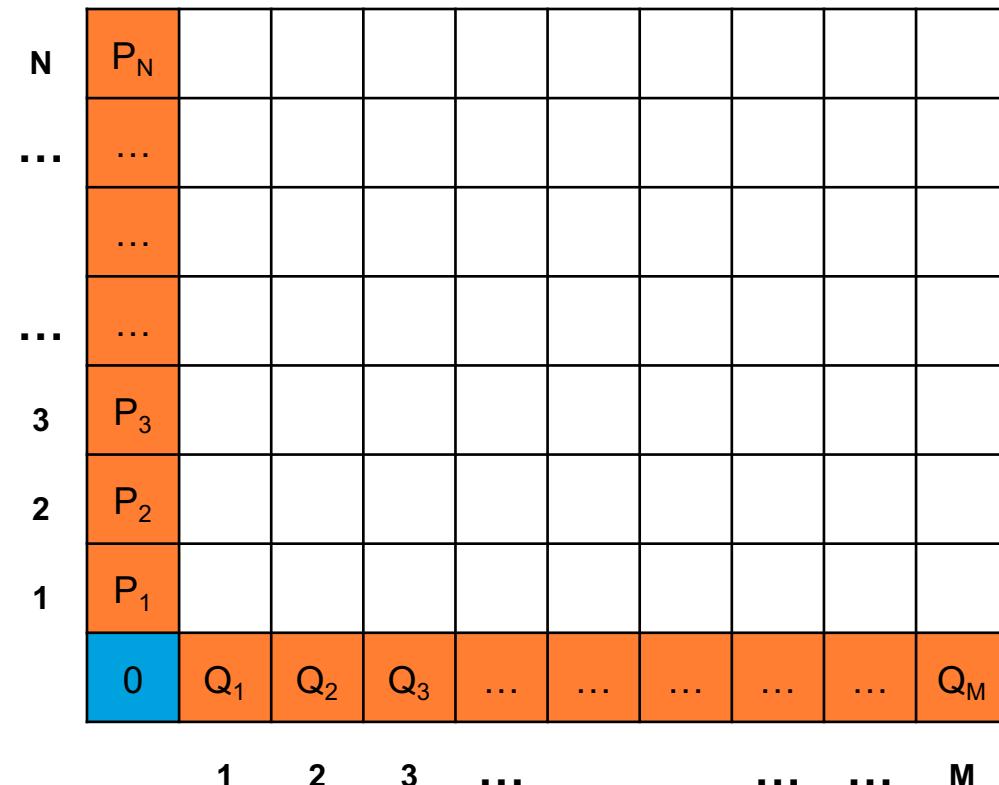


}

Dynamic Time Warping (DTW)

| Complexity: $O(M,N)$

| Dynamic Programming
based implementation

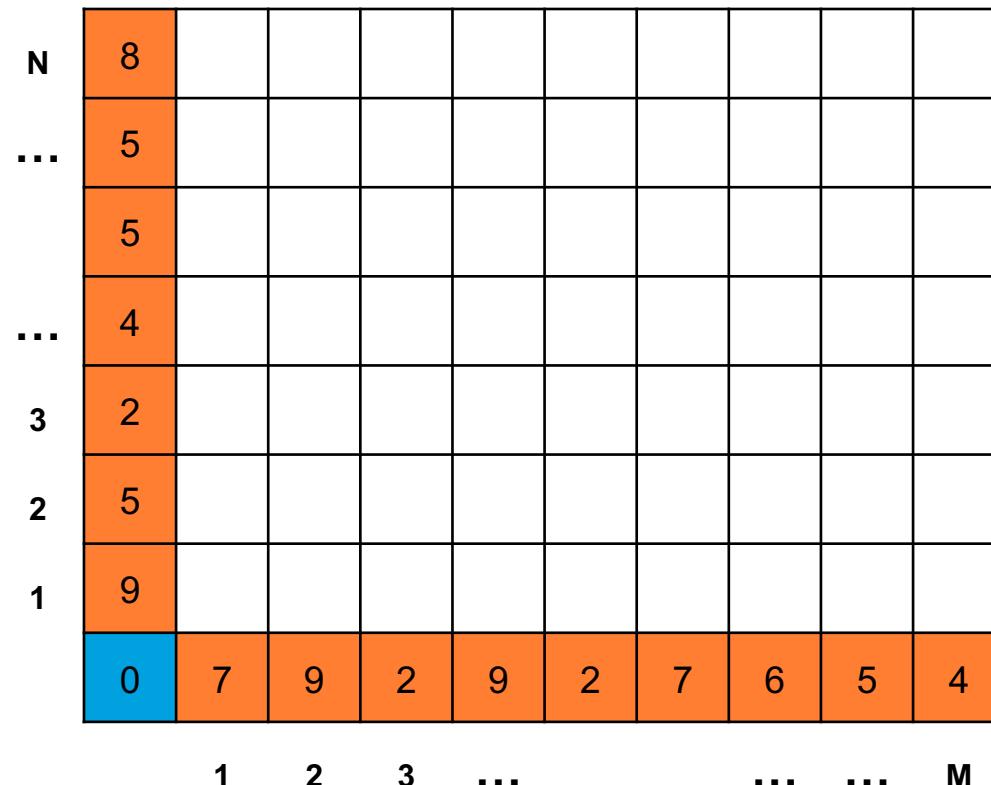


Dynamic Time Warping (DTW)

| Complexity: $O(M,N)$

| Dynamic Programming
based implementation

7	9	2	9	2	7	6	5	4
9	5	2	4	5	5	8		



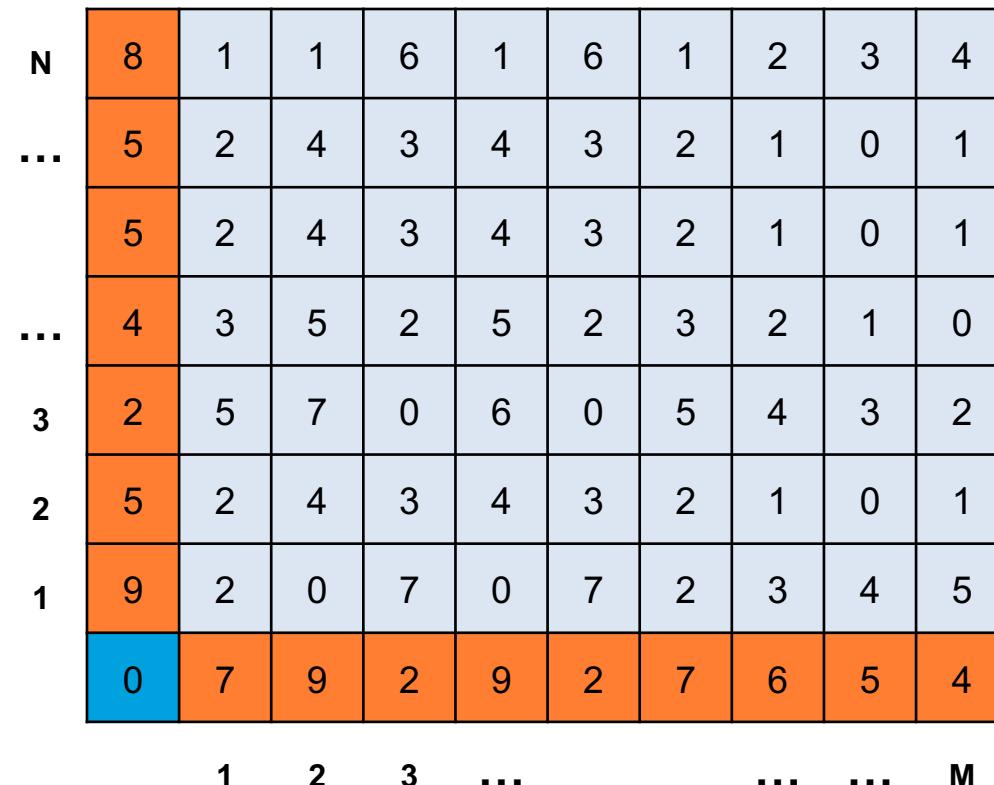
Dynamic Time Warping (DTW)

| Complexity: $O(M,N)$

| Dynamic Programming
based implementation

7	9	2	9	2	7	6	5	4
---	---	---	---	---	---	---	---	---

9	5	2	4	5	5	8
---	---	---	---	---	---	---



Dynamic Time Warping (DTW)

| Complexity: $O(M,N)$

| Dynamic Programming
based implementation

7	9	2	9	2	7	6	5	4
9	5	2	4	5	5	8		

N	8	17	1	6	1	6	1	2	3	4
...	5	16	4	3	4	3	2	1	0	1
...	5	14	4	3	4	3	2	1	0	1
...	4	12	5	2	5	2	3	2	1	0
3	2	9	7	0	6	0	5	4	3	2
2	5	4	4	3	4	3	2	1	0	1
1	9	2	2	9	9	16	18	21	25	30
	0	7	9	2	9	2	7	6	5	4
	1	2	3	M				

Dynamic Time Warping (DTW)

| Complexity: $O(M,N)$

| Dynamic Programming
based implementation

7	9	2	9	2	7	6	5	4
9	5	2	4	5	5	8		

N	8	17	1	6	1	6	1	2	3	4
...	5	16	4	3	4	3	2	1	0	1
...	5	14	4	3	4	3	2	1	0	1
...	4	12	5	2	5	2	3	2	1	0
3	2	9	7	0	6	0	5	4	3	2
2	5	4	6	3	4	3	2	1	0	1
1	9	2	2	9	9	16	18	21	25	30
	0	7	9	2	9	2	7	6	5	4
	1	2	3	M				

Dynamic Time Warping (DTW)

| Complexity: $O(M,N)$

| Dynamic Programming
based implementation

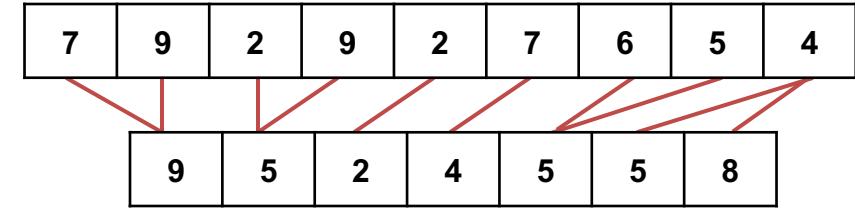
7	9	2	9	2	7	6	5	4
9	5	2	4	5	5	8		

N	8	17	17	19	14	20	15	16	17	18
...	5	16	18	13	14	14	15	14	14	14
...	5	14	16	10	11	13	13	13	13	14
...	4	12	14	7	10	11	12	14	15	15
3	2	9	11	5	11	9	14	18	18	20
2	5	4	6	5	9	12	14	15	15	16
1	9	2	2	9	9	16	18	21	25	30
0	7	9	2	9	2	7	6	5	4	
	1	2	3	M	

Dynamic Time Warping (DTW)

| Complexity: $O(M,N)$

| Dynamic Programming
based implementation

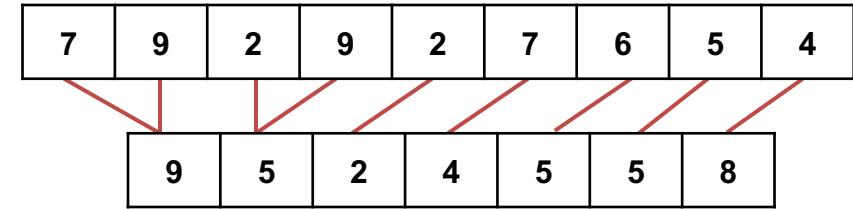


N	8	17	17	19	14	20	15	16	17	18
...	5	16	18	13	14	14	15	14	14	14
...	5	14	16	10	11	13	13	13	13	14
...	4	12	14	7	10	11	12	14	15	15
3	2	9	11	5	11	9	14	18	18	20
2	5	4	6	5	9	12	14	15	15	16
1	9	2	2	9	9	16	18	21	25	30
	0	7	9	2	9	2	7	6	5	4
	1	2	3	M	

Dynamic Time Warping (DTW)

| Complexity: $O(M,N)$

| Dynamic Programming
based implementation

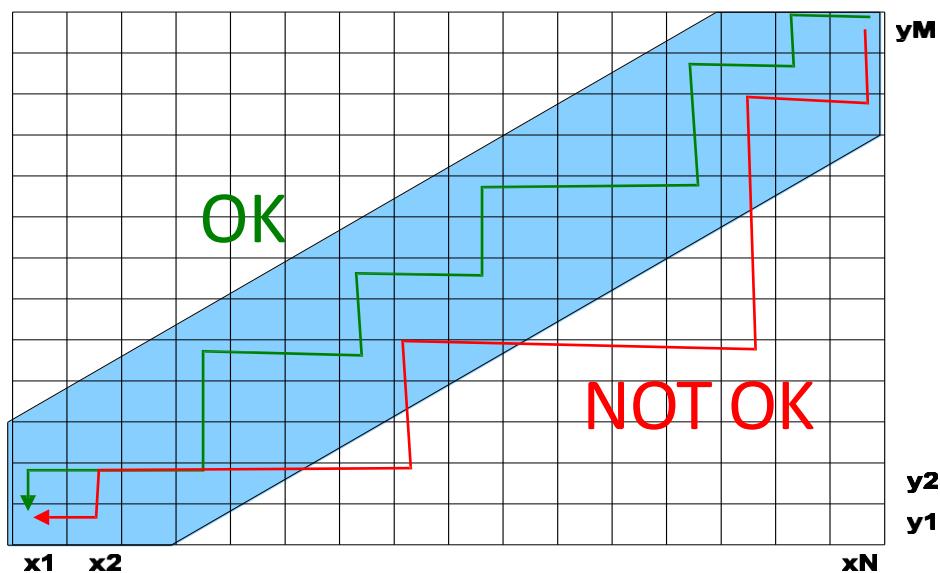


N	8	17	17	19	14	20	15	16	17	18
...	5	16	18	13	14	14	15	14	14	14
...	5	14	16	10	11	13	13	13	13	14
...	4	12	14	7	10	11	12	14	15	15
3	2	9	11	5	11	9	14	18	18	20
2	5	4	6	5	9	12	14	15	15	16
1	9	2	2	9	9	16	18	21	25	30
	0	7	9	2	9	2	7	6	5	4
	1	2	3	M	

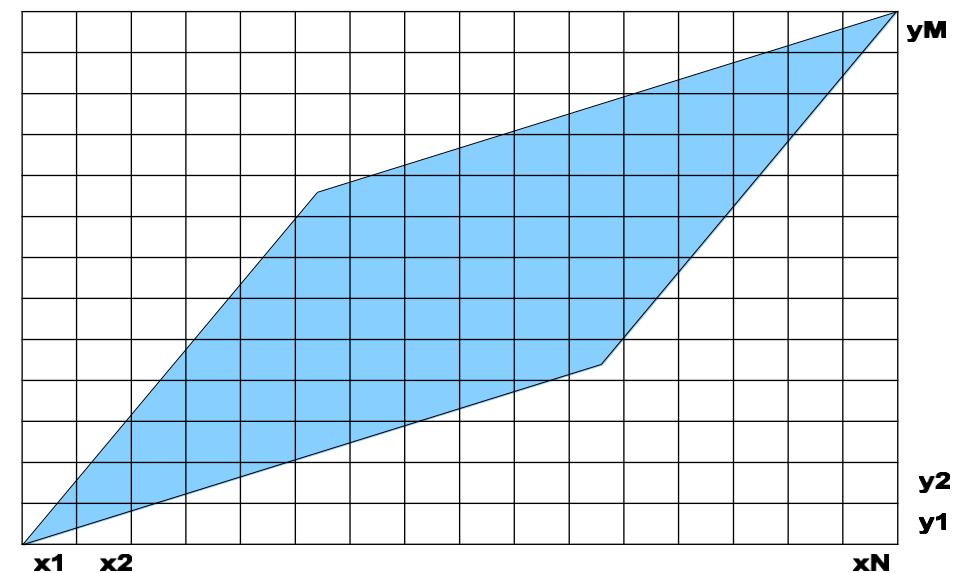
Reducing the Cost of DTW

To reduce the $O(NM)$ cost of filling the grid various heuristics impose **constraints** on the grid regions through which the warp paths can pass.

Sakoe-Chiba band [1]



Itakura parallelogram [2]

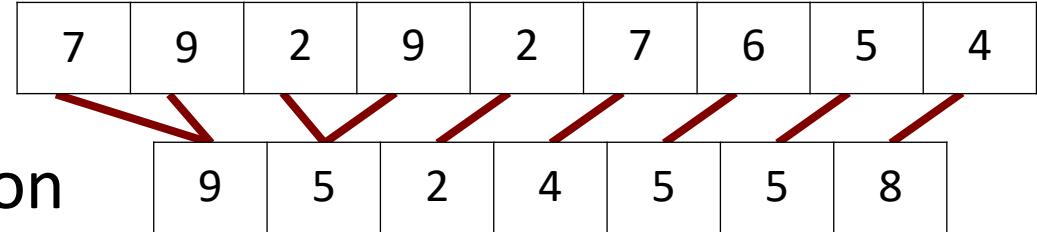


1 Dynamic Programming Algorithm Optimisation for Spoken Word Recognition, 1978

2 F. Itakura. Minimum prediction residual principle applied to speech recognition, 1975

Dynamic Time Warping (DTW)

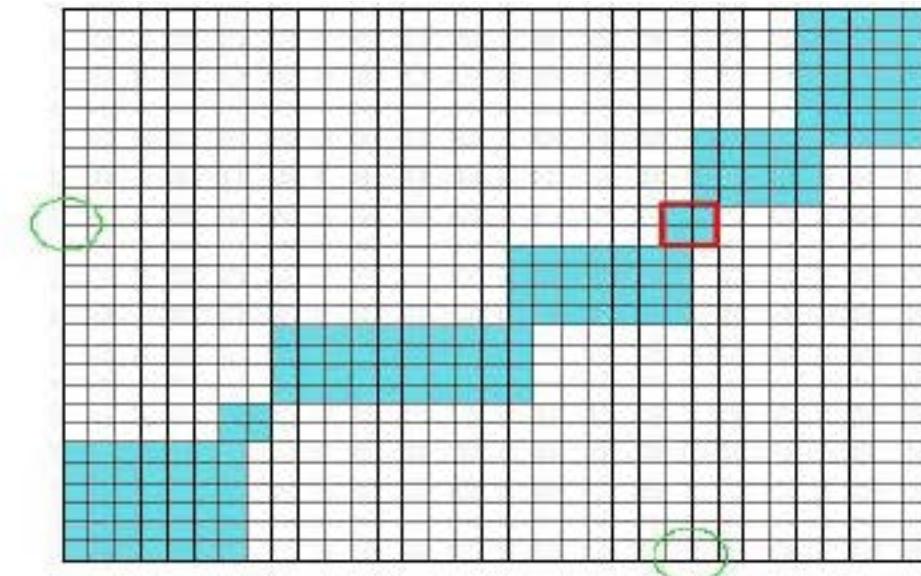
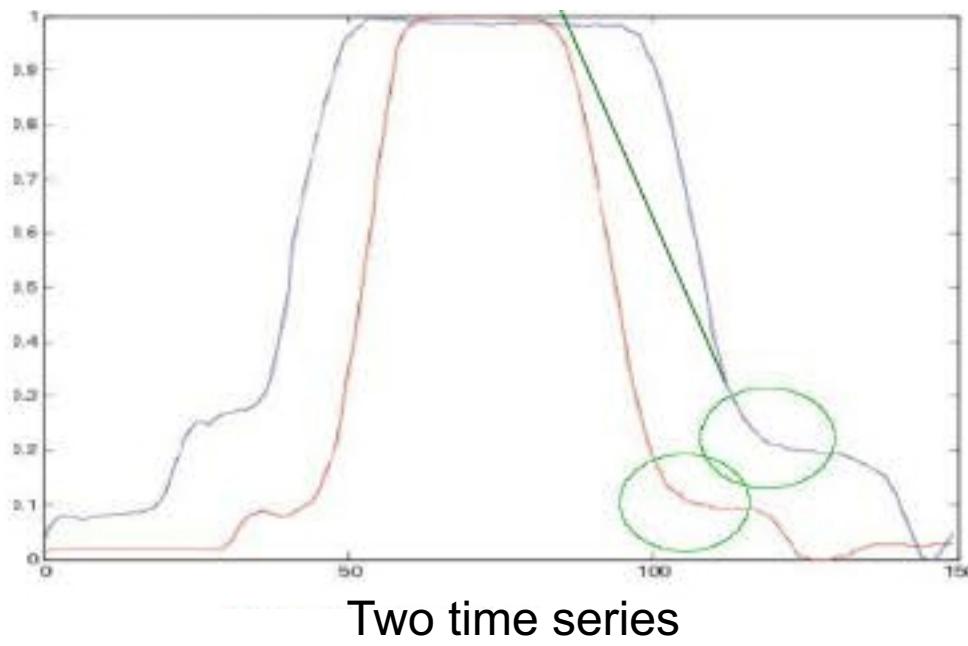
- Complexity: $O(M,N)$
- Dynamic Programming based implementation



N	8	17	17	19	14	20	15	16	17	18
...	5	16	18	13	14	14	15	14	14	14
...	5	14	16	10	11	13	13	13	13	14
...	4	12	14	7	10	11	12	14	15	15
3	2	9	11	5	11	9	14	18	18	20
2	5	4	6	5	9	12	14	15	15	16
1	9	2	2	9	9	16	18	21	25	30
0	7	9	2	9	2	7	6	5	4	
	1	2	3	M		

Time series often carry **temporal features** that can be used for identifying **locally relevant constraints** to eliminate redundant work in an **adaptive** manner.

Sample (matching) structural features of two time series



Adaptive constraints on the DTW grid

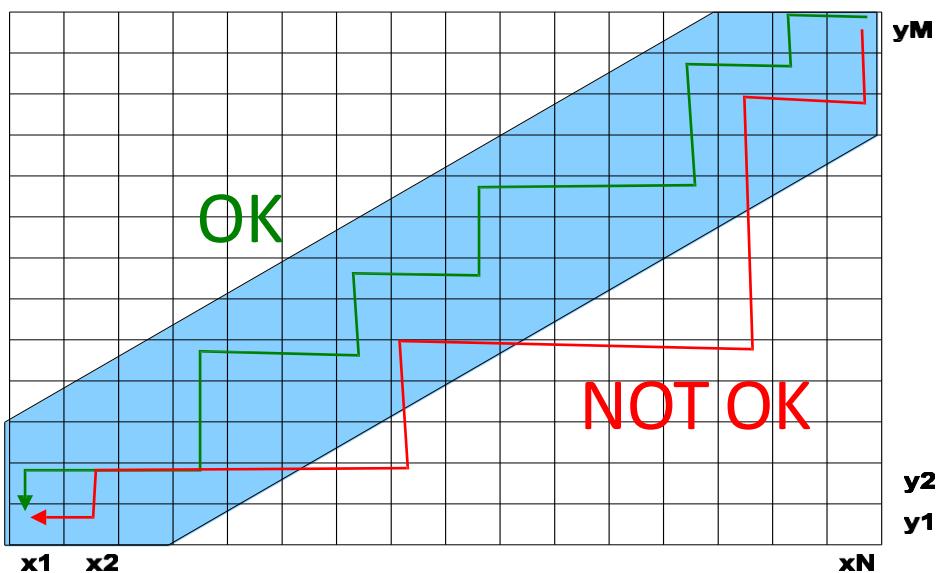
1 Dynamic Programming Algorithm Optimisation for Spoken Word Recognition, 1978

2 F. Itakura. Minimum prediction residual principle applied to speech recognition, 1975

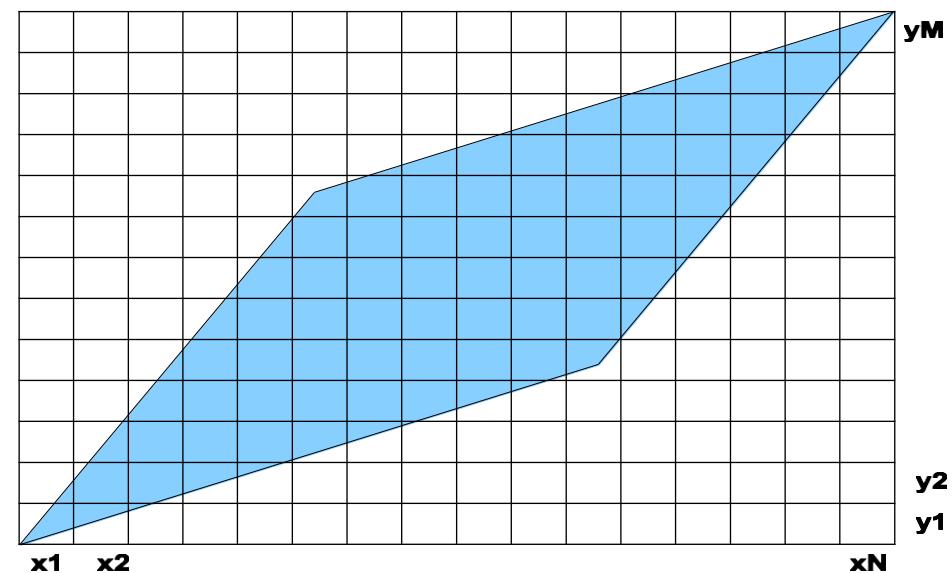
Reducing the Cost of DTW

To reduce the $O(NM)$ cost of filling the grid various heuristics impose **constraints** on the grid regions through which the warp paths can pass.

Sakoe-Chiba band [1]



Itakura parallelogram [2]



1 Dynamic Programming Algorithm Optimisation for Spoken Word Recognition, 1978

2 F. Itakura. Minimum prediction residual principle applied to speech recognition, 1975

Time series are similar to sequences

| A **string or sequence**, $S = (c_1, c_2, \dots c_N)$, is a finite sequence of symbols.

abcbbbaabbaabcbbbbaabbc

| A **time series**, $T = (d_1, d_2, \dots, d_N)$, is a finite sequence of data values.



<https://trends.google.com/trends/explore?q=bg%20data>

SAX (Symbolic Aggregate AppRoXimation)

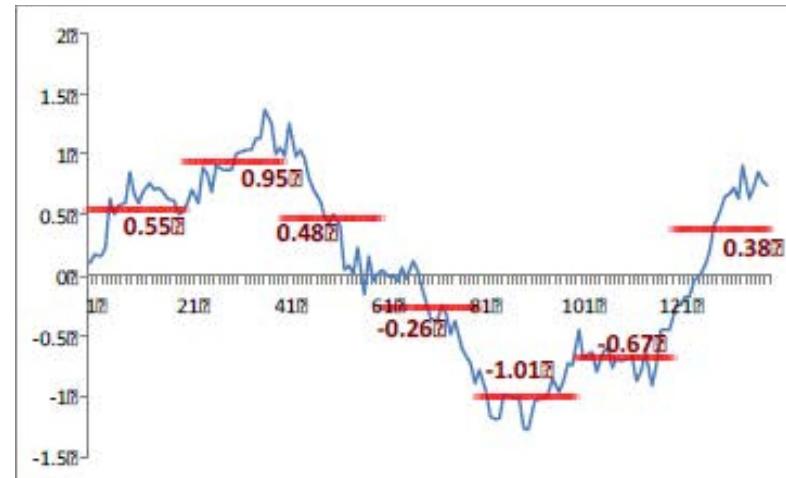
- | Time series are similar to sequences
- | Can we transform a time series into a compact sequence representation?



SAX (Symbolic Aggregate AppRoXimation)

- | Time series are similar to sequences
- | Transform a time series into a compact sequence representation

- Divide the time series into w-length (non-overlapping) windows
- For each window,
 - compute the average amplitude

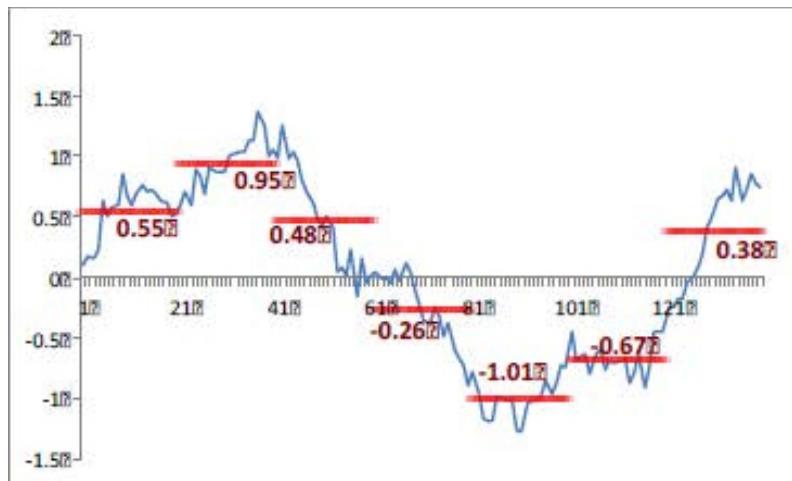


SAX (Symbolic Aggregate AppXimation)

| Time series are similar to sequences

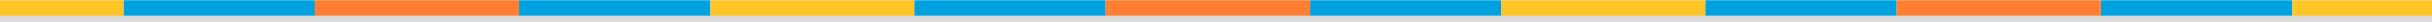
| Transform a time series into a compact sequence representation

- Divide the time series into w-length (non-overlapping) windows
- For each window,
 - compute the average amplitude
 - assign a symbol from a dictionary with s symbols representing this average amplitude



Range	Symbol
[0.9,2]	A
[0.3,0.9]	B
[0.0,0.3]	C
[0.0,-0.3]	D
[-0.3,-0.9]	E
[-0.9,-2]	F

SAX (Symbolic Aggregate AppXimation)



| Time series are similar to sequences

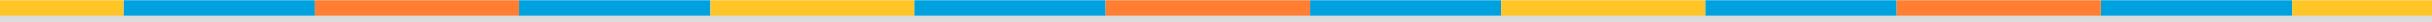
| Transform a time series into a compact sequence representation

- Divide the time series into w-length (non-overlapping) windows
- For each window,
 - compute the average amplitude
 - assign a symbol from a dictionary with s symbols representing this average amplitude

| Note that, SAX reduces

- **Temporal resolution**, by dividing the string into windows ($\text{length} \sim N/w$)
- **Amplitude resolution**, by using only one of the s symbols per window

Summary



- Synchronized measures (Euclidean distance, correlation) are relatively cheap to compute, but may not account for shifts, stretches, and other temporal misalignments
- Asynchronous approaches (edit distance, DTW) can account for misalignments, but are expensive. Solution include
 - constraining the search space
 - creation of reduced representations of the time series



Sequences and Time Series

Time Series Motifs

K. Selçuk Candan, Professor of Computer Science and Engineering

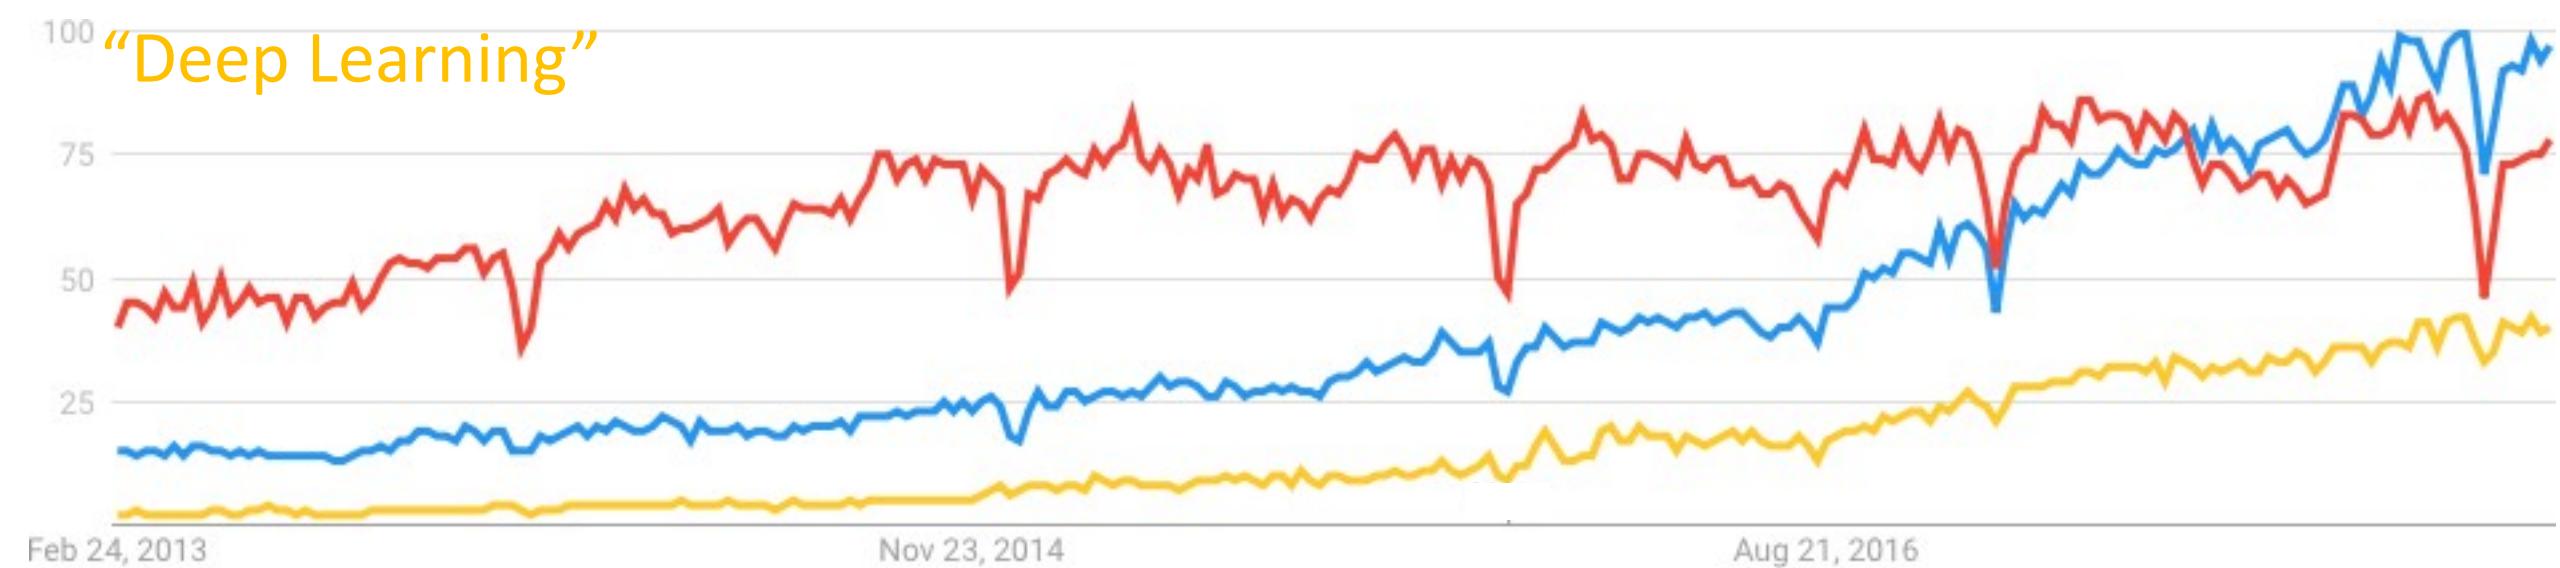


Motifs

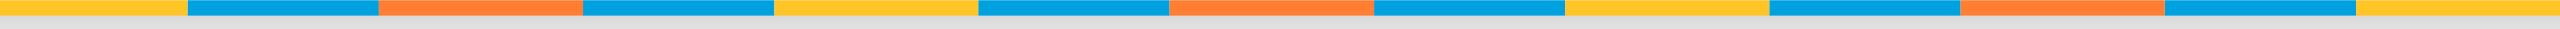
“Big Data”

“Machine Learning”

“Deep Learning”



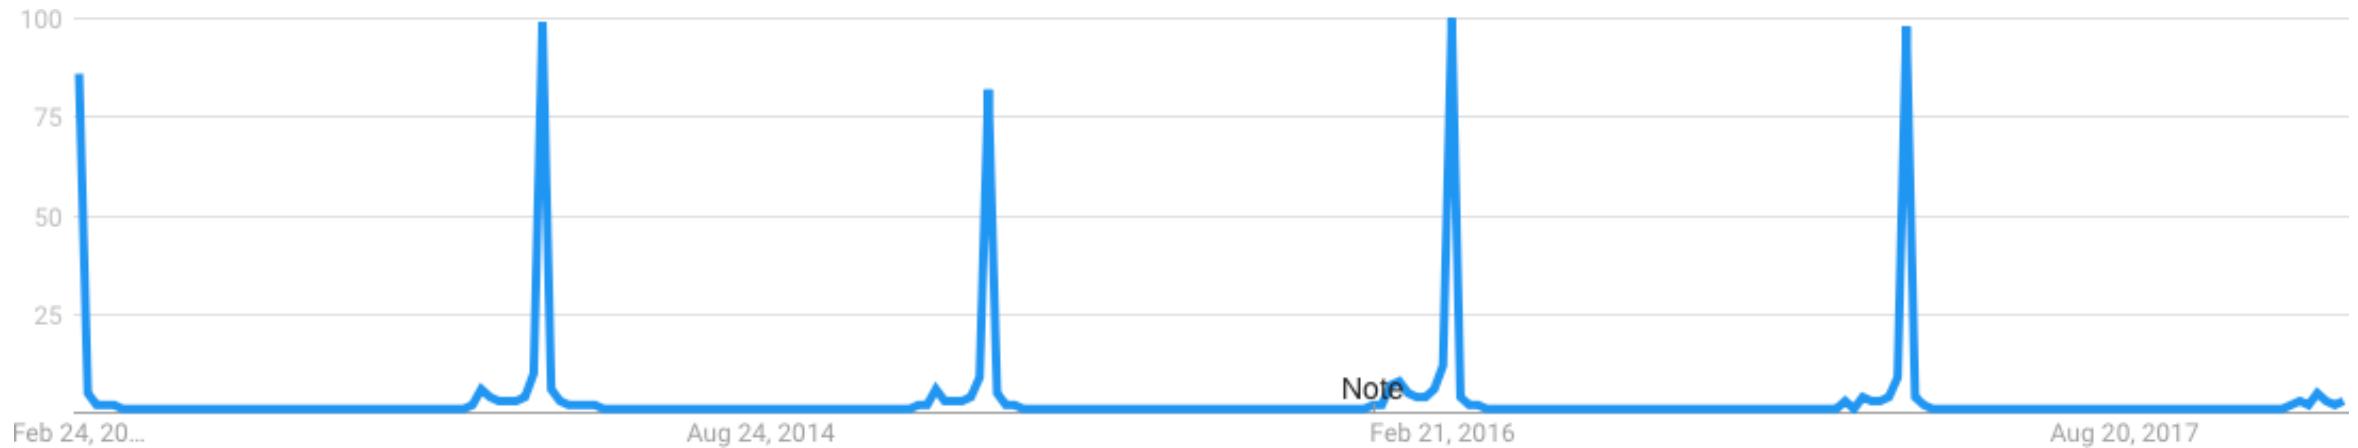
Motif Search Algorithm



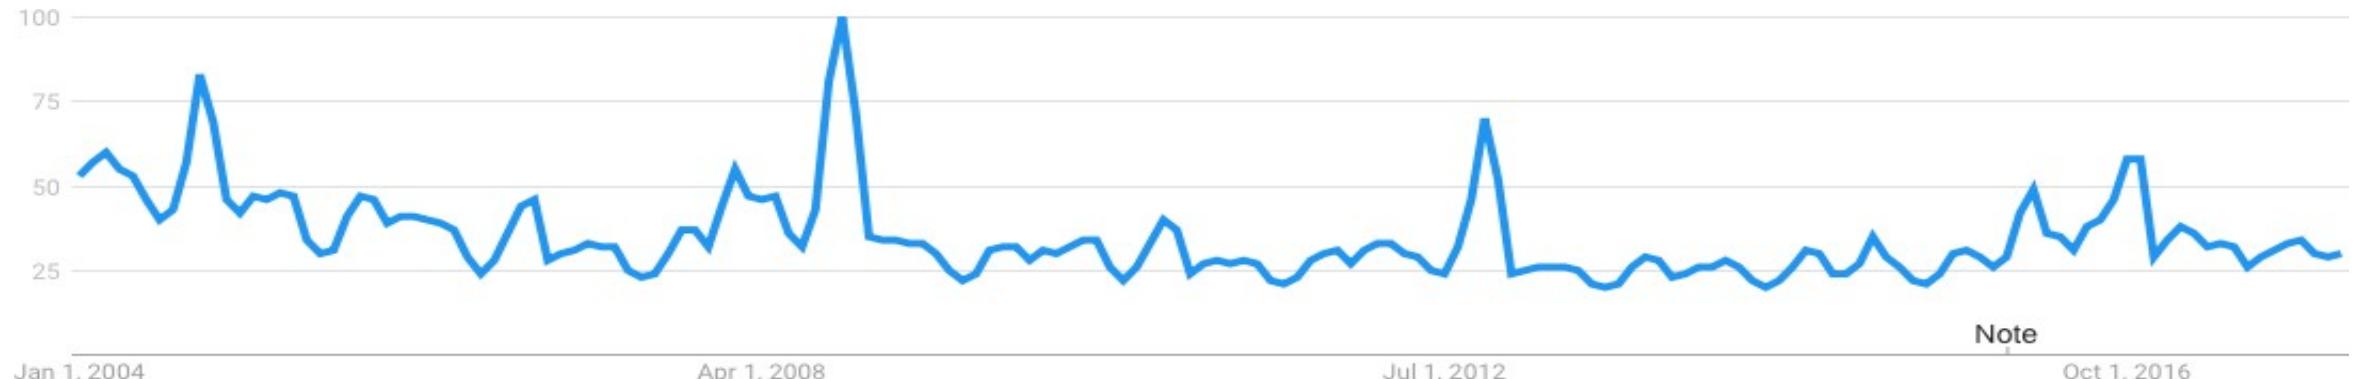
- General approach
 - Enumerate subsequences of the time series (with a given length)
 - Apply a clustering algorithm to identify groups of similar subsequences

Motif Quality

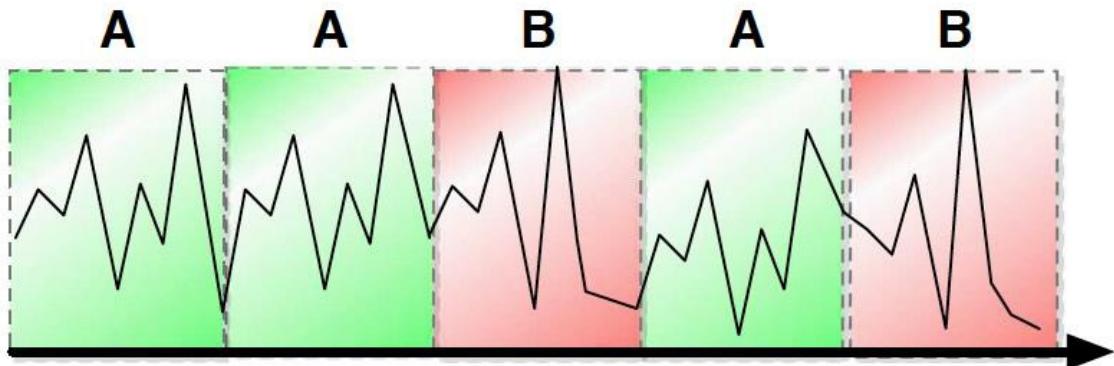
“Oscars”



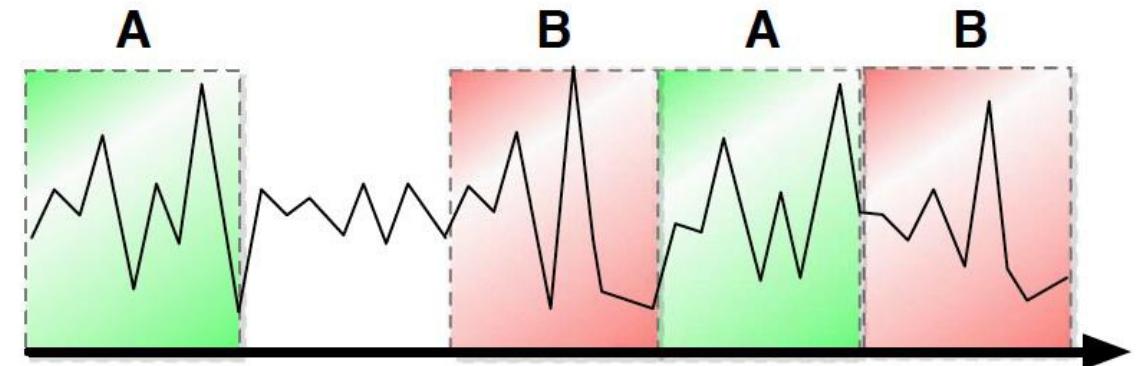
“Politics”



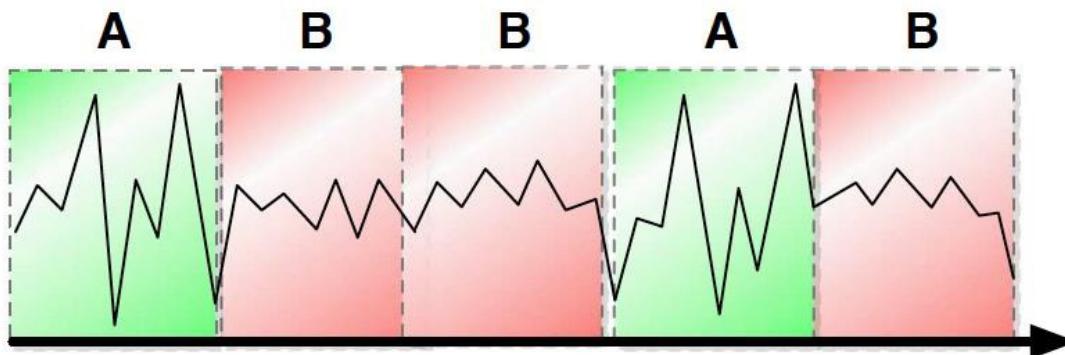
Motifs



(a) motif *A* has a higher support

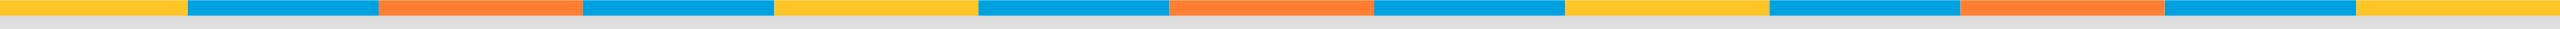


(b) motif *A* has smaller distortion/higher purity



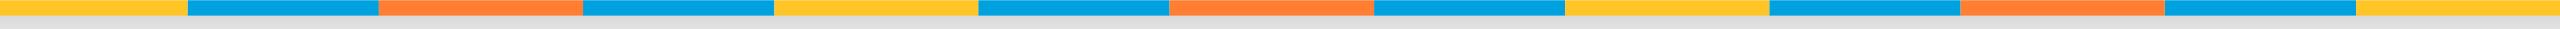
(c) motif *A* has higher dynamicity

Motif Search Algorithm



- General approach
 - Enumerate subsequences of the time series (with a given length)
 - Apply a clustering algorithm to identify groups of similar subsequences
 - Eliminate clusters with
 - too few subsequences (not enough support)
 - too imperfect matches (not well defined motifs)
 - too low dynamicity (patterns that do not carry much useful information or are not interesting)

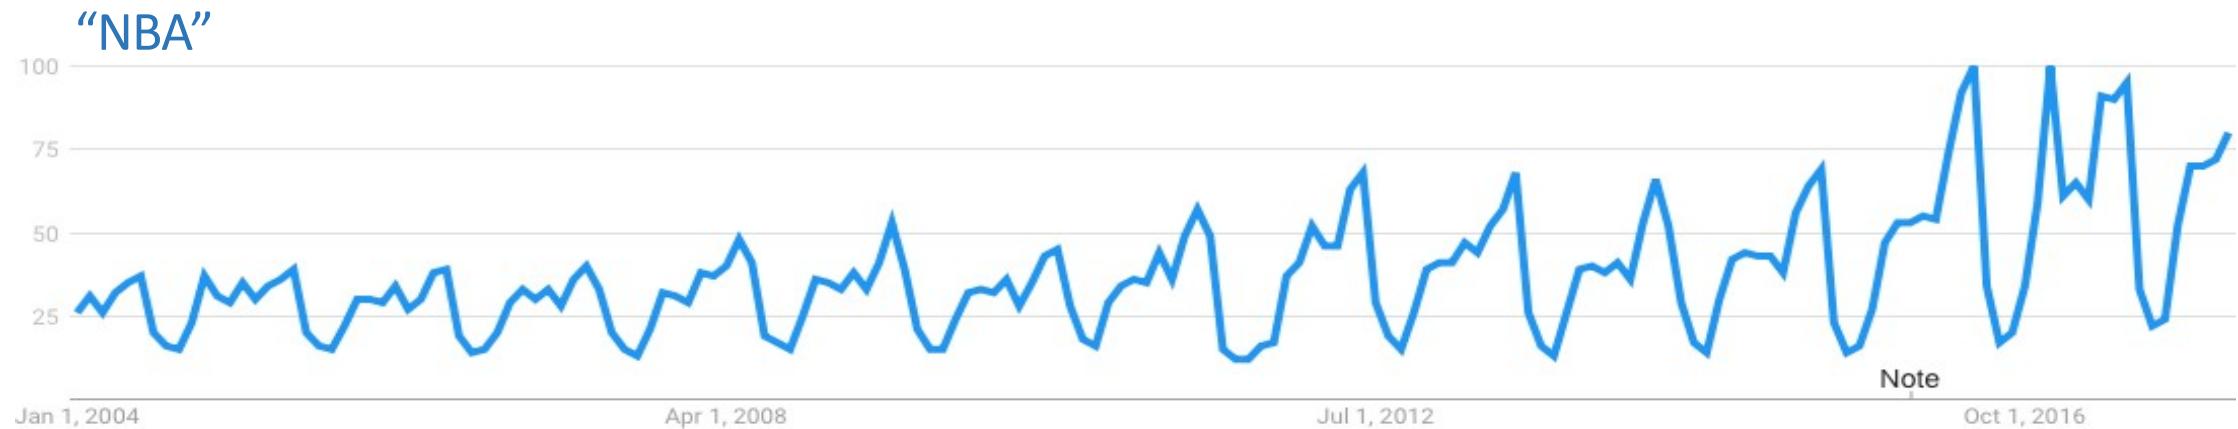
Motif Search Algorithm



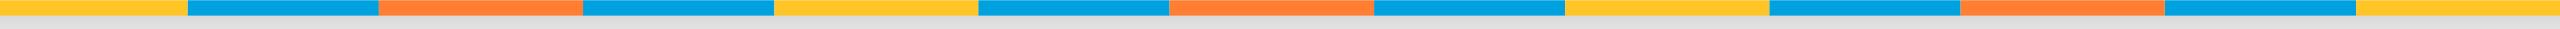
- General approach
 - Enumerate subsequences of the time series (with a given length)
 - all subsequences
 - subsequences that are sufficiently different from their immediate neighborhood
 - subsequences that have high dynamicity
 - subsequences that have at least one other similar subsequence in the series
 - Apply a clustering algorithm to identify groups of similar subsequences
 - Eliminate clusters with
 - too few subsequences (not enough support)
 - too imperfect matches (not well defined motifs)
 - too low dynamicity (patterns that do not carry much useful information or are not interesting)

Other questions

- Do all instances of the motif need to be of the same (average) amplitude?



Motif Search Algorithm

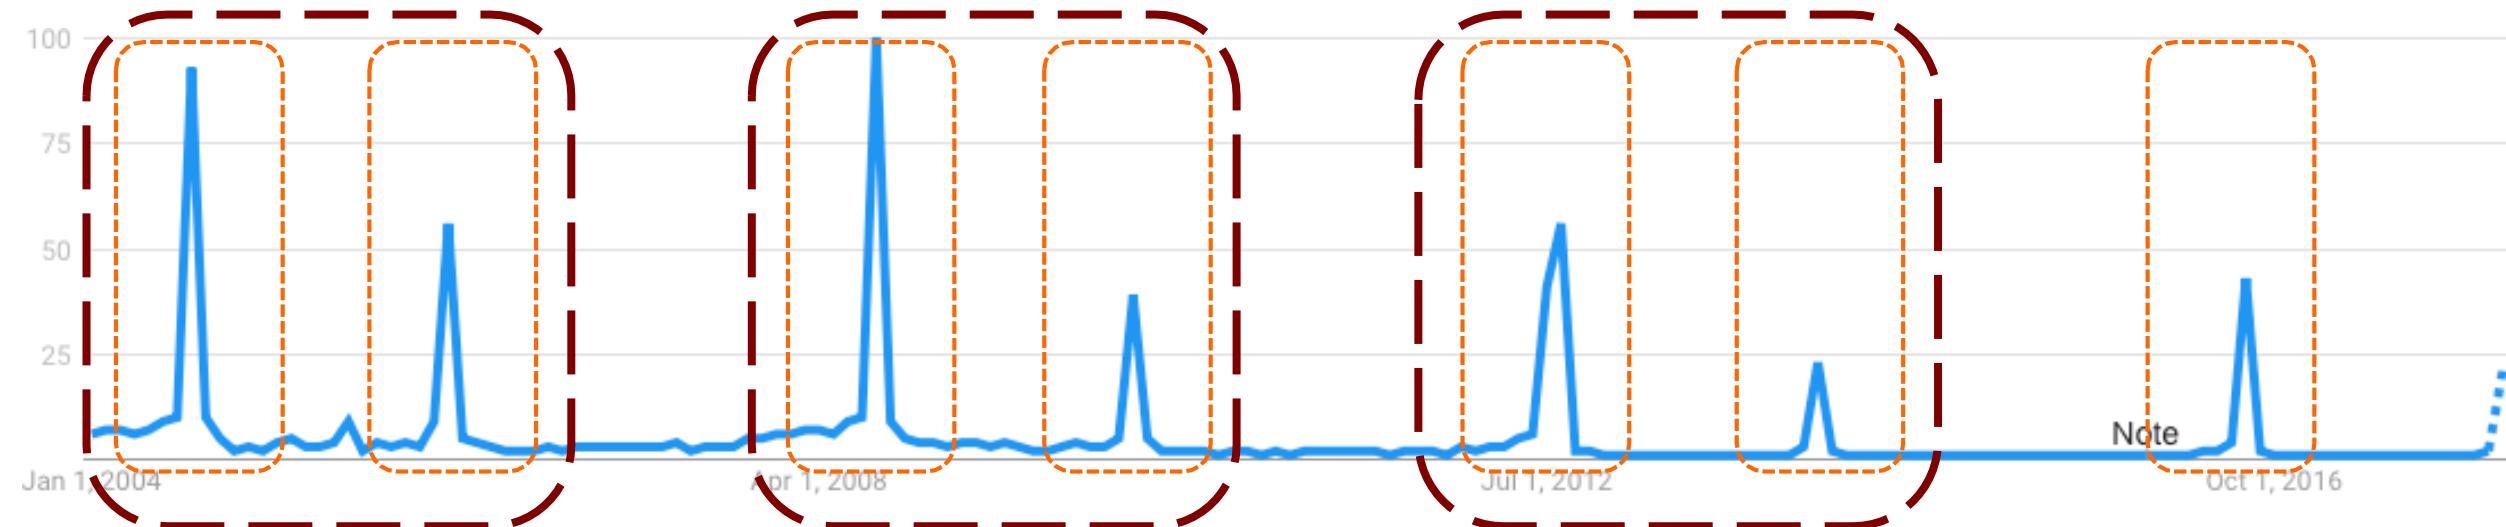


- General approach
 - Enumerate subsequences of the time series (with a given length)
 - all subsequences
 - subsequences that are sufficiently different from their immediate neighborhood
 - subsequences that have high dynamicity
 - subsequences that have at least one other similar subsequence in the series
 - Apply a clustering algorithm to identify groups of similar subsequences
 - allow for **amplitude differences**
 - Eliminate clusters with
 - too few subsequences (not enough support)
 - too imperfect matches (not well defined motifs)
 - too low dynamicity (patterns that do not carry much useful information or are not interesting)

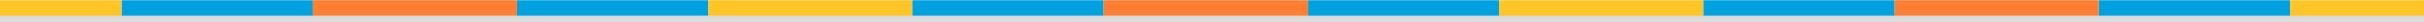
Other questions

- What is the motif length?
- Do all instances of a motif need to have the same length?

“Olympics”



Motif Search Algorithm



- General approach
 - Enumerate subsequences of the time series (**with different lengths**)
 - all subsequences
 - subsequences that are sufficiently different from their immediate neighborhood
 - subsequences that have high dynamicity
 - subsequences that have at least one other similar subsequence in the series
 - Apply a clustering algorithm to identify groups of similar subsequences
 - allow for amplitude differences
 - **account for length differences**
 - Eliminate clusters with
 - too few subsequences (not enough support)
 - too imperfect matches (not well defined motifs)
 - too low dynamicity (patterns that do not carry much useful information or are not interesting)

Summary

- Motif search
 - Divide series into sub-sequences
 - Eliminate un-interesting sub-sequences
 - Apply clustering algorithm
 - Eliminate un-interesting clusters
- Challenge
 - There are many sub-sequences
 - How to define what is an interesting sub-sequence?
 - How to define what is an interesting cluster?